

A Taxonomy of Communication Networks

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

09/16/2021

Outline

- *Admin. and recap*
- A brief introduction to the Internet:
 - present
- Challenges of Internet networks and apps
- A taxonomy of communication networks

Admin.

- If you haven't filled out the survey, please go to the class website to do so

Recap

- A protocol defines the **format** and the **order** of messages exchanged between two or more communicating entities, as well as the **actions** taken on the transmission or receipt of a message or other **events**.
- Key Internet milestones and their implications:
 - ARPANET is sponsored by ARPA → design should survive failures
 - The initial IMPs (routers) were made by a small company → keep the network simple
 - Many networks → internetworking: need a network to connect networks
 - Commercialization → architecture supporting decentralized, autonomous systems

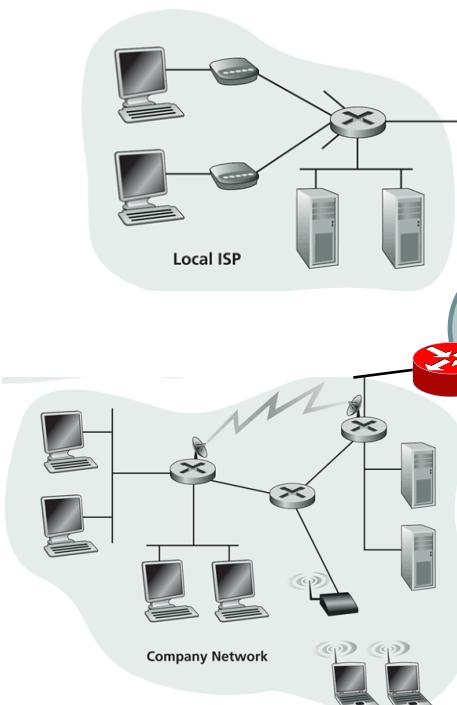
Outline

- Admin. and recaps
- *A brief introduction to the Internet*
 - past
 - *present*

Internet Physical Infrastructure

Residential access

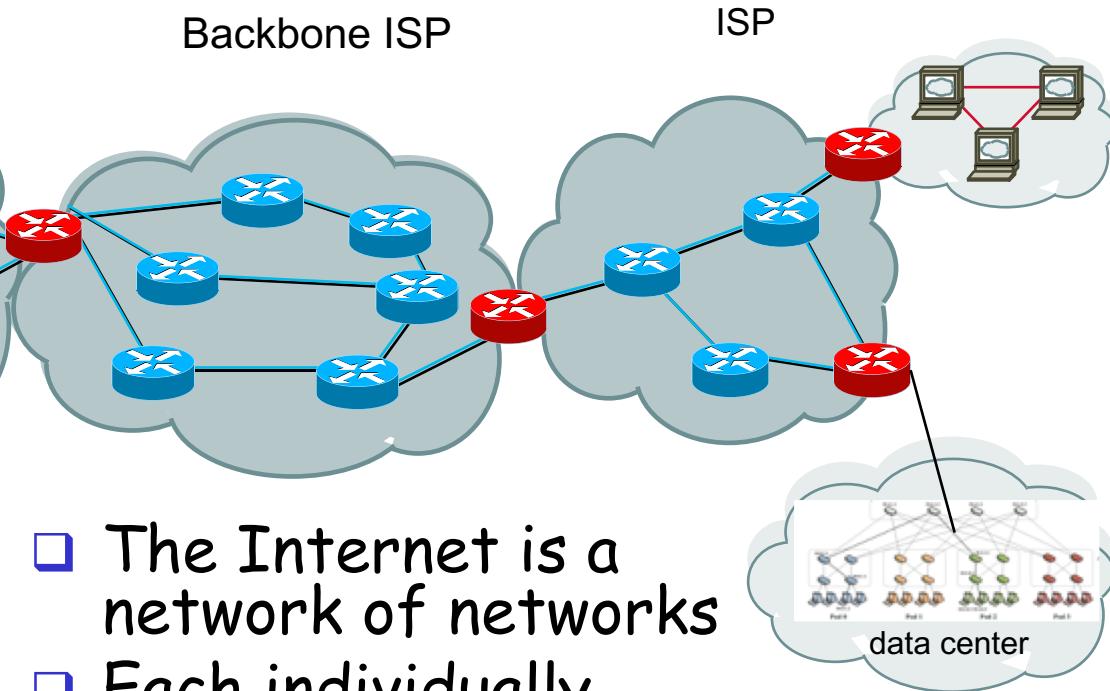
- Cable, Fiber, DSL, Wireless



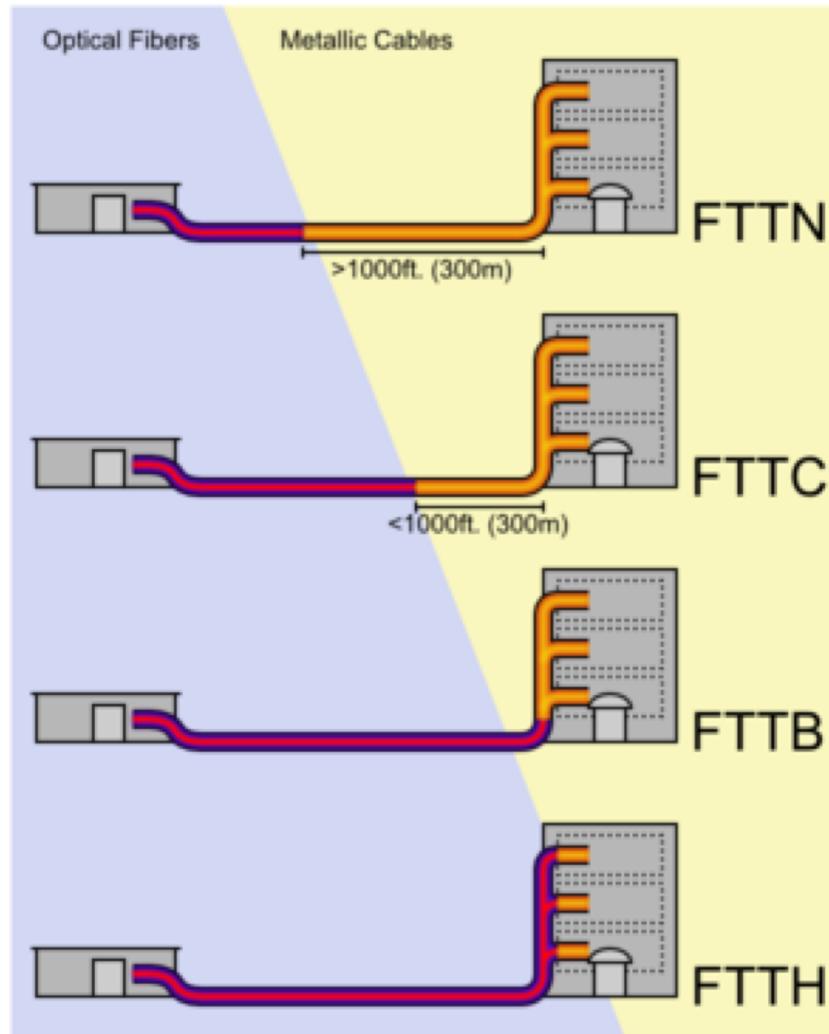
Campus access,
e.g.,

- Ethernet
- Wireless

- The Internet is a network of networks
- Each individually administrated network is called an Autonomous System (AS)



Access: Fiber to the x

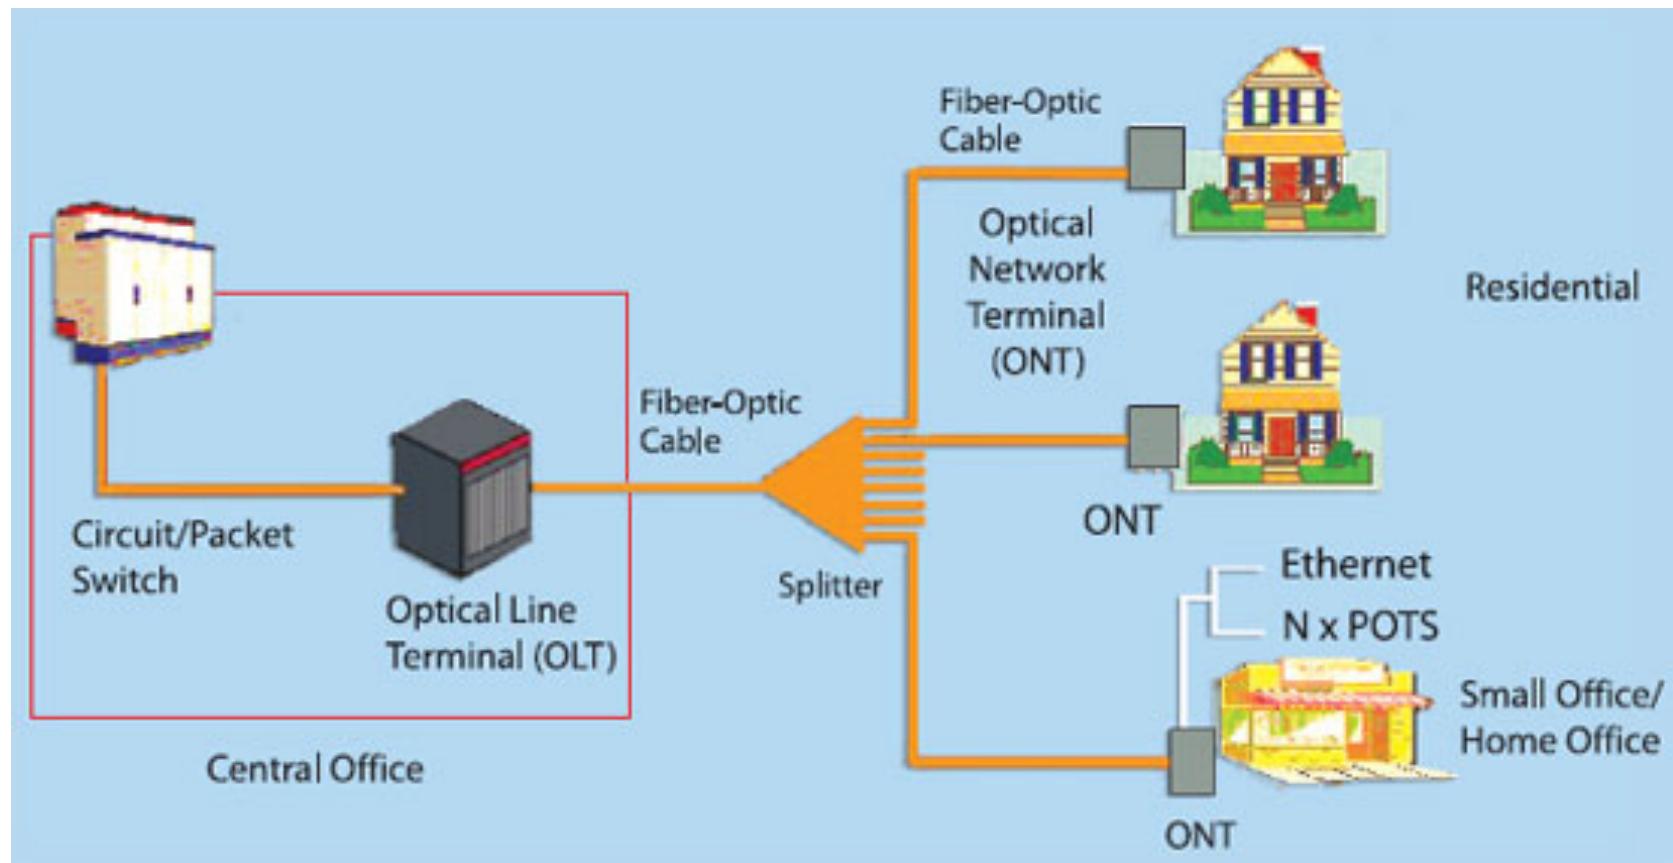


Access: Fiber to the Premises (FTTP)

- Deployed by Verizon, AT&T, Google,
- One of the largest comm. construction projects

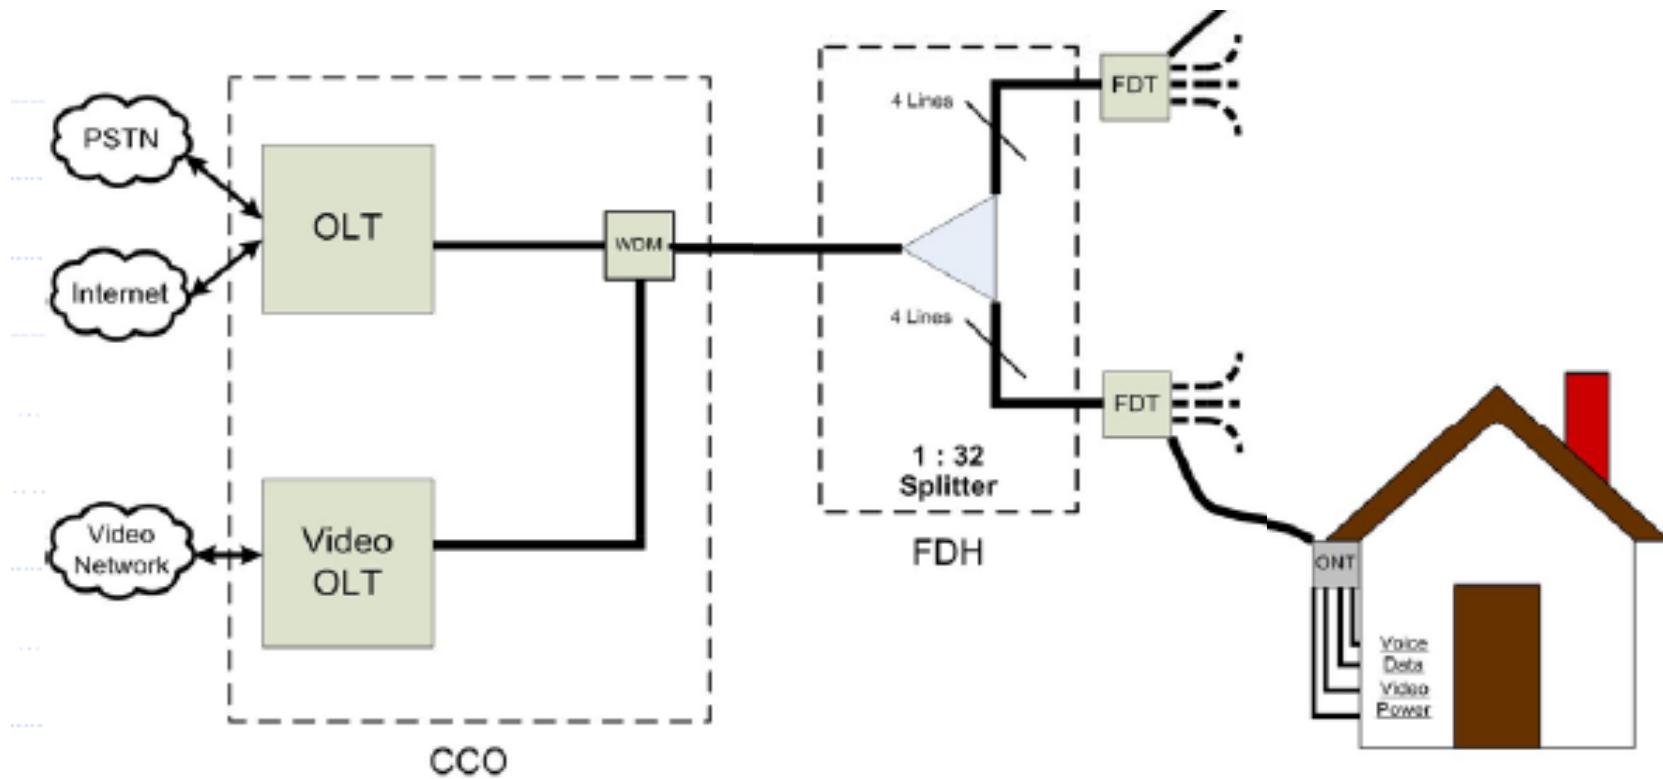


FTTP Architecture

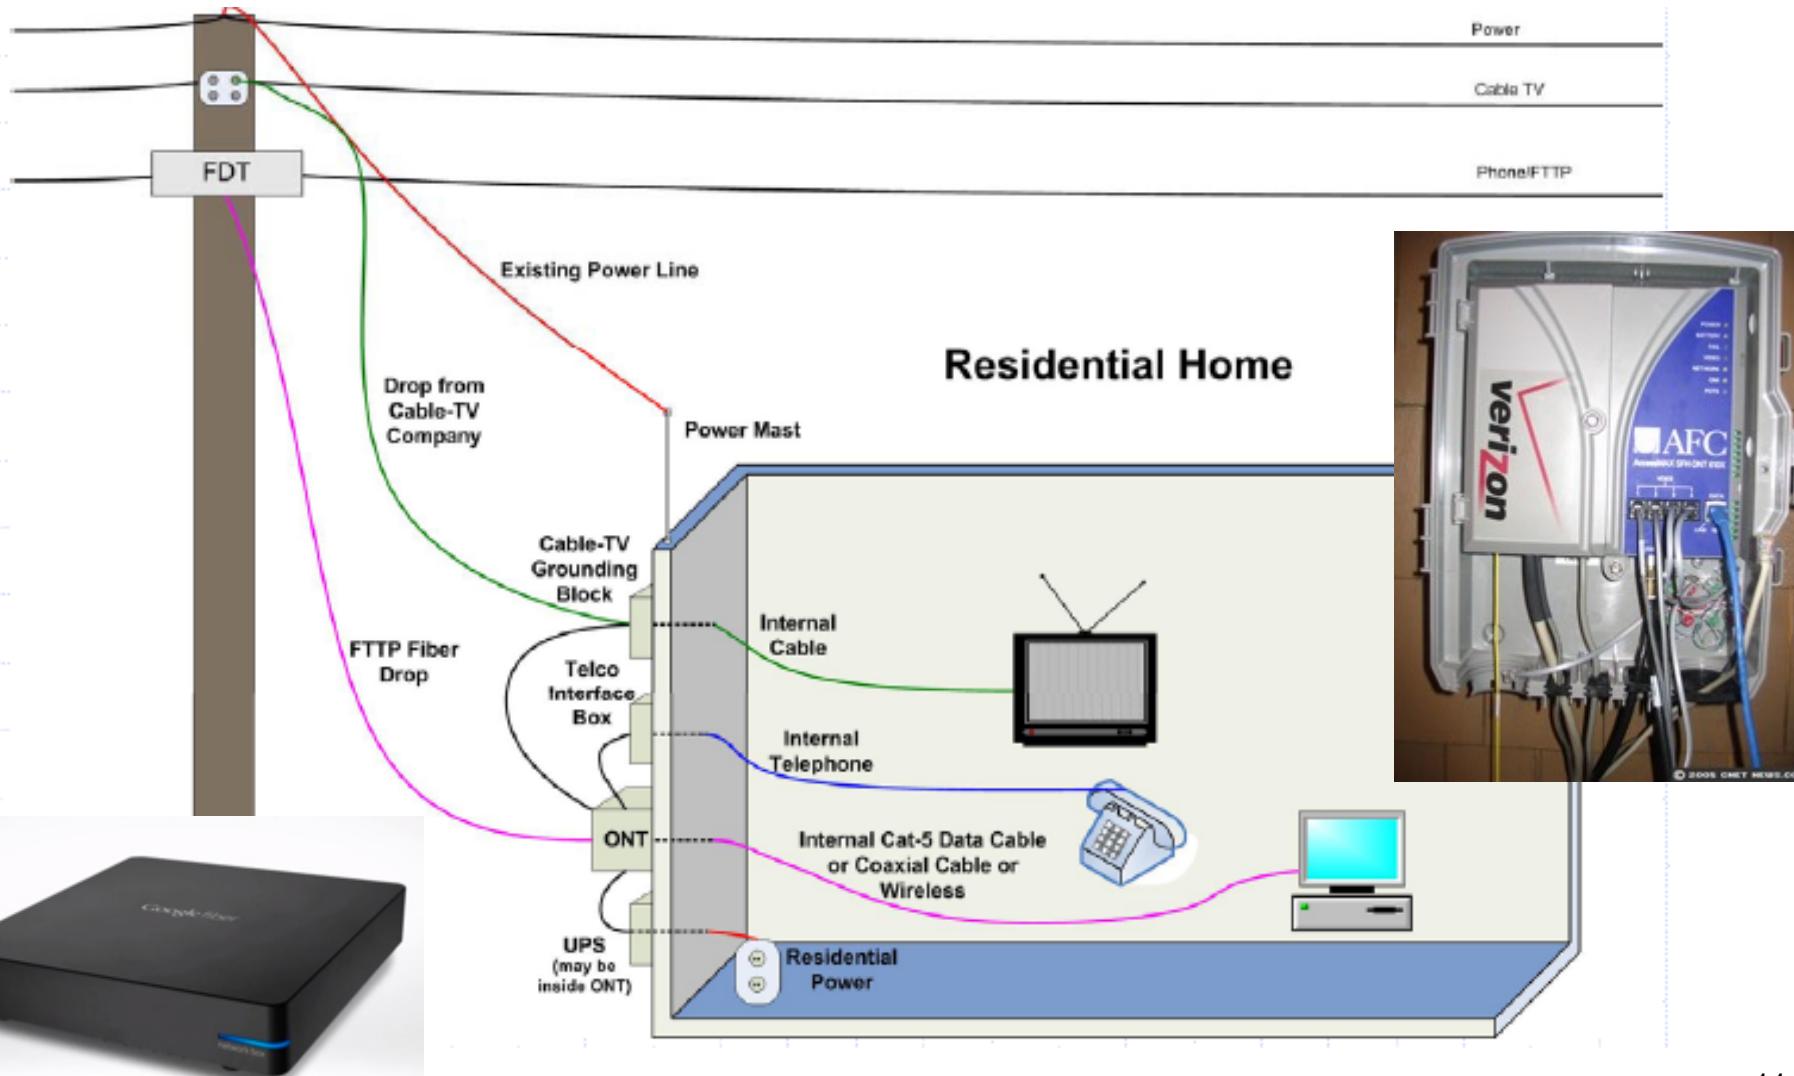


FTTP Architecture

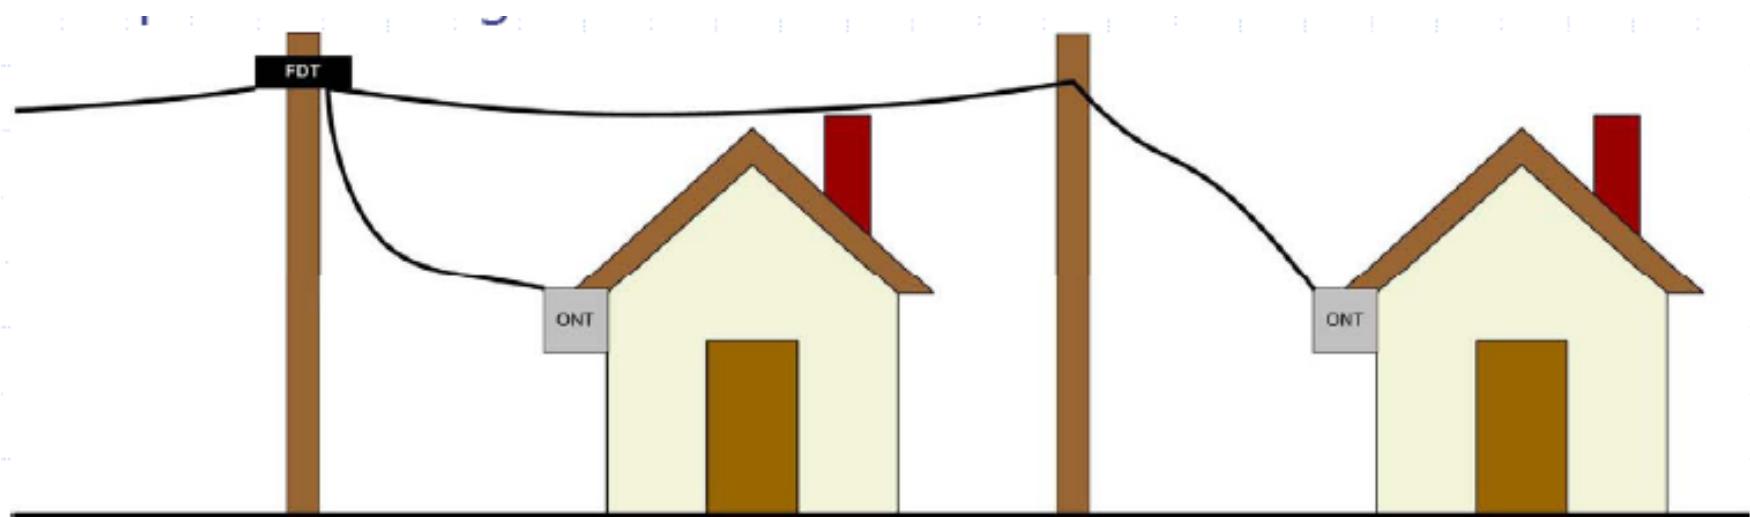
- r Optical Network Terminal (ONT) box outside dwelling or business
- r Fiber Distribution Terminal (FDT) in poles or pedestals
- r Fiber Distribution Hub (FDH) at street cabinet
- r Optical Line Terminal (OLT) at central office



FTTP Architecture: To Home



FTTP Architecture: Fiber Distribution Terminal (FDT)



FTTP Architecture: Central to Fiber Distribution Hub (FDH)



- Backbone fiber ring on primary arterial streets (brown)
- Local distribution fiber plant (red) meets backbone at cabinet



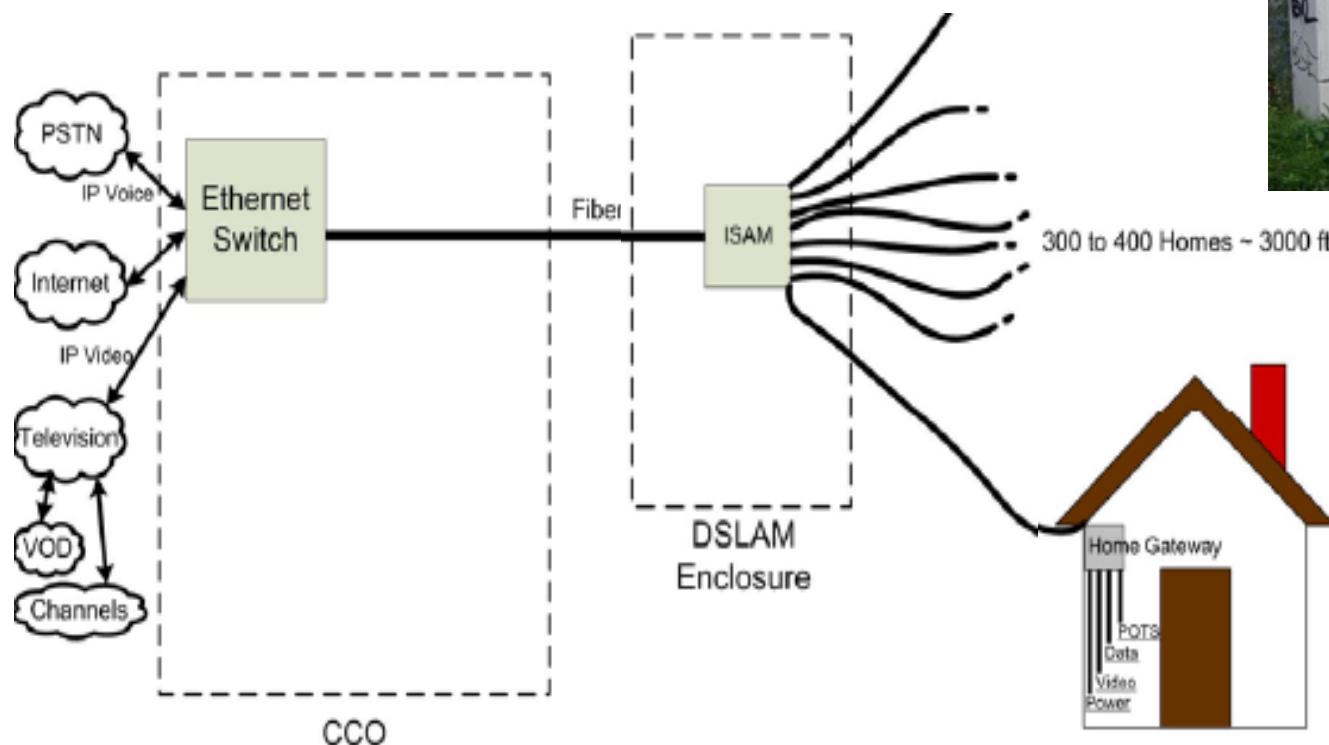
FDH

Access: DSL

- Compared with FTTP, copper from cabinet (DSLAM) to home



DSLAM



Access: Wireless

A large, metallic, ribbed hot air balloon is centered in the frame, floating over a landscape of rolling hills and mountains. The balloon is oriented vertically, with its basket at the bottom. The background shows a vast, greenish-yellow field in the foreground, leading to a range of mountains in the distance under a clear blue sky.

BALLOON-POWERED INTERNET
FOR EVERYONE

Access: Wireless

Starlink explained: Everything you should know about Elon Musk's satellite internet venture

The billionaire SpaceX CEO is launching satellites into orbit and promising to deliver high-speed broadband internet to as many users as possible.

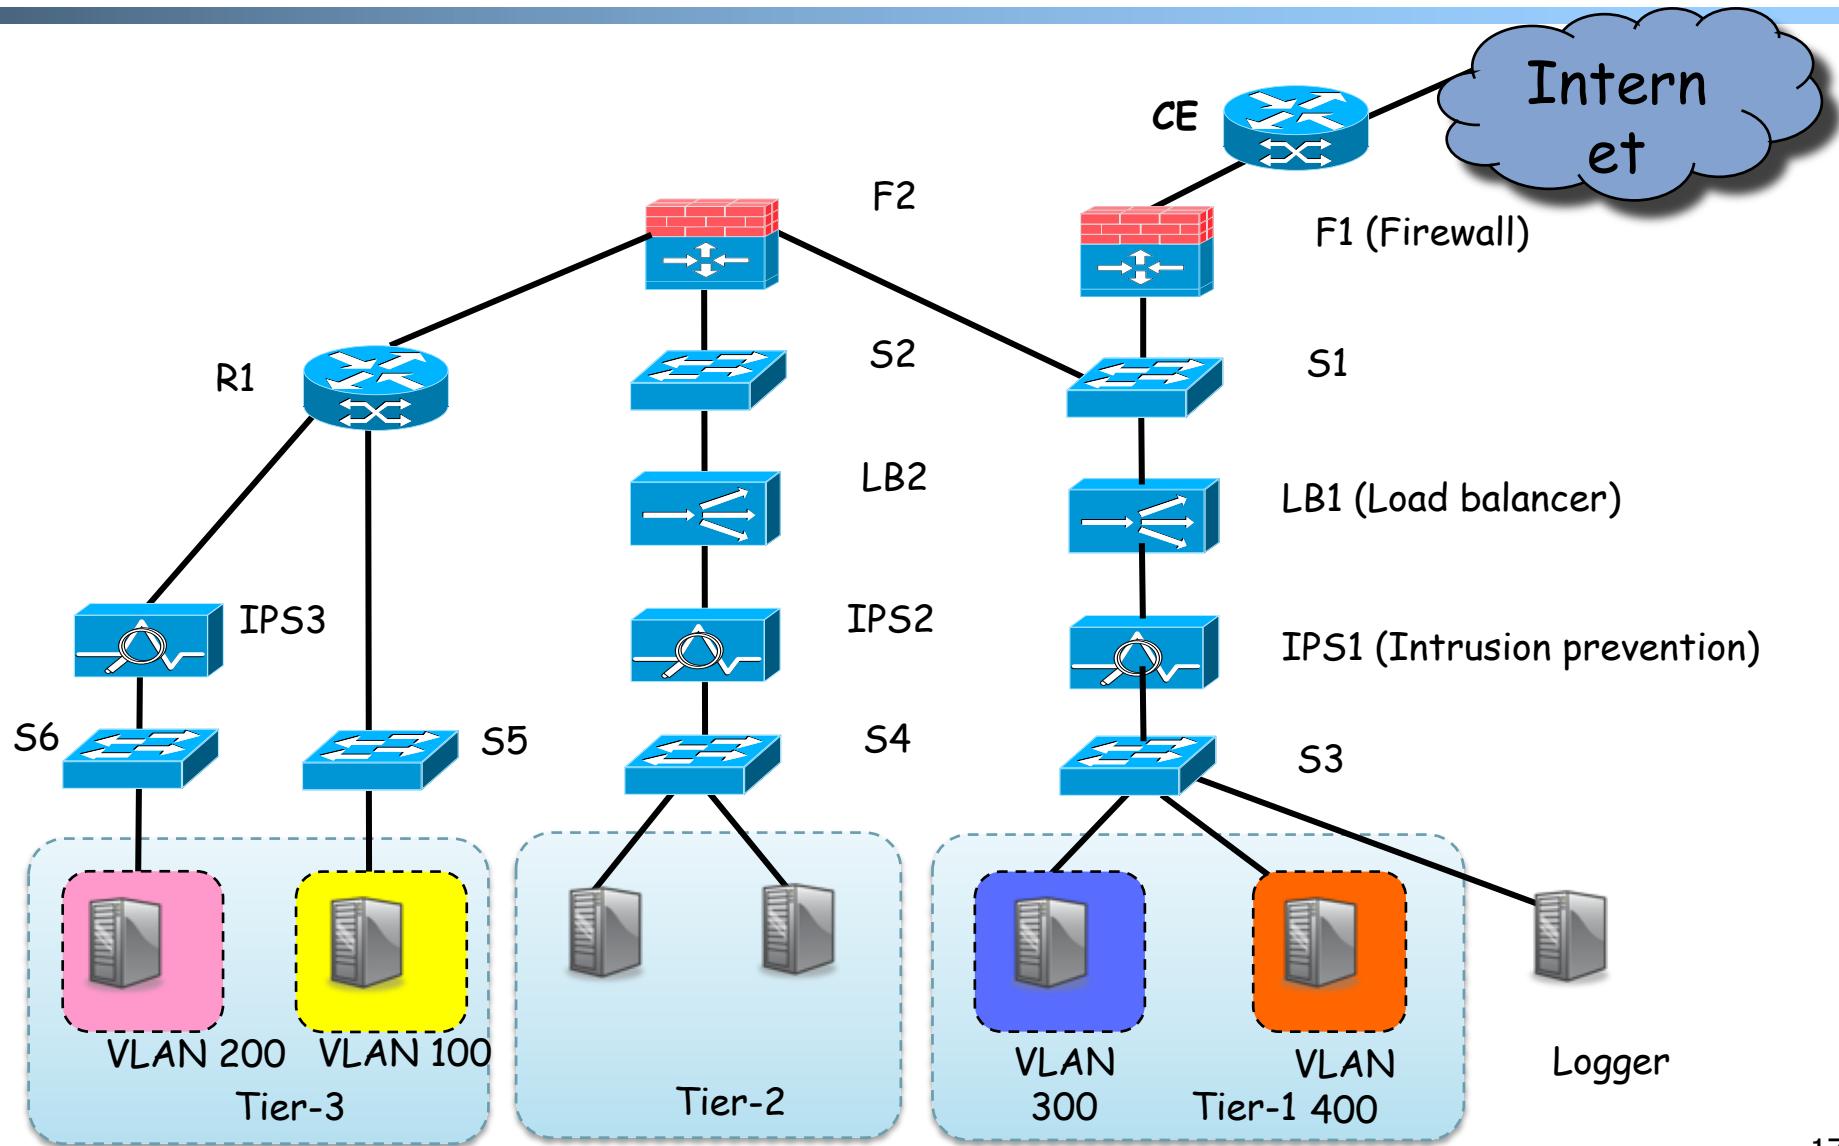


Ry Crist Aug. 24, 2021 5:15 p.m. PT

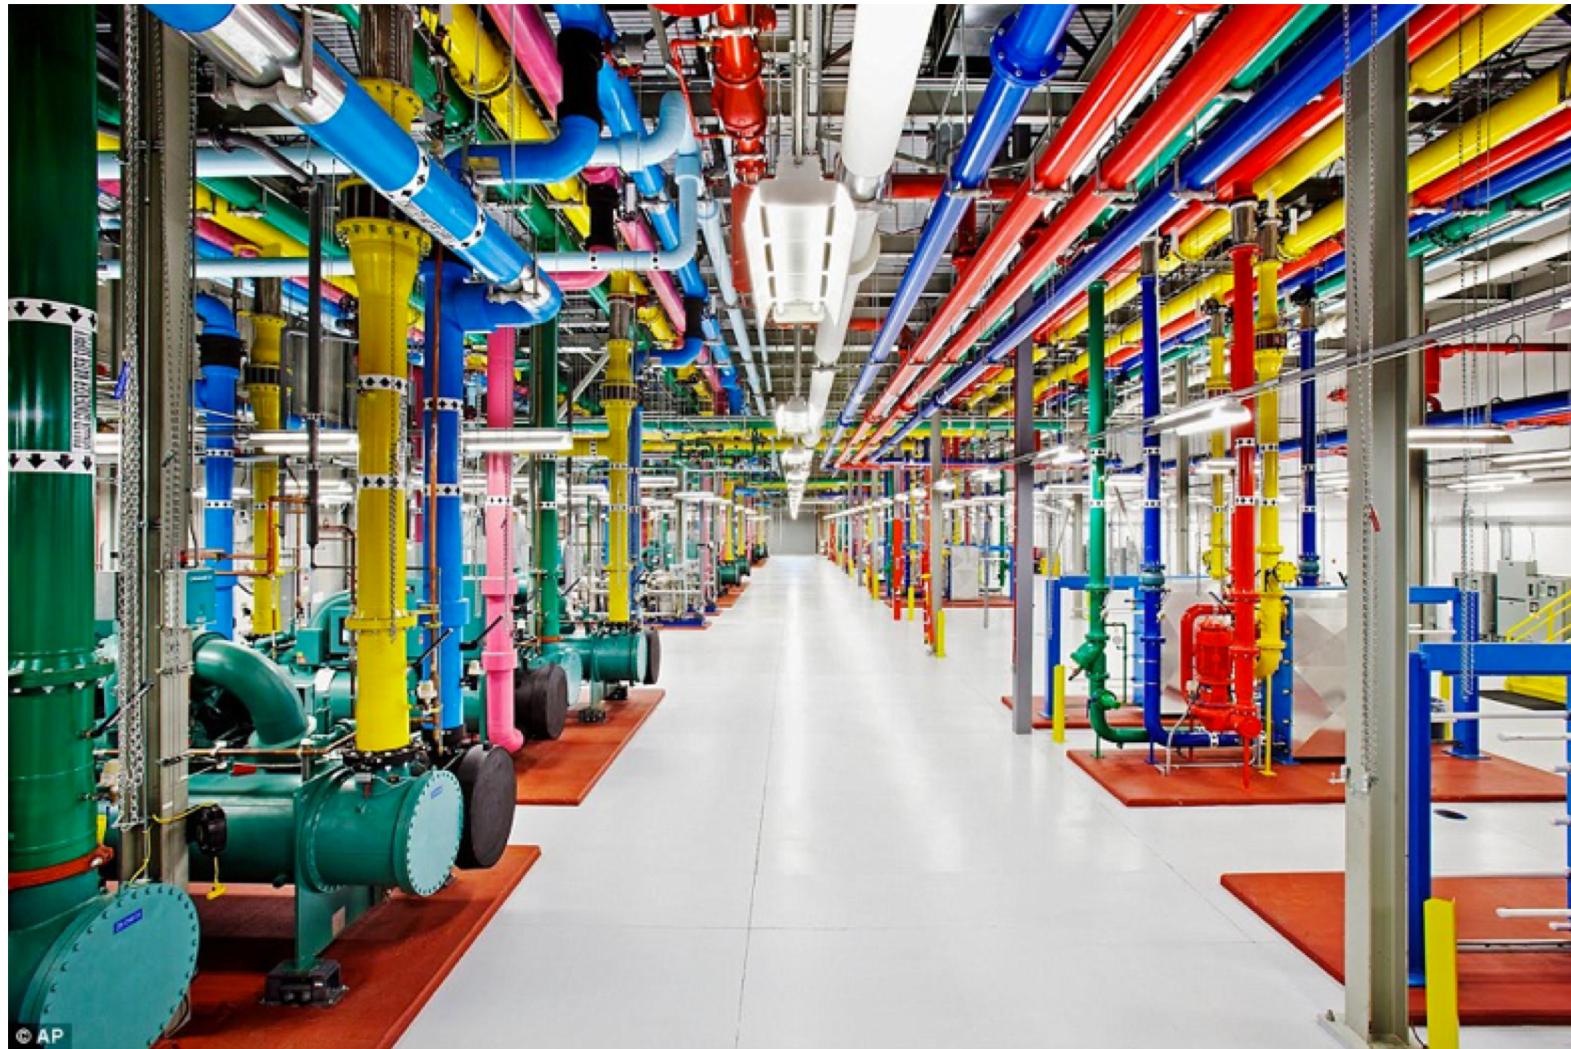
▶ LISTEN - 13:07



Campus Network

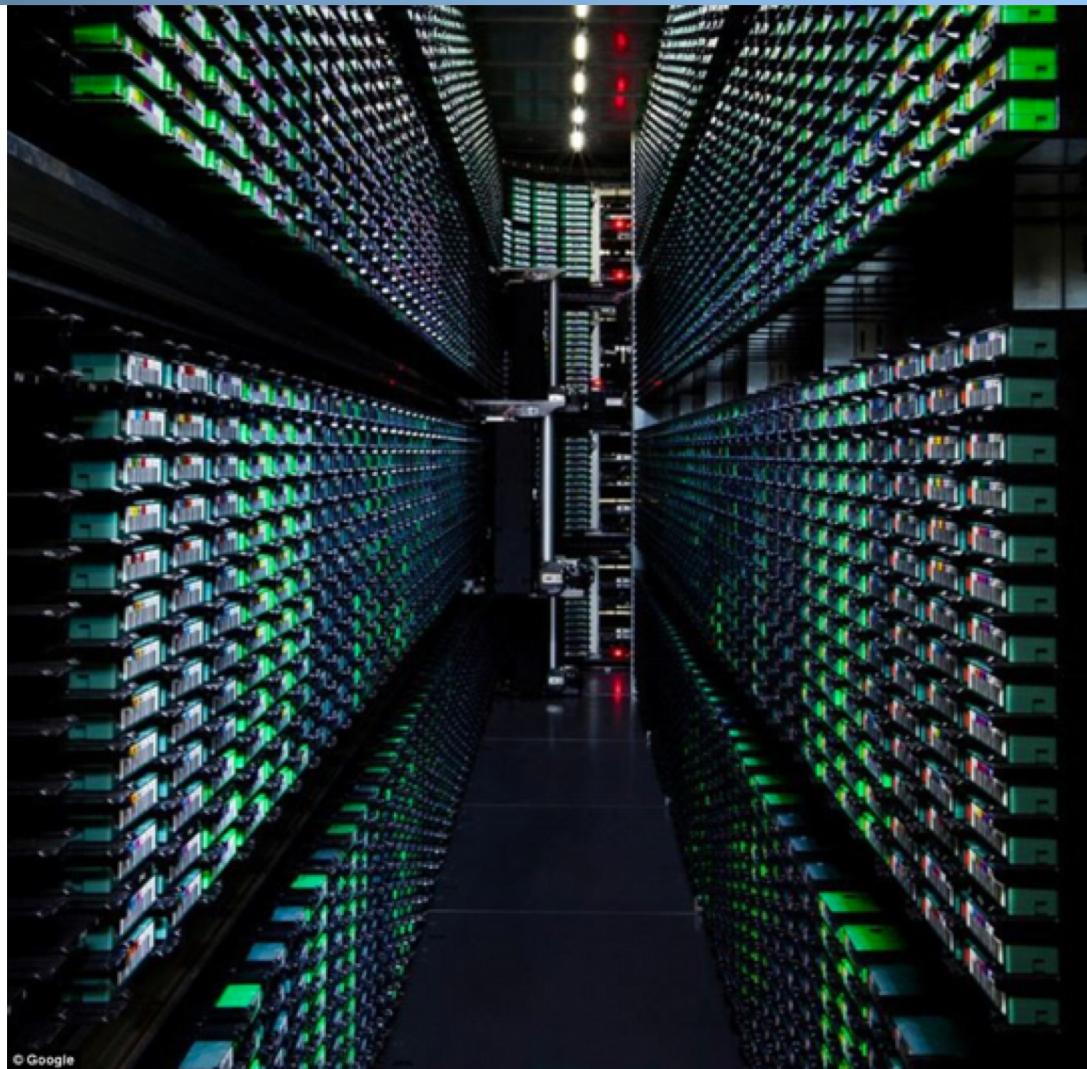


Data Center Networks



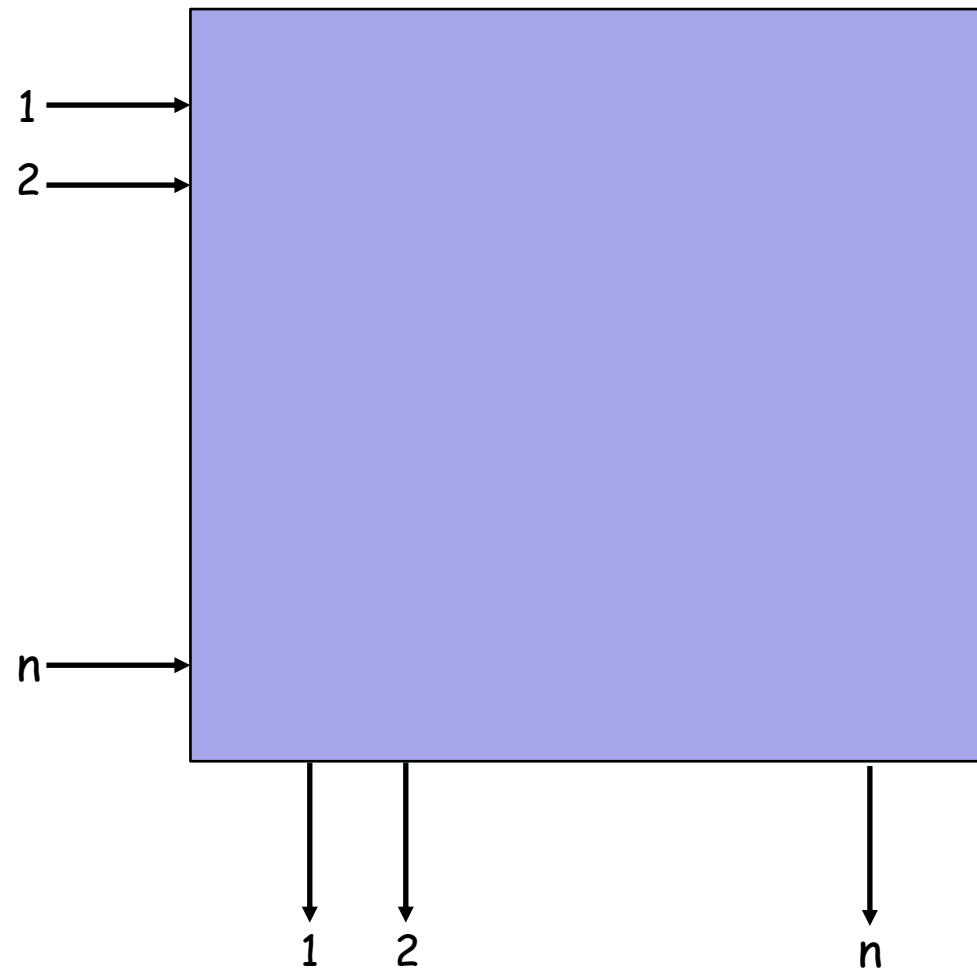
[http://www.dailymail.co.uk/sciencetech/article-3369491/Google's-plan-world-Search-engine-build-half-billion-dollar-data-center-US.html](http://www.dailymail.co.uk/sciencetech/article-3369491/Google-s-plan-world-Search-engine-build-half-billion-dollar-data-center-US.html)

Data Center Networks

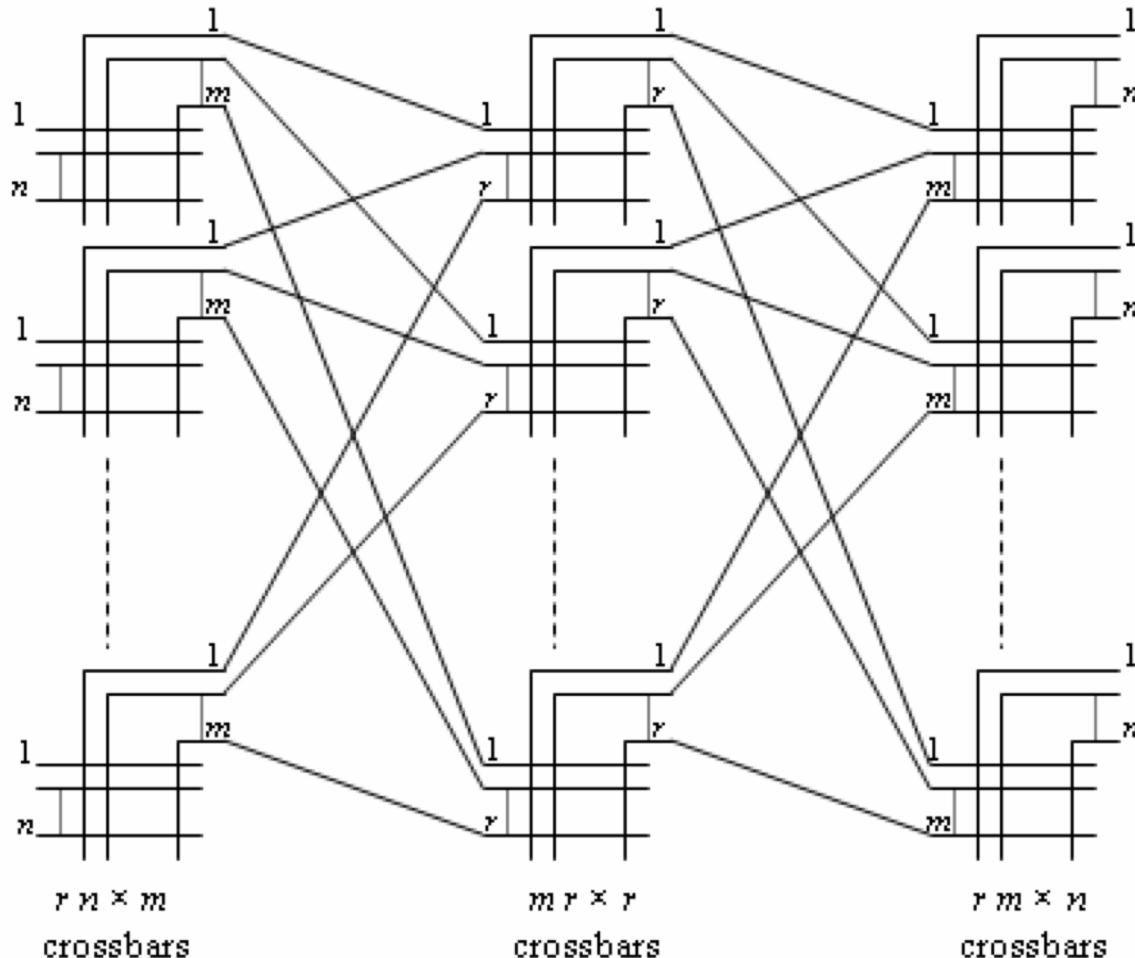


<http://www.dailymail.co.uk/sciencetech/article-3369491/Google-s-plan-world-Search-engine-build-half-billion-dollar-data-center-US.html>

Foundation of Data Center Networks



Foundation of Data Center Networks: Clos Networks

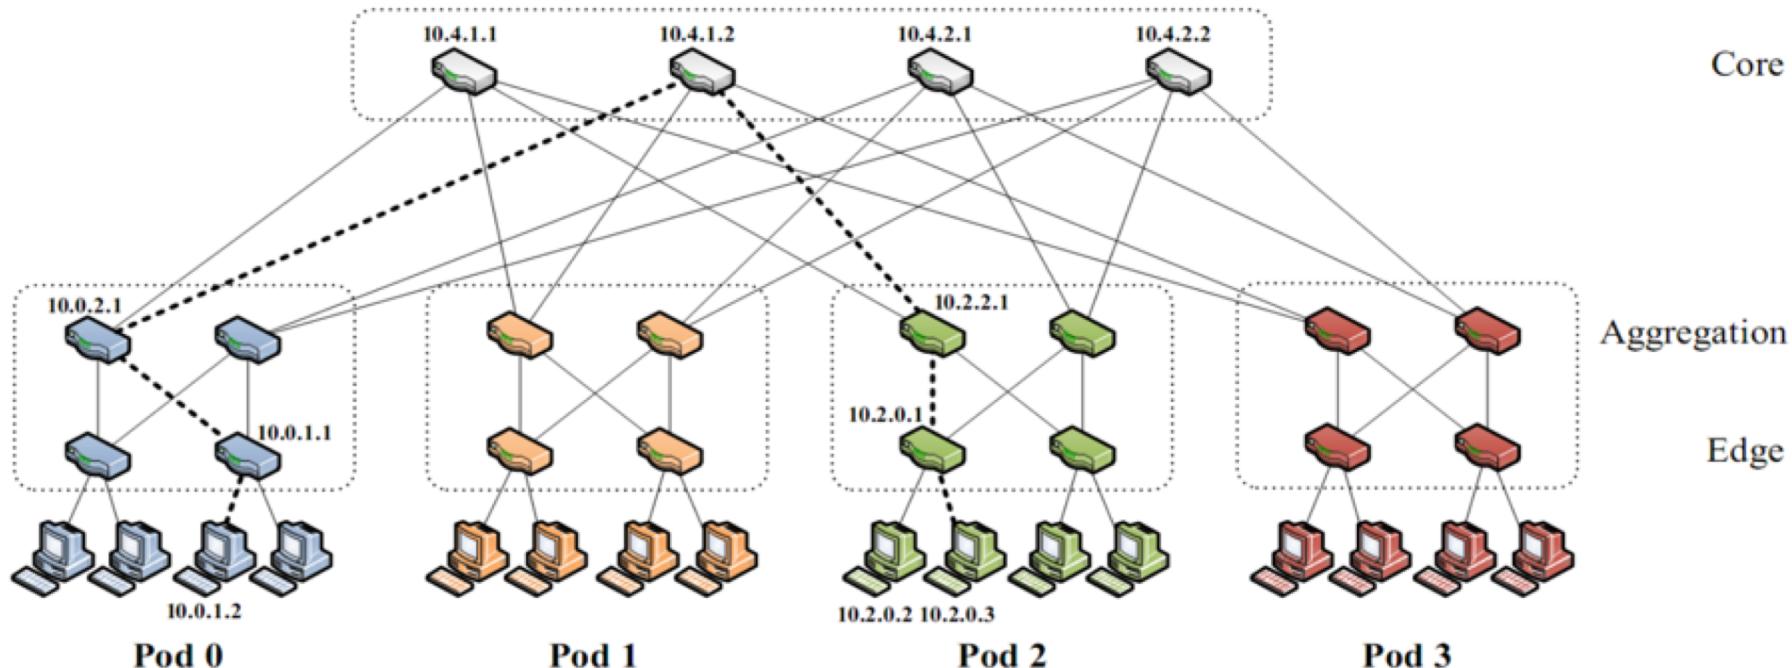


Q: How big is m so that each new call can be established w/o moving current calls?

Challenge to the class:
If you can move existing calls, it is only $m \geq n$.

Data Center Networks: Fat-tree Networks

- K-ary fat tree: three-layer topology (edge, aggregation and core)
 - k pods w/ each pod consists of $(k/2)^2$ servers & 2 layers of $k/2$ k-port switches
 - each edge switch connects to $k/2$ servers & $k/2$ aggr. switches
 - each aggr. switch connects to $k/2$ edge & $k/2$ core switches
 - $(k/2)^2$ core switches: each connects to k pods



Q: How large a network can k-ary support using k-port switches?

Data Center Networks

- r For example, Google Jupiter at 1 Pbits/sec bisection bw: 100,000 servers at 10G each

Datacenter Generation	First Deployed	Merchant Silicon	ToR Config	Aggregation Block Config	Spine Block Config	Fabric Speed	Host Speed	Bisection BW
Four-Post CRs	2004	vendor	48x1G	-	-	10G	1G	2T
Firehose 1.0	2005	8x10G 4x10G (ToR)	2x10G up 24x1G down	2x32x10G (B)	32x10G (NB)	10G	1G	10T
Firehose 1.1	2006	8x10G	4x10G up 48x1G down	64x10G (B)	32x10G (NB)	10G	1G	10T
Watchtower	2008	16x10G	4x10G up 48x1G down	4x128x10G (NB)	128x10G (NB)	10G	nx1G	82T
Saturn	2009	24x10G	24x10G	4x288x10G (NB)	288x10G (NB)	10G	nx10G	207T
Jupiter	2012	16x40G	16x40G	8x128x40G (B)	128x40G (NB)	10/40G	nx10G/ nx40G	1.3P

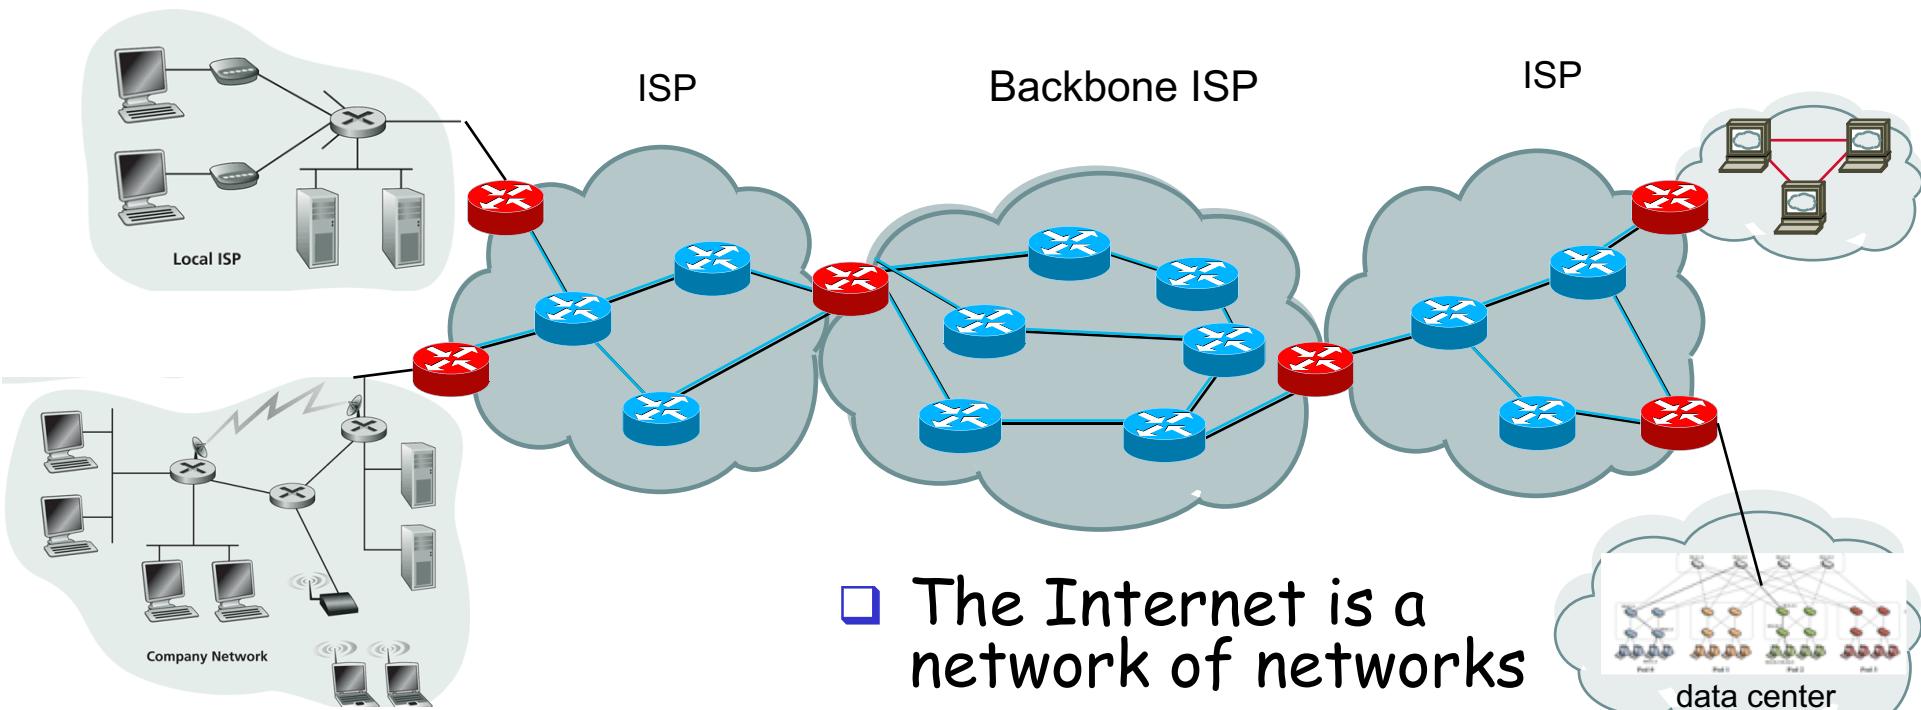
<http://googlecloudplatform.blogspot.com/2015/06/A-Look-Inside-Googles-Data-Center-Networks.html>

<http://conferences.sigcomm.org/sigcomm/2015/pdf/papers/p183.pdf>

Recall: Internet Physical Infrastructure

Residential access, e.g.,

- Cable, Fiber, DSL, Wireless



Campus access, e.g.,

- Ethernet, Wireless

- The Internet is a network of networks
- Each individually administrated network is called an Autonomous System (AS)

Yale Internet Connection

traceroute www.tsinghua.edu.cn

```
1 college.net.yale.internal (172.28.201.65) 1.440 ms 1.227 ms 1.453 ms
2 10.1.1.13 (10.1.1.13) 1.359 ms 1.153 ms 1.173 ms
3 level3-10g-asr.net.yale.internal (10.1.4.40) 2.786 ms 6.110 ms 2.547 ms
4 cen-10g-yale.net.yale.internal (10.1.3.102) 2.646 ms 3.242 ms 2.576 ms
5 * * *
6 enrt064hhh-9k-te0-3-0-5.net.cen.ct.gov (67.218.83.254) 5.169 ms 3.797 ms 6.891 ms
7 198.71.46.215 (198.71.46.215) 3.615 ms 3.742 ms 3.931 ms
8 et-10-0-0.1180.rtsw.newy32aoa.net.internet2.edu (198.71.46.214) 6.661 ms 6.532 ms 6.310 ms
9 et-4-0-0.4079.sdn-sw.phil.net.internet2.edu (162.252.70.103) 8.658 ms 8.714 ms 8.666 ms
10 et-1-1-0.4079.rtsw.wash.net.internet2.edu (162.252.70.119) 11.787 ms 30.111 ms 11.900 ms
11 et-8-1-0.4079.sdn-sw.ashb.net.internet2.edu (162.252.70.62) 12.428 ms 16.654 ms 15.862 ms
12 et-7-1-0.4079.rtsw.chic.net.internet2.edu (162.252.70.61) 28.898 ms 28.999 ms 28.908 ms
13 et-3-1-0.4070.rtsw.kans.net.internet2.edu (198.71.47.207) 40.084 ms 39.958 ms 39.695 ms
14 et-8-0-0.4079.sdn-sw.denv.net.internet2.edu (162.252.70.10) 50.195 ms 50.562 ms 50.258 ms
15 et-4-1-0.4079.rtsw.salt.net.internet2.edu (162.252.70.9) 59.707 ms 60.261 ms 59.762 ms
16 et-7-0-0.4079.sdn-sw.lasv.net.internet2.edu (162.252.70.30) 67.555 ms 67.539 ms 67.312 ms
17 et-4-1-0.4079.rtsw.losa.net.internet2.edu (162.252.70.29) 72.419 ms 72.428 ms 72.376 ms
...
```

Internet2



INTERNET2 NETWORK INFRASTRUCTURE TOPOLOGY

OCTOBER 2014

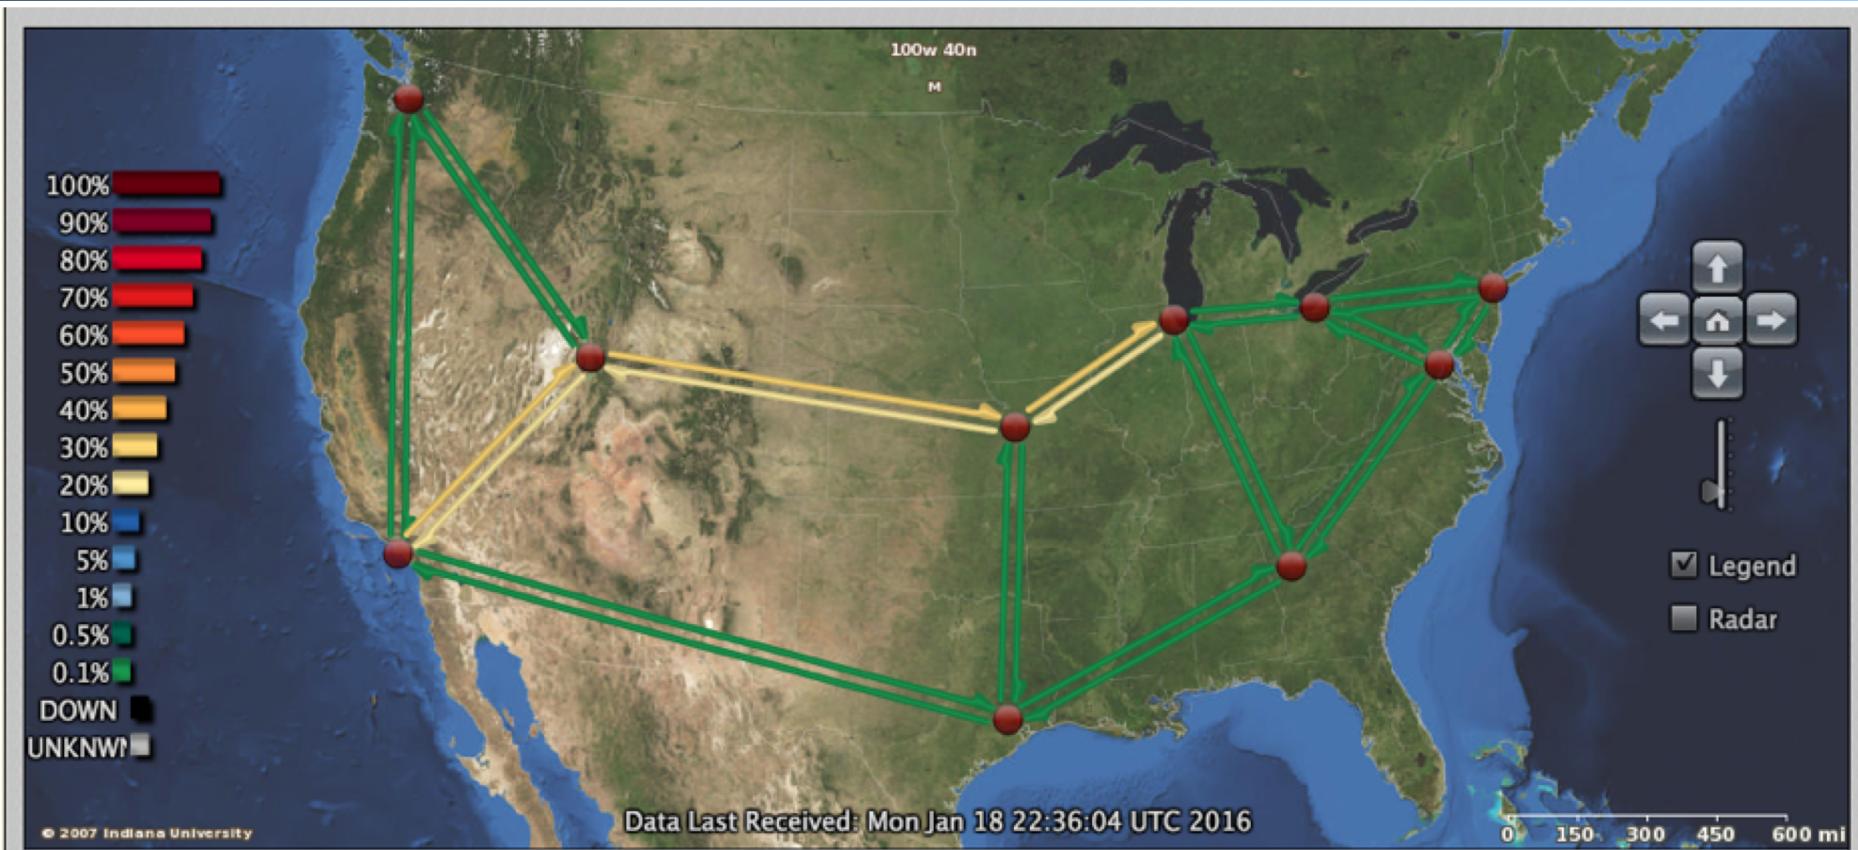


INTERNET2 NETWORK BY THE NUMBERS

12	JUNIOR MFGD ROLERS SUPPORTING LAYER 3 SERVICE
34	BANDONE AND JUNIOR SWITCHES SUPPORTING LAYER 2 SERVICE
66	CUSTOM COLLOCATION FACILITIES
86	AMPLIFICATION BACKS
527	MILES OF NEWLY ACQUIRED DARK FIBER
83	TYPES OF OPTICAL CAPACITY
193	GIGS OF HYBRID LAYER 2 AND LAYER 3 CAPACITY
522	CENTRALIZED SDN/RNF NETWORK ELEMENTS
246	MILES PARTNERED CAPACITY WITH ZTE COMMUNICATIONS IN SUPPORT OF THE NORTHERN U.S. REGION



Internet2



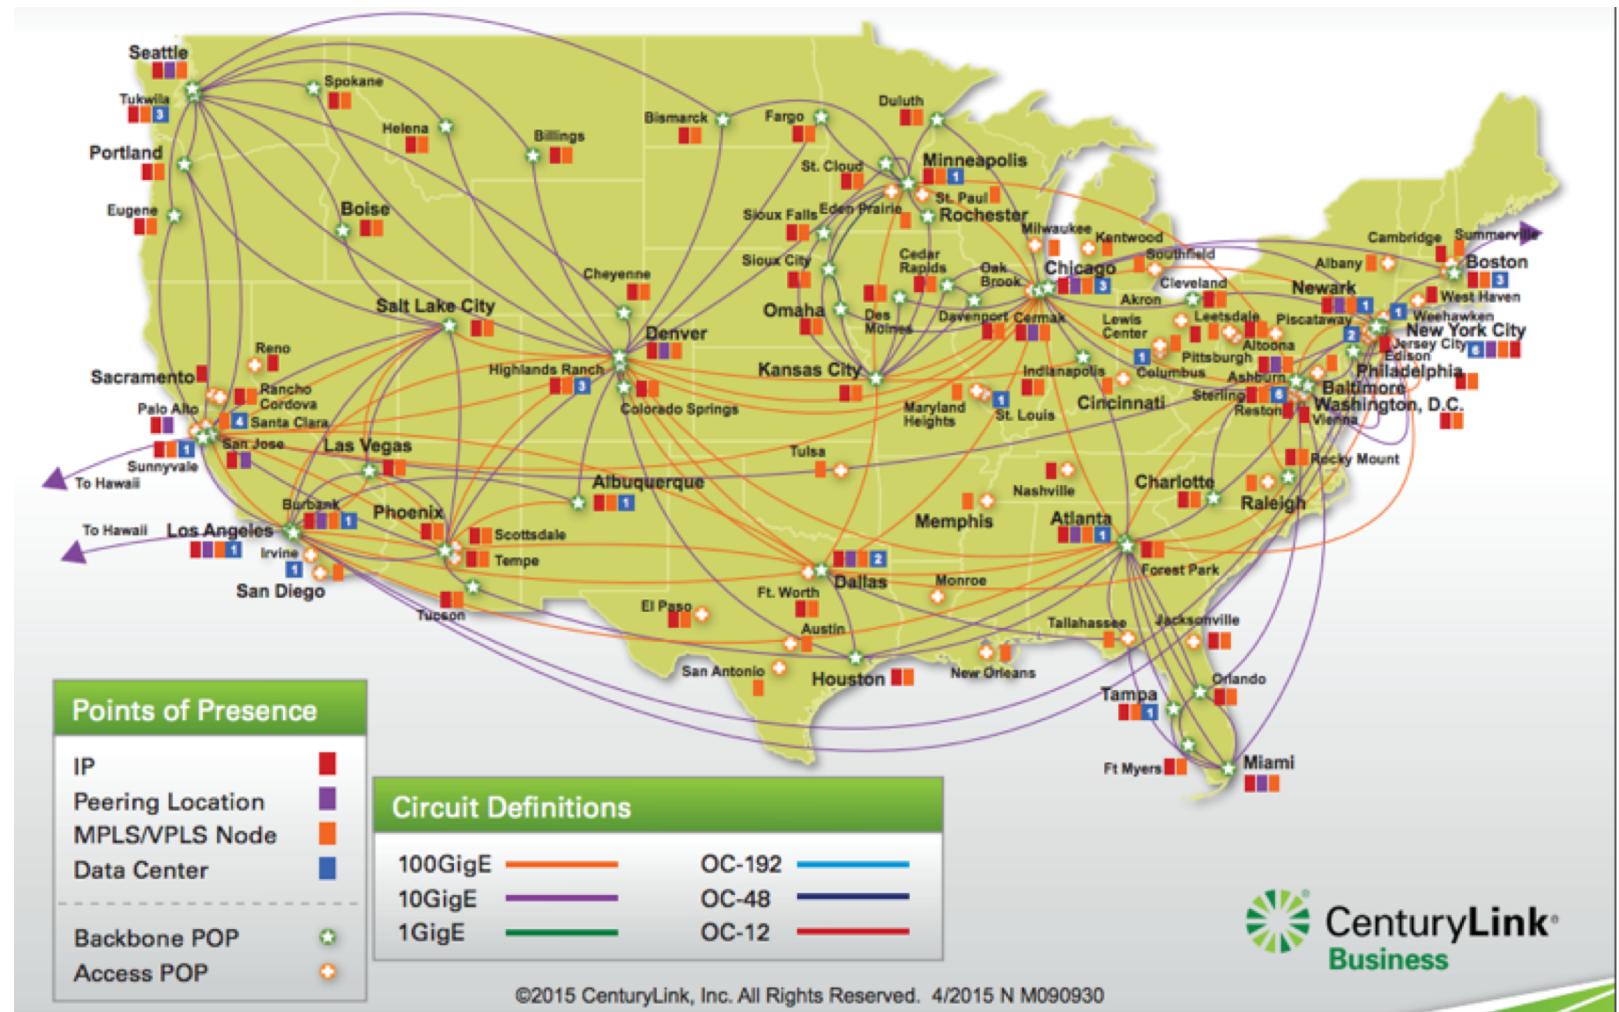
http://atlas.grnoc.iu.edu/atlas.cgi?map_name=Internet2%20IP%20Layer

XMU Internet Connection

Try traceroute from XMU to

- www.microsoft.com
- www.baidu.com
- www.sina.com.cn
- www.taobao.com

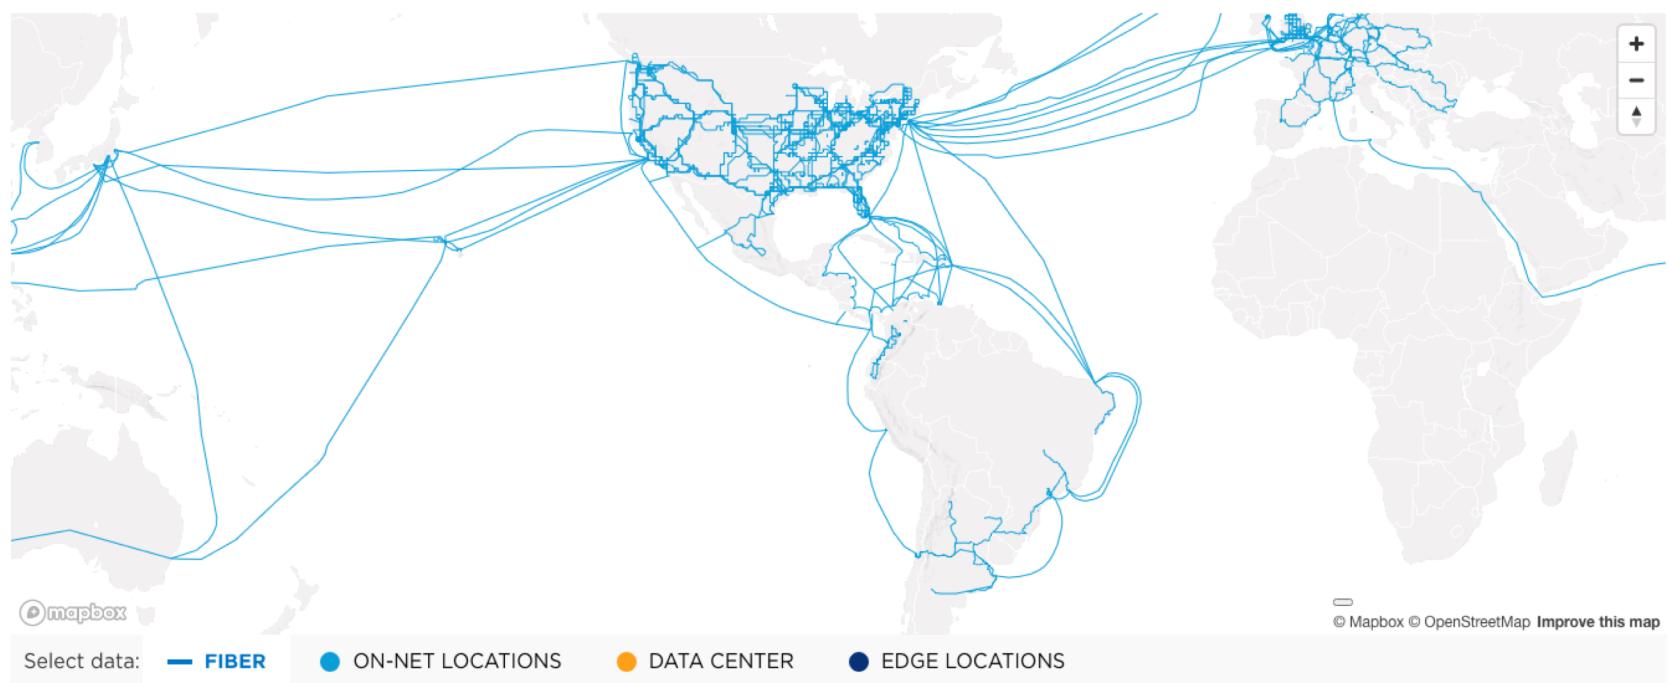
Qwest (CenturyLink) Network Maps



<http://www.centurylink.com/business/asset/network-map/ip-mpls-network-nm090930.pdf>

<http://www.centurylink.com/business/resource-center/network-maps/>

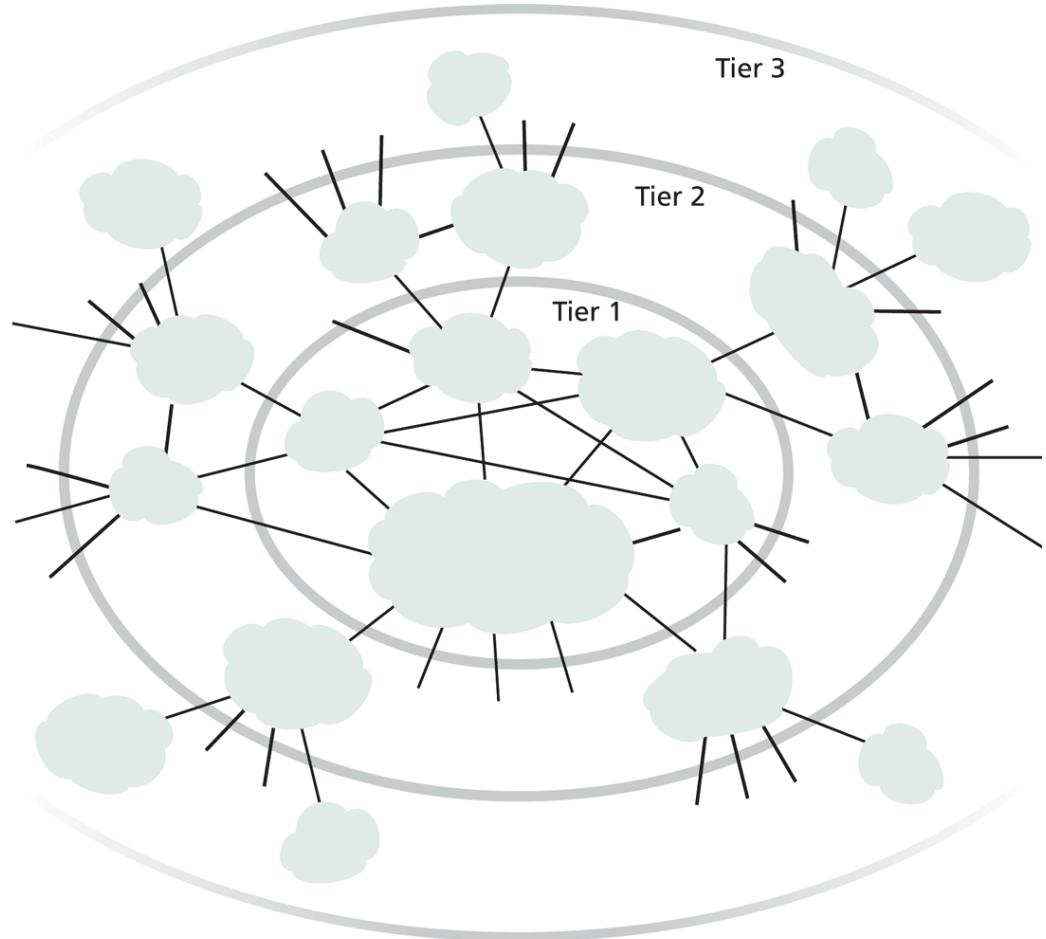
Level3 (now part of LUMEN) Network Map



<https://www.lumen.com/en-us/resources/network-maps.html>

Internet ISP Connectivity

- Roughly hierarchical
 - Divided into tiers
 - Tier-1 ISPs are also called backbone providers, e.g., AT&T, Verizon, Sprint, Level 3, Qwest
- An ISP runs (private) Points of Presence (PoP) where its customers and other ISPs connect to it
- ISPs also connect at (public) Internet Exchange Point (IXP)
 - public peering



Outline

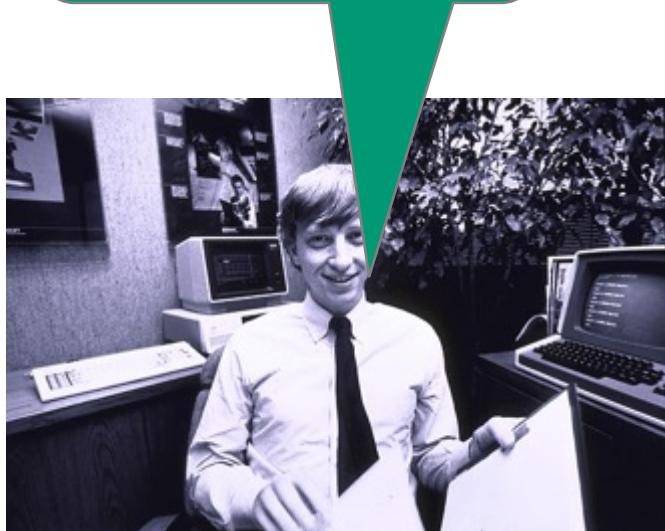
- Admin. and recaps
- *A brief introduction to the Internet*
 - past
 - *present*
 - *topology*
 - *traffic*

Internet (Consumer) Traffic

Consumer Internet Traffic, 2012–2017								CAGR 2012–2017
	2012	2013	2014	2015	2016	2017		
By Network (PB per Month)								
Fixed	25,529	32,097	39,206	47,035	56,243	66,842		21%
Mobile	684	1,239	2,223	3,774	6,026	9,131		68%
By Subsegment (PB per Month)								
Internet video	14,818	19,855	25,800	32,962	41,916	52,752		29%
Web, email, and data	5,173	6,336	7,781	9,542	11,828	14,494		23%
File sharing	6,201	7,119	7,816	8,266	8,478	8,667		7%
Online gaming	22	26	32	39	48	59		22%
By Geography (PB per Month)								
Asia Pacific	9,033	11,754	14,887	18,707	23,458	29,440		27%
North America	6,834	8,924	11,312	14,188	17,740	21,764		26%
Western Europe	5,086	5,880	6,804	7,810	9,197	10,953		17%
Central and Eastern Europe	2,194	2,757	3,433	4,182	5,015	5,897		22%
Latin America	2,656	3,382	4,049	4,588	5,045	5,487		16%
Middle East and Africa	410	640	944	1,334	1,816	2,432		43%
Total (PB per Month)								
Consumer Internet traffic	26,213	33,337	41,429	50,809	62,269	75,973		24%

Internet Traffic in Perspective

640K ought to be enough
for anybody.



1 Petabyte
1,000 Terabytes or
250,000 DVDs

1 Exabyte
1,000 Petabytes or
250 million DVDs

1 Zettabyte
1,000 Exabytes or
250 billion DVDs

1 Yottabyte
1,000 Zettabytes or
250 trillion DVDs

480 Terabytes

A digital library of all of the world's catalogued books in all languages

100 Petabytes

The amount of data produced in a single minute by the new particle collider at CERN

5 Exabytes

A text transcript of all words ever spoken†

100 Exabytes

A video recording of all the meetings that took place last year across the world

400 Exabytes

The amount of data that crossed the Internet in 2012 alone

1 Zettabyte

The amount of data that has traversed the Internet since its creation

300 Zettabytes

The amount of visual information conveyed from the eyes to the brain of the entire human race in a single year‡

20 Yottabytes

A holographic snapshot of the earth's surface

† Roy Williams, "Data Powers of Ten," 2000

‡ Based on a 2006 estimate by the University of Pennsylvania School of Medicine that the retina transmits information to the brain at 10 Mbps.

Outline

- Admin. and recaps
- A brief introduction to the Internet: past and present
- *Challenges of Internet networks and apps*

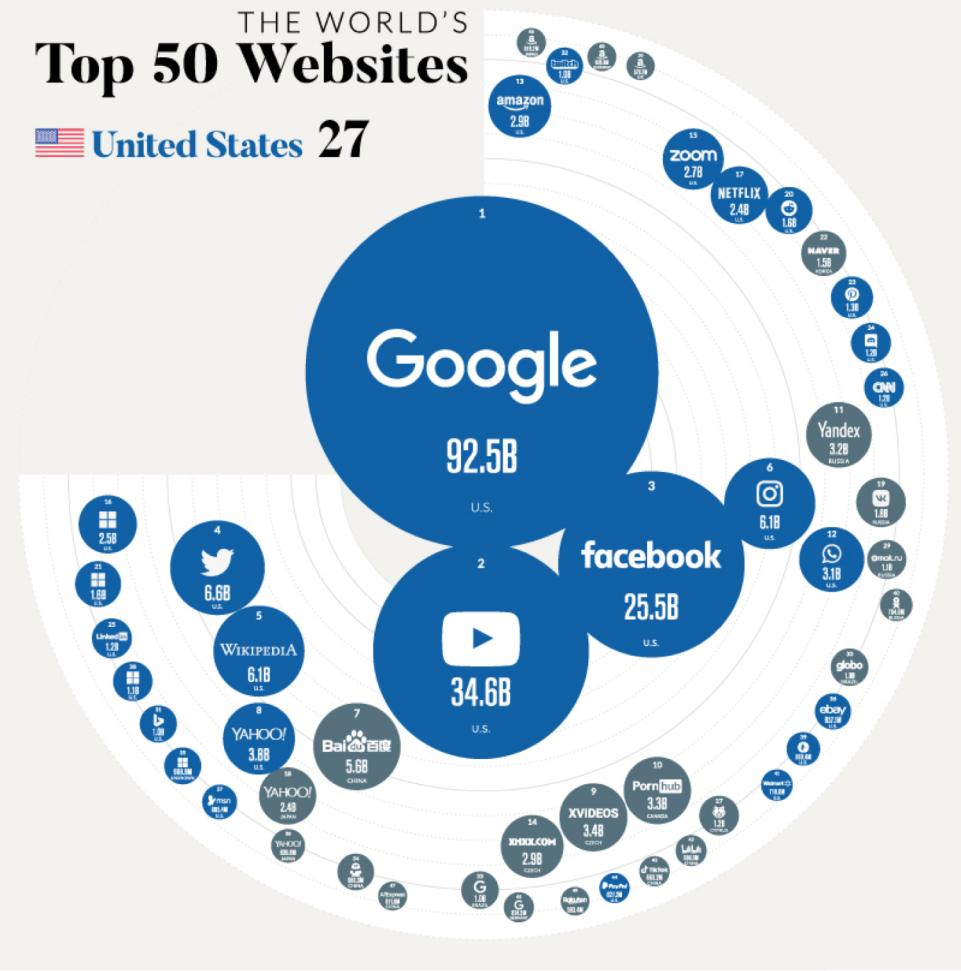
Scale



“Developers who have worked at the small scale might be asking themselves why we need to bother when we could just use some kind of out-of-the-box solution. For small-scale applications, this can be a great idea. We save time and money up front and get a working and serviceable application. The problem comes at larger scales—there are no off-the-shelf kits that will allow you to build something like Amazon... There’s a good reason why the largest applications on the Internet are all bespoke creations: no other approach can create massively scalable applications within a reasonable budget.”

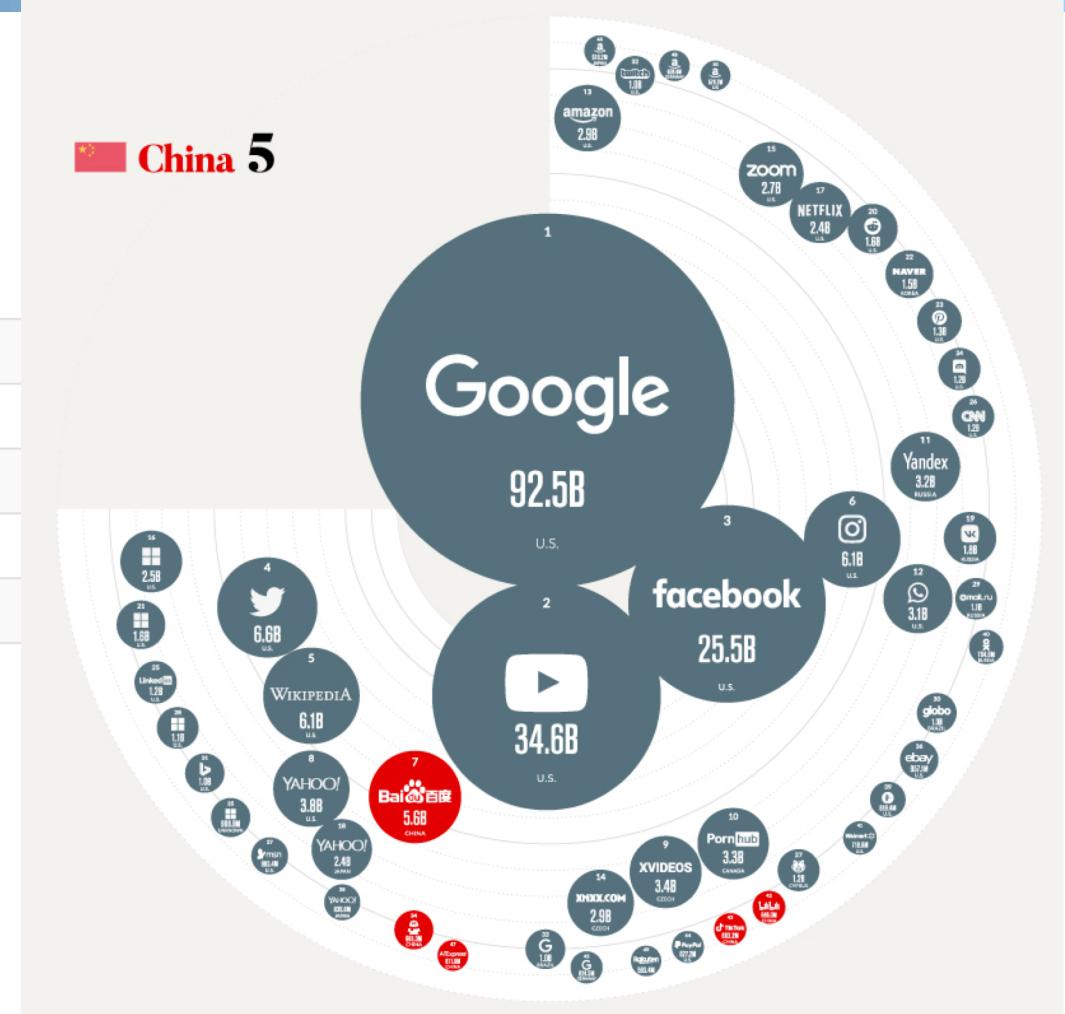
Largest Internet Sites in the World

1	Google.com	92.5B
2	Youtube.com	34.6B
3	Facebook.com	25.5B
4	Twitter.com	6.6B
5	Wikipedia.org	6.1B
6	Instagram.com	6.1B
8	Yahoo.com	3.8B
12	Whatsapp.com	3.1B
13	Amazon.com	2.9B
15	Zoom.us	2.7B
16	Live.com	2.5B
17	Netflix.com	2.4B
20	Reddit.com	1.6B
21	Office.com	1.6B
23	Pinterest.com	1.3B
24	Discord.com	1.2B
25	Linkedin.com	1.2B
26	Cnn.com	1.2B



Largest Internet Sites in the World

7	Baidu.com	5.6B
34	QQ.com	981.3M
42	Bilibili.com	686.0M
43	Tiktok.com	663.2M
47	Aliexpress.com	611.0M



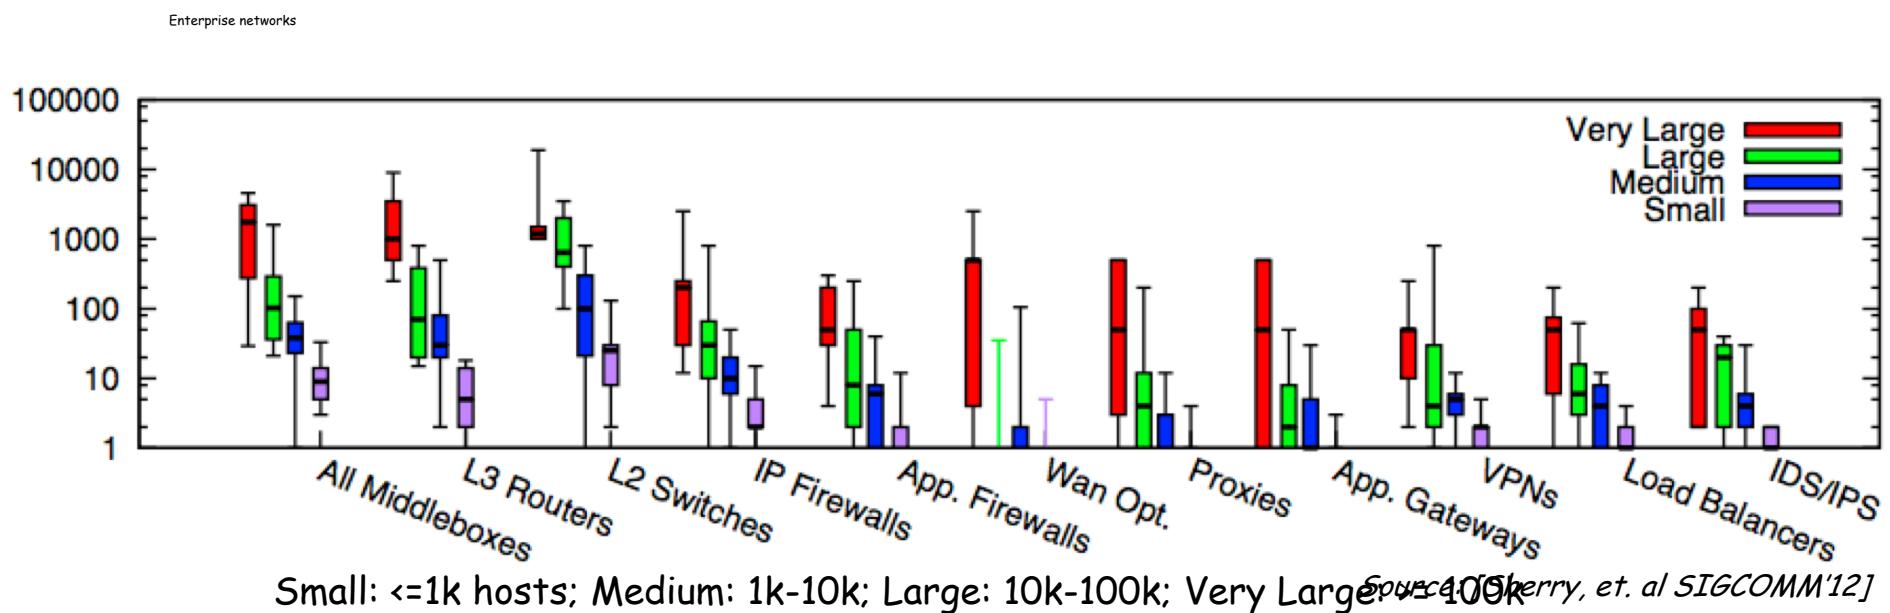
General Complexity



- Complexity in highly organized systems arises primarily from design strategies intended to create robustness to uncertainty in their environments and component parts.
 - Scalability is robustness to changes to the size and complexity of a system as a whole.
 - Evolvability is robustness of lineages to large changes on various (usually long) time scales.
 - Reliability is robustness to component failures.
 - Efficiency is robustness to resource scarcity.
 - Modularity is robustness to component rearrangements.

Core: Simple Forwarding to Network Functions

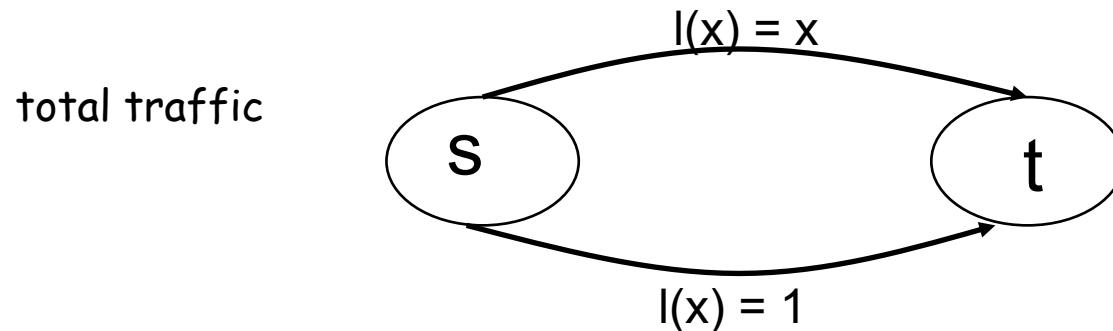
- Modern networks contain diverse types of equipment beyond simple routing/forwarding



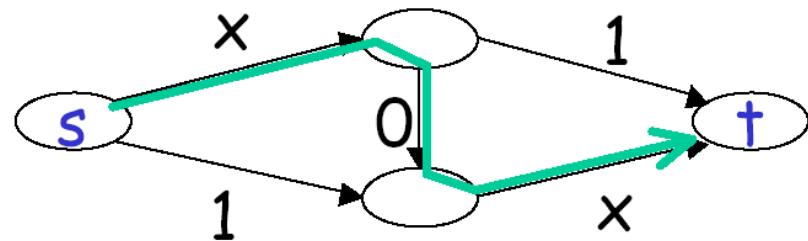
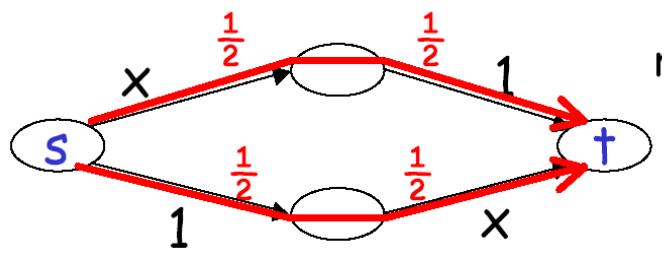
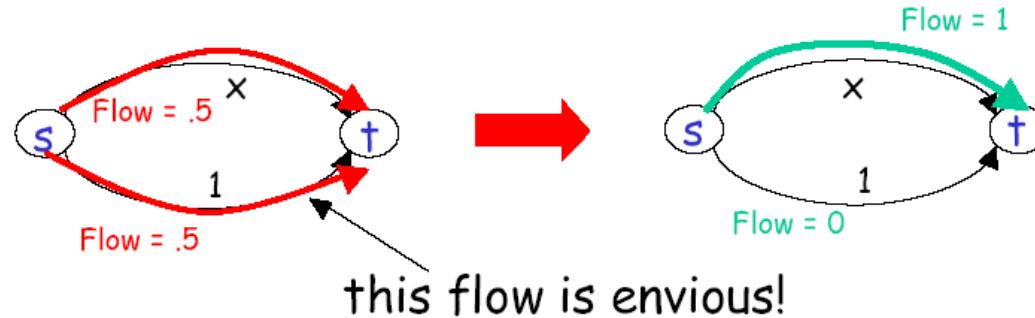
Centralized vs Decentralized (Price of Anarchy)



- Autonomous ("Selfish") App: Assume each link has a latency function $l_e(x)$: latency of link e when x amount of traffic goes through e :



Autonomous ("Selfish") App



Braess's paradox

Distributed vs Centralized



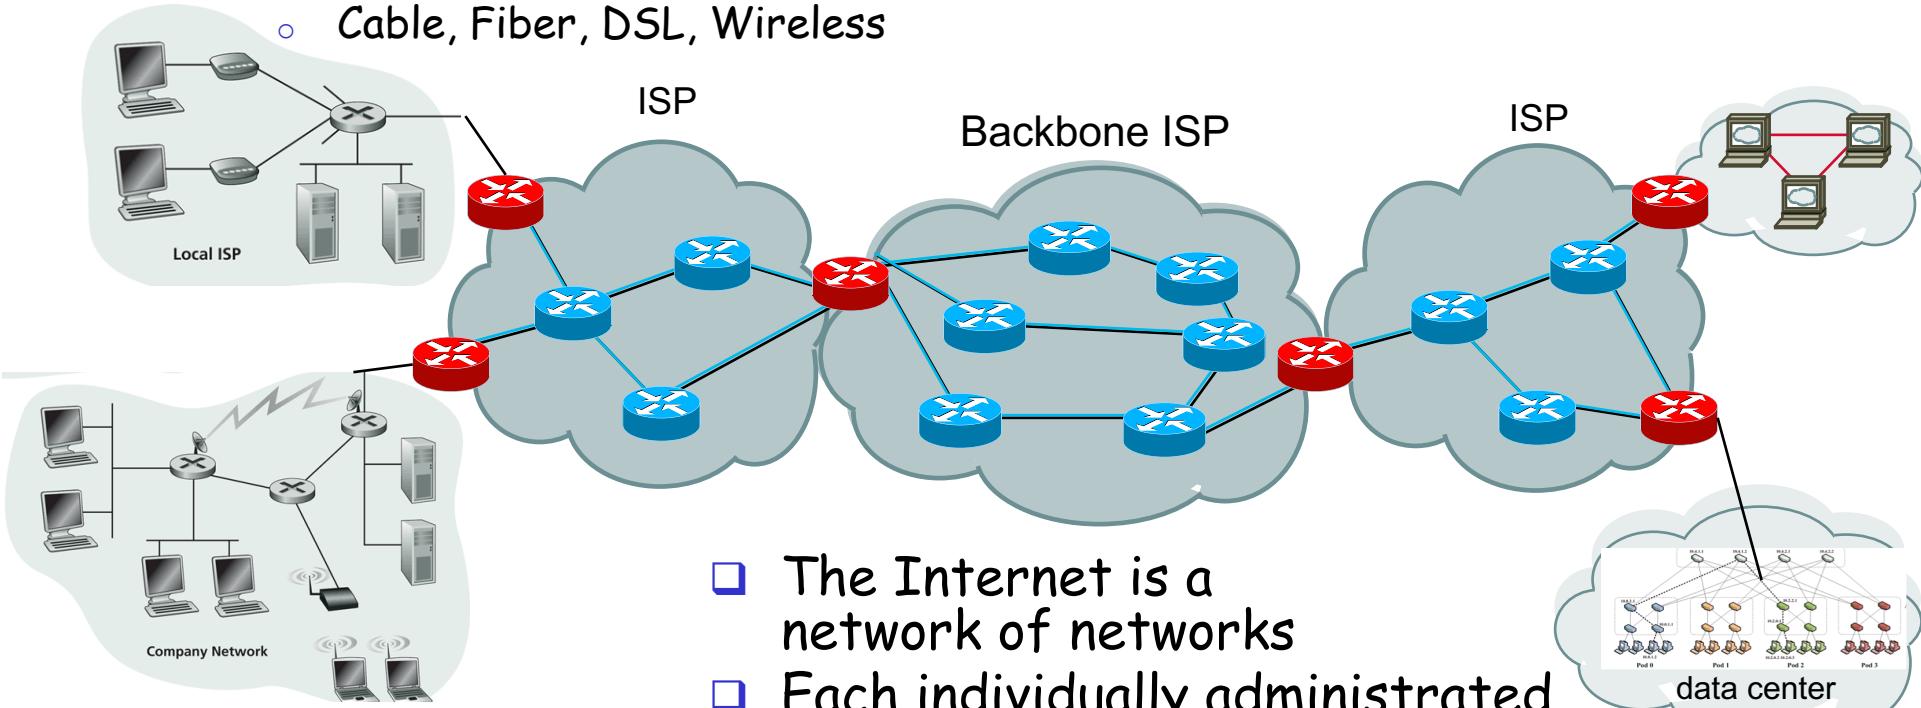
- Distributed computing is hard, e.g.,
 - FLP Impossibility Theorem
 - Arrow's Impossibility Theorem

- Achieved good design for only few specific tasks (e.g., state distribution, leader election). Hence, a trend in networking is Software Defined Networking, which is a way of moving away from generic distributed computing, by focusing on utilizing the few well-understood primitives, in particular logically centralized state.

Recall: Internet Physical Infrastructure

Residential access, e.g.,

- Cable, Fiber, DSL, Wireless



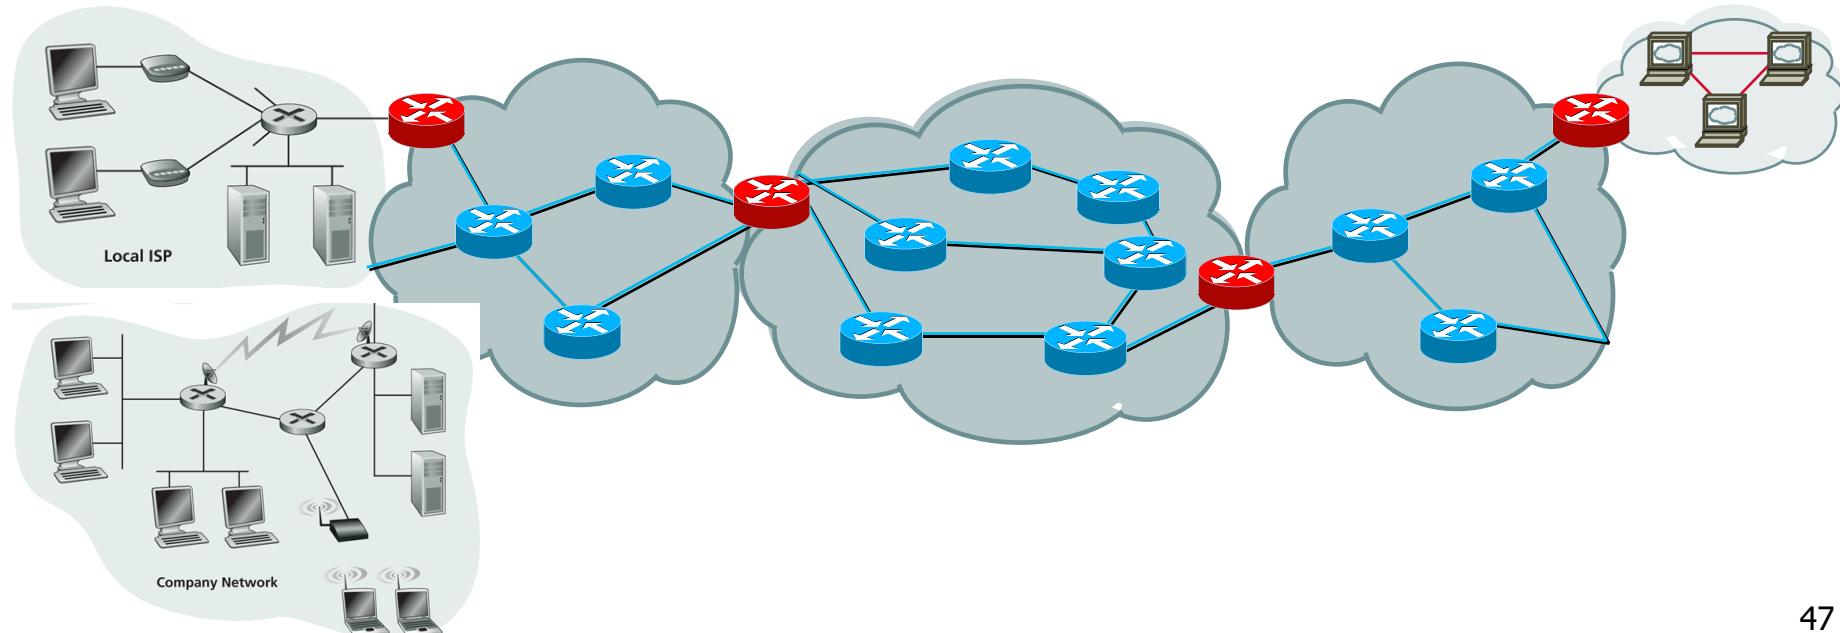
Campus access, e.g.,

- Ethernet, Wireless

- ❑ The Internet is a network of networks
- ❑ Each individually administrated network is called an Autonomous System (AS)
 - ~ 58000 ASes; Avg 5.7 hops;
(<http://bgp.potaroo.net/as2.0/bgp-active.html>)

Roadmap

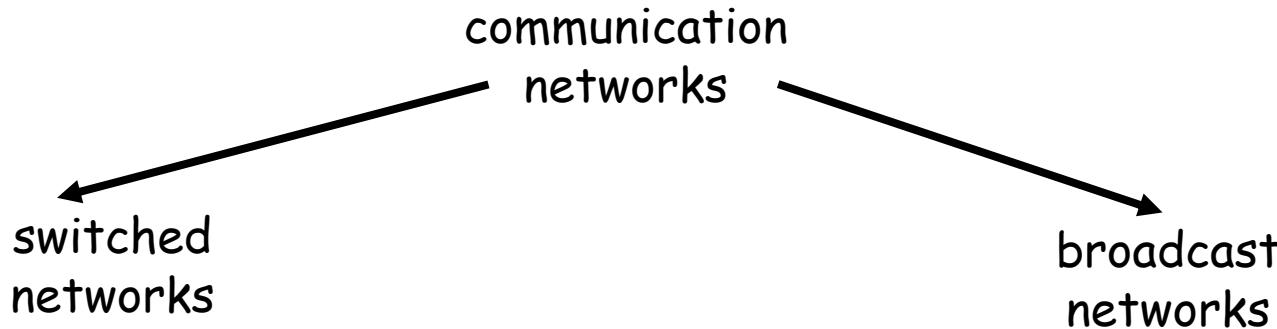
- So far we have looked at only the topology and physical connectivity of the Internet: a mesh of computers interconnected via various physical media
- **A basic question:** how are data (the bits) transferred through communication networks?



Outline

- Admin. and recaps
- A brief introduction to the Internet: past and present
- Challenges of Internet networks and apps
- *A taxonomy of communication networks*

Taxonomy of Communication Networks



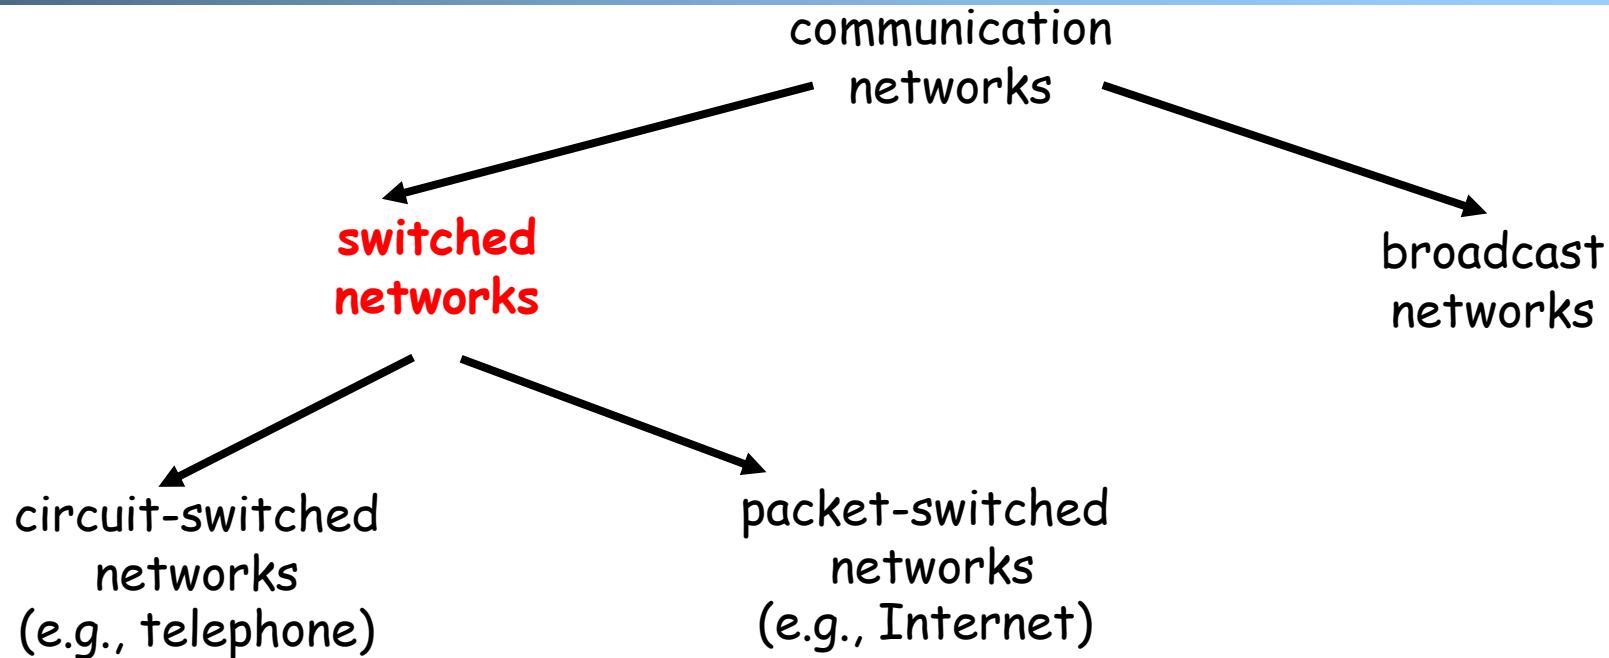
□ Broadcast networks

- nodes share a common channel; information transmitted by a node is received by all other nodes in the network
- examples: TV, radio

□ Switched networks

- information is transmitted to a small sub-set (usually only one) of the nodes

A Taxonomy of Switched Networks



- **Circuit switching:** dedicated circuit per call/session:
 - e.g., telephone, cellular voice
- **Packet switching:** data sent thru network in discrete “chunks”
 - e.g., Internet, cellular data

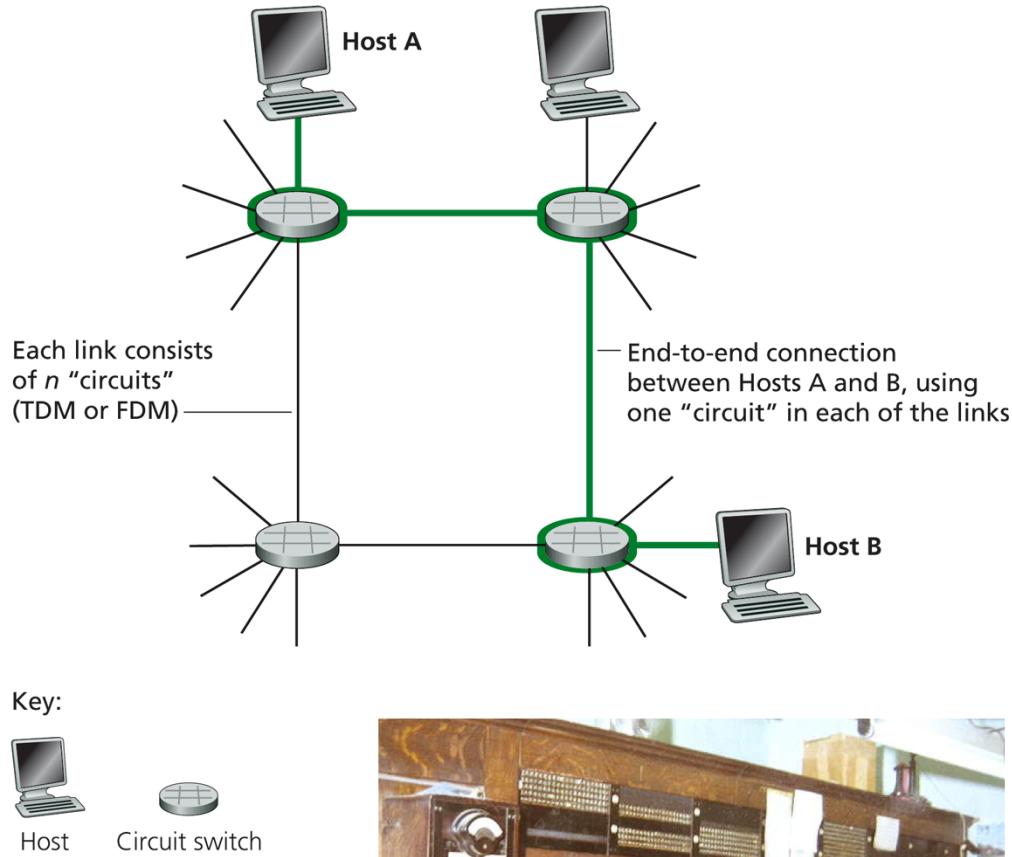
Outline

- Admin. and recaps
- A brief introduction to the Internet: past and present
- Challenges of Internet networks and apps
- *A taxonomy of communication networks*
 - *circuit switched networks*

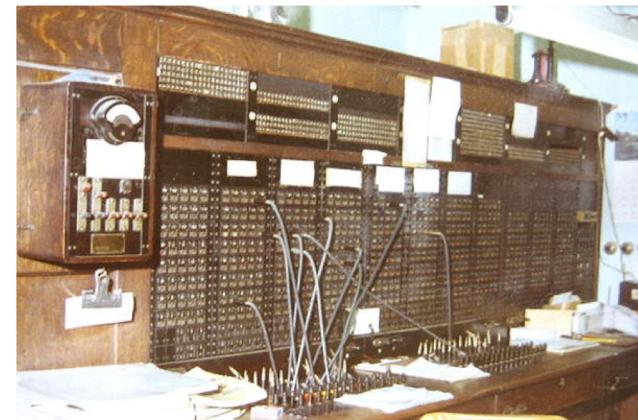
Circuit Switching

- Each link has a number of “circuits”
 - sometime we refer to a “circuit” as a channel or a line

- An end-to-end connection reserves one “circuit” at each link

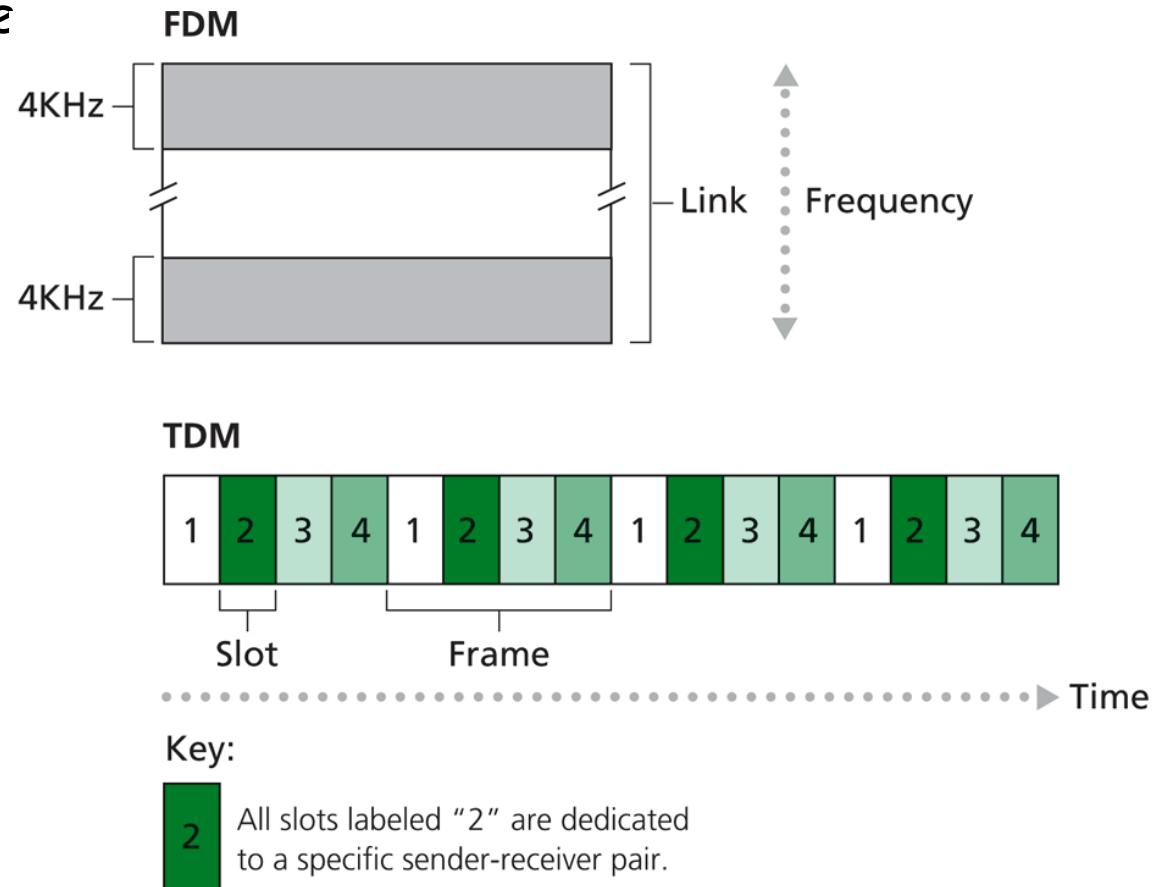


First commercial telephone switchboard was opened in 1878 to serve the 21 telephone customers in New Haven



Circuit Switching: Resources/Circuits (Frequency, Time and others)

- Divide link resource into “circuits”
 - frequency division multiplexing (FDM)
 - time division multiplexing (TDM)
 - others such as code division multiplexing (CDM), color/lambda division



Circuit Switching: The Process

- Three phases
 - circuit establishment
 - data transfer
 - circuit termination

Timing Diagram of Circuit Switching

