
Network Applications:

DNS;

Network App Programming--UDP

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

10/05/2021

Outline

- Admin. and recap
- DNS
 - High-level design
 - Details
 - Extensions/alternatives
- Network app programming
 - UDP

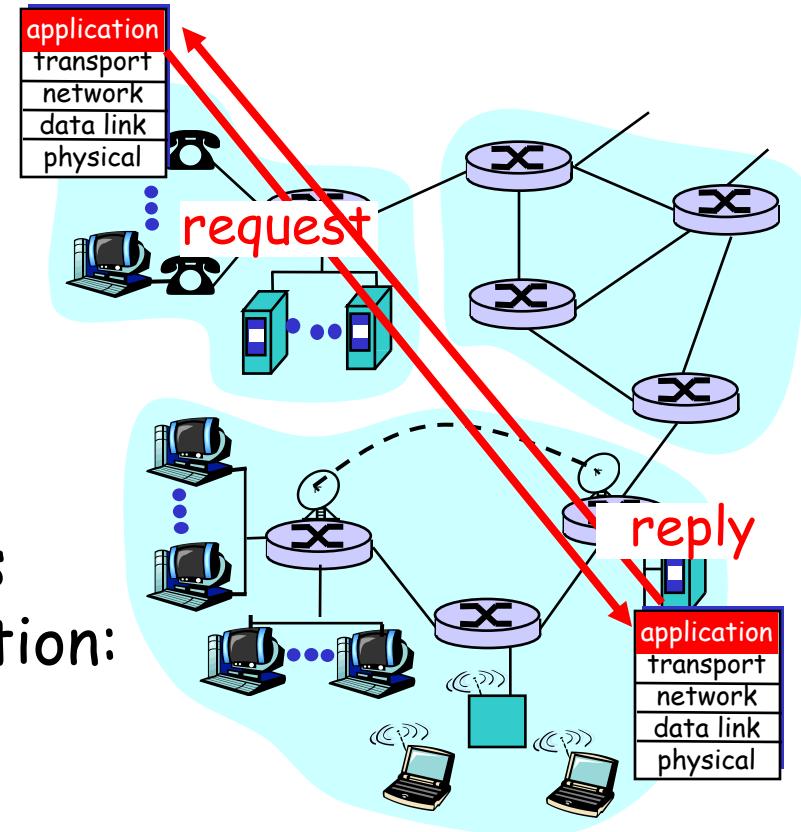
Admin

- Assignment Two linked on the schedule page
 - Oct. 28, in class or by email to the instructor

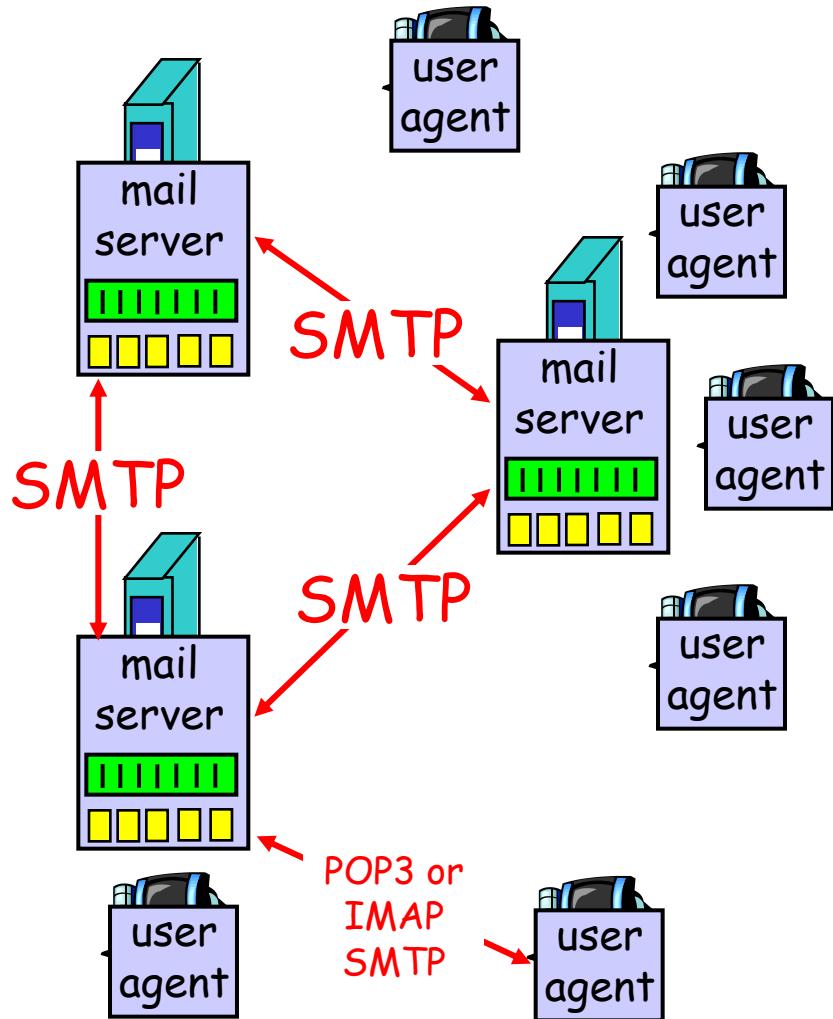
Recap: Client-Server Paradigm

- ❑ The basic paradigm of network applications is the client-server (C-S) paradigm

- ❑ Some key design questions to ask about a C-S application:
 - extensibility
 - scalability
 - robustness
 - security



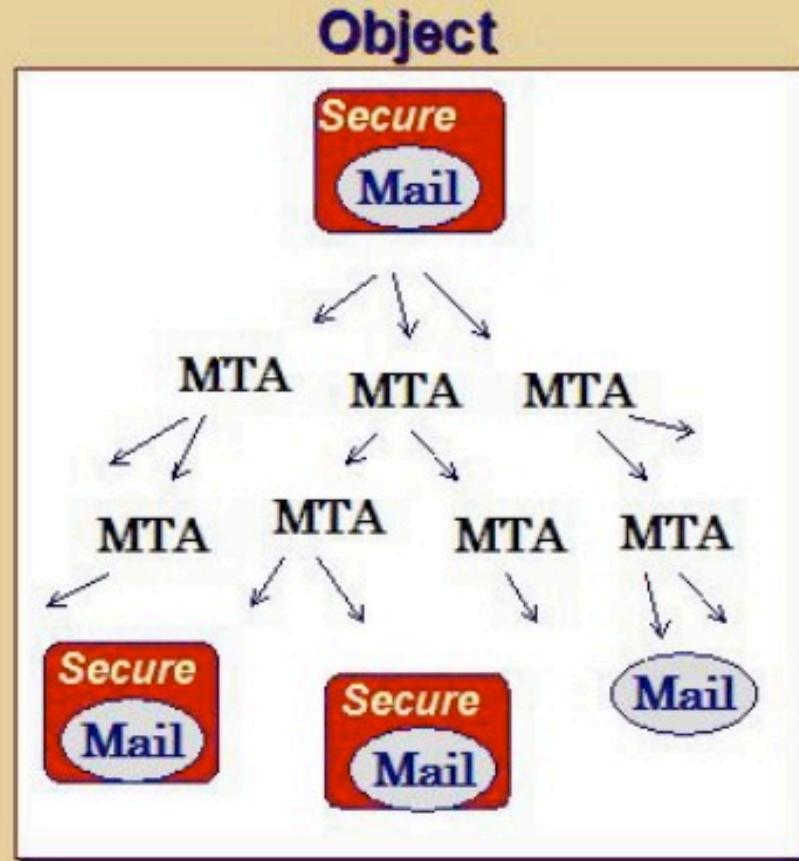
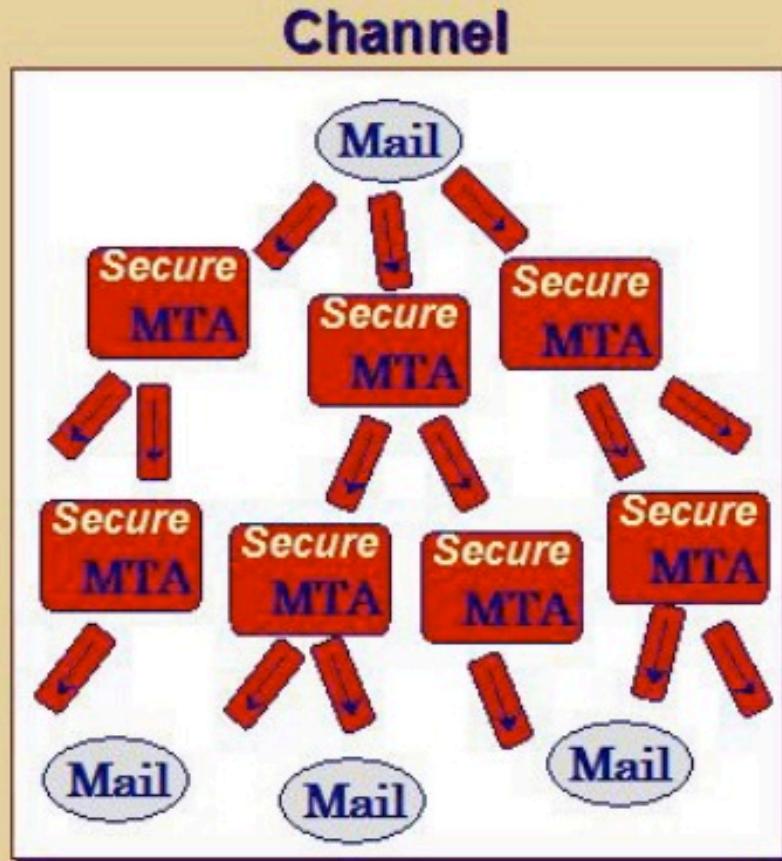
Recap: Email App



Some key design features of Email

- **Separate protocols for different functions**
 - email access (e.g., POP3, IMAP)
 - email transport (SMTP)
- **Separation of envelop and message body (end-to-end arguments)**
 - envelop: simple/basic requests to implement transport control;
 - message body: fine-grain control through ASCII header and message body
 - MIME type as self-describing data type
- **Status code** in response makes message easy to parse

Recap: Email Authentication Approaches



Sender Policy Frame (SPF)

DomainKeys Identified Mail (DKIM)
Authenticated Results Chain (ARC)

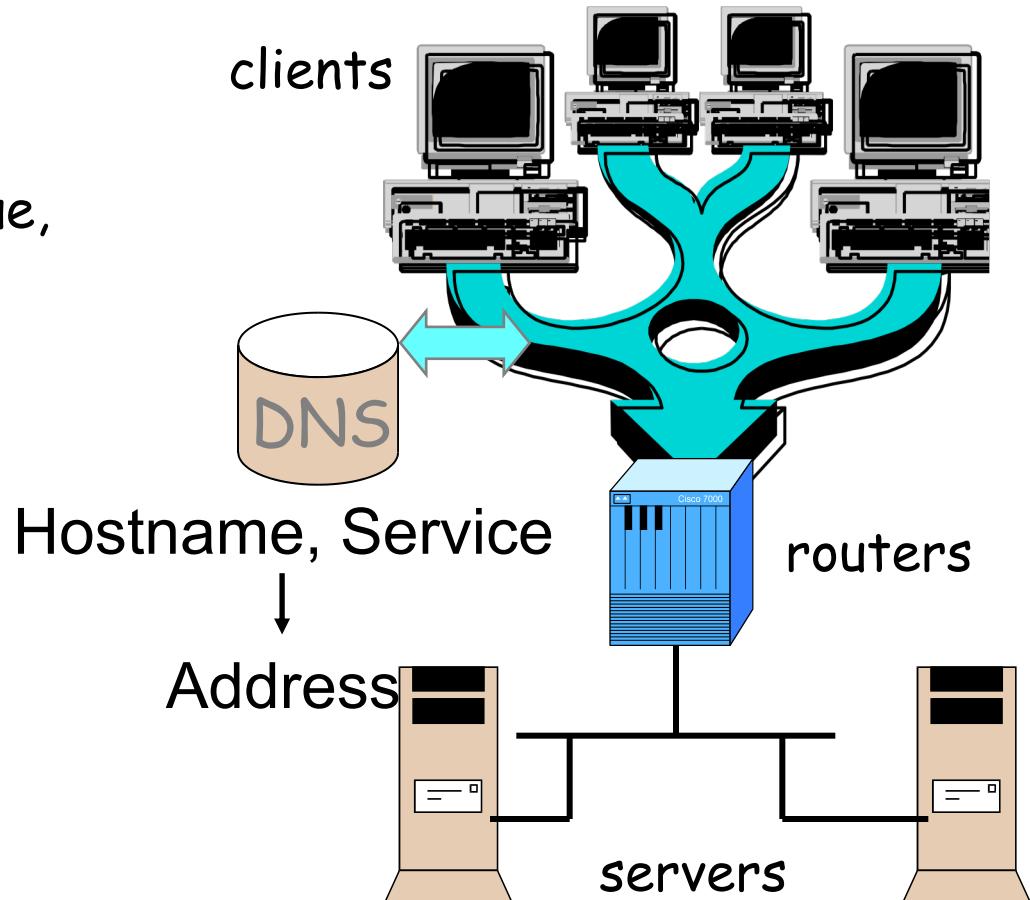
Summary: Some Key Remaining Issues about Email

- Basic: How to find the email server of a domain?
- Scalability/robustness: how to find multiple servers for the email domain?
- Security
 - SPF: How does SPF know if its neighbor MTA is a permitted sender of the domain?
 - DKIM: How does DKIM retrieve the public key of the author domain?

Recap: Domain Name System (DNS)

□ Function

- map between (domain name, service) to value,
e.g.,
 - (xmu.edu.cn, addr)
→ 210.34.0.35
 - (xmu.edu.cn, email)
→ cmsn1.xmu.edu.cn



Recap: DNS Records

DNS: stores resource records (**RR**)

RR format: (**name**, **type**, **value**, **ttl**)

- Type=A
 - **name** is hostname
 - **value** is IP address
- Type=NS
 - **name** is domain (e.g. xmu.edu.cn)
 - **value** is the name of the authoritative name server for this domain
- Type=TXT
 - general txt
- Type=CNAME
 - **name** is an alias of a “canonical” (real) name
 - **value** is canonical name
- Type=MX
 - **value** is hostname of mail server associated with **name**
- Type=SRV
 - general extension for services
- Type=PTR
 - a pointer to another name

Recap: Observations

- MX can return multiple servers
- DNS may rotate the servers in answer
- Address can also return multiple addresses
- SPF is encoded as the txt type

Outline

- Admin. and recap
- DNS
 - High-level design
 - *Details*

DKIM Example

- Send email from hotmail and check message

DKIM Example

- DKIM / ARC:

Msg: ARC-Message-Signature: i=1; a=rsa-sha256;
c=relaxed/relaxed; d=microsoft.com; s=arcselector9901;
h=From:Date:Subject:Message-ID:Content-Type:MIME-
Version;
bh=b091TxHI+4MjgAusrfg0EWGiDmvQ5hZRZ/aqb1MKLY8
=; ...

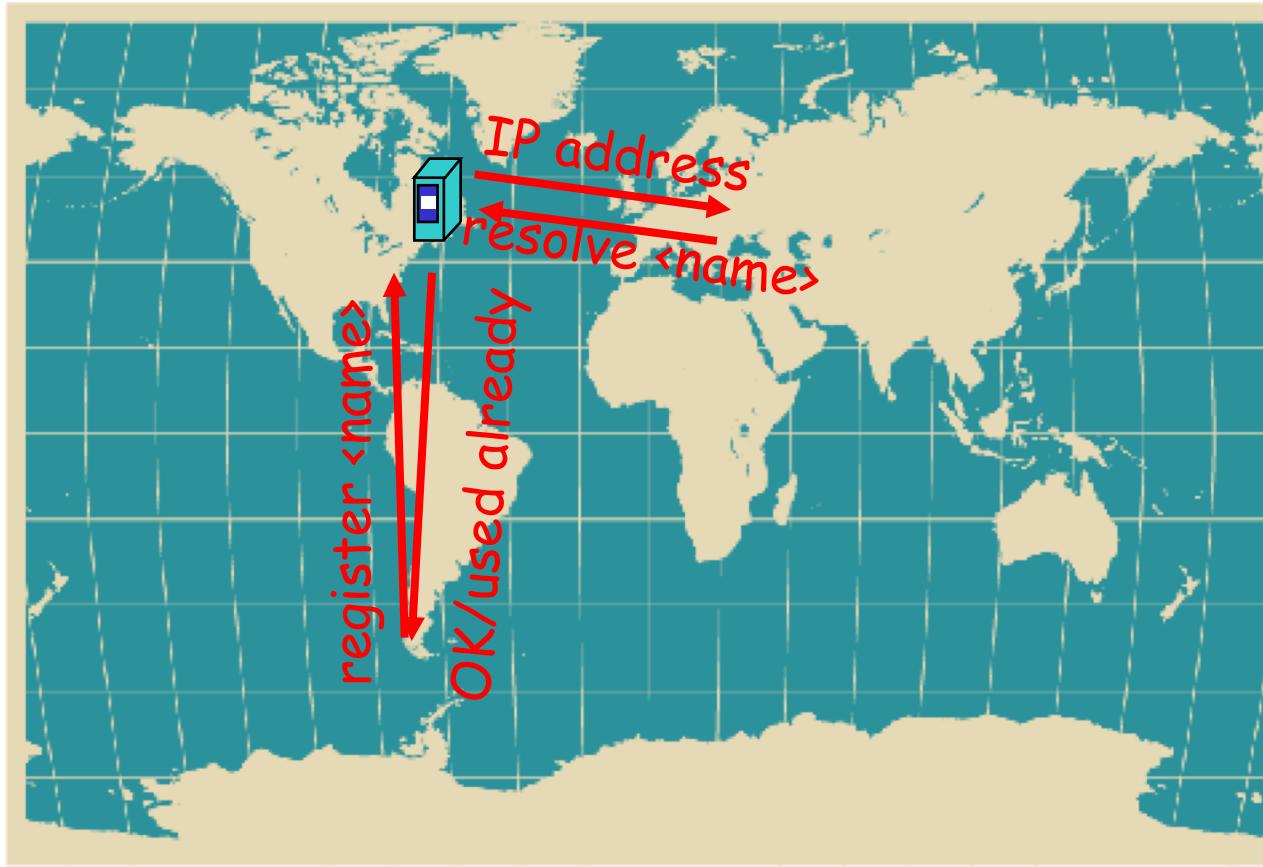
DKIM-Signature: v=1; a=rsa-sha256; c=relaxed/relaxed;
d=hotmail.com; s=selector1; h=From:Date:Subject:Message-
ID:Content-Type:MIME-Version:X-MS-Exchange-
SenderADCheck;...

- Query: dig arcselector9901._domainkey.microsoft.com txt
- DKIM introduces a session key to allow multiple public keys
 - <session>._domainkey.<domain>

DNS Design: Dummy Design

- DNS itself can be considered as a client-server system as well
- How about a dummy design: introducing one super Internet DNS server?

THE DNS server of the Internet

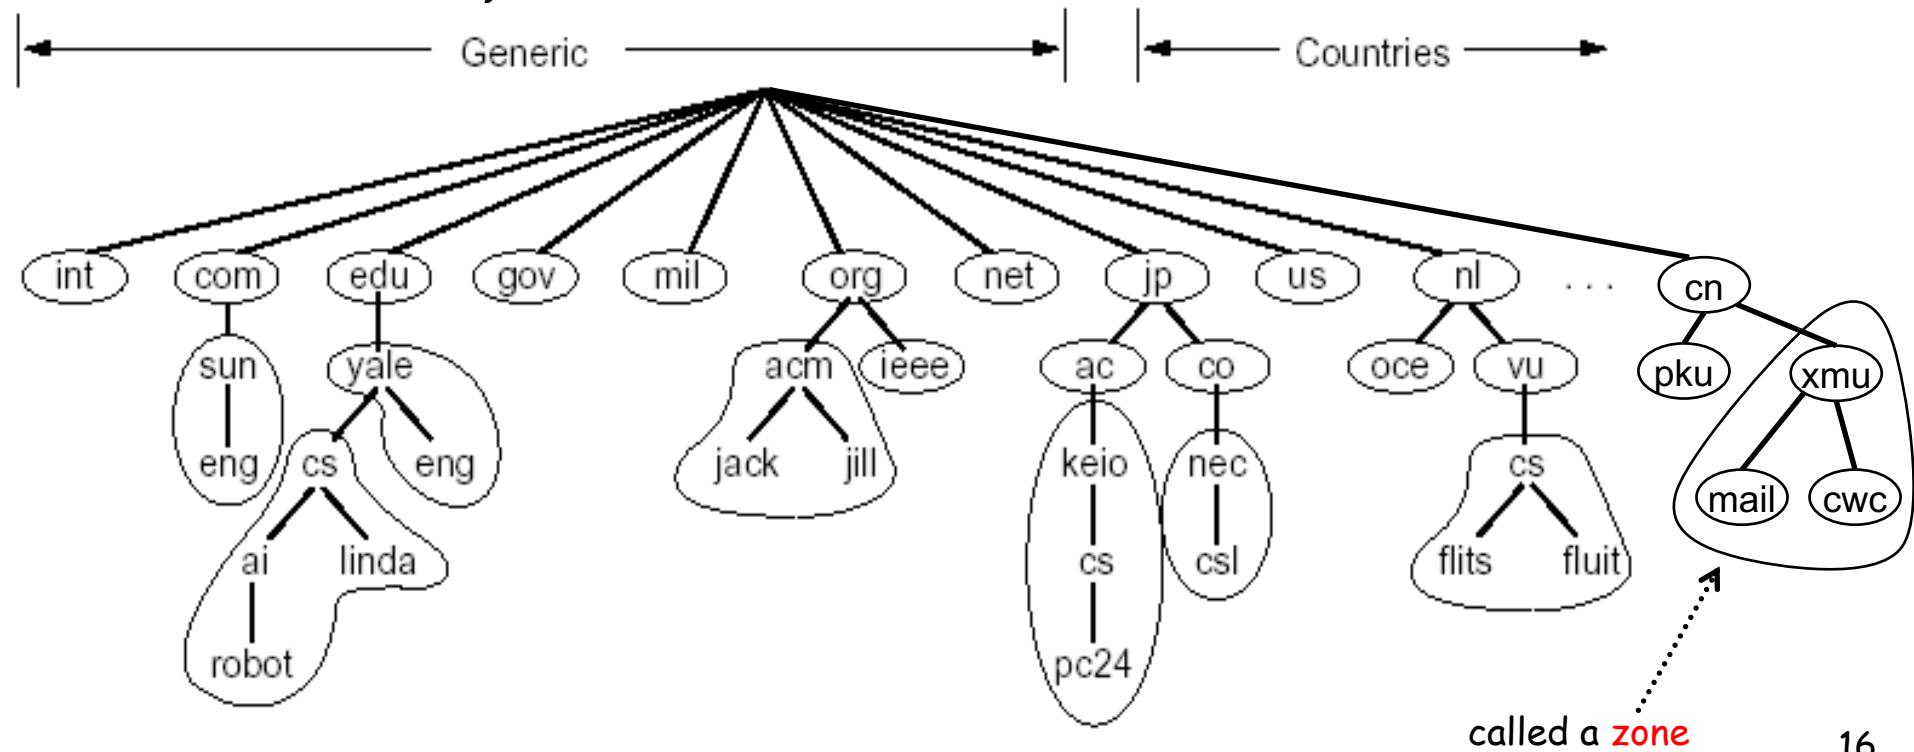


Problems of a Single DNS Server

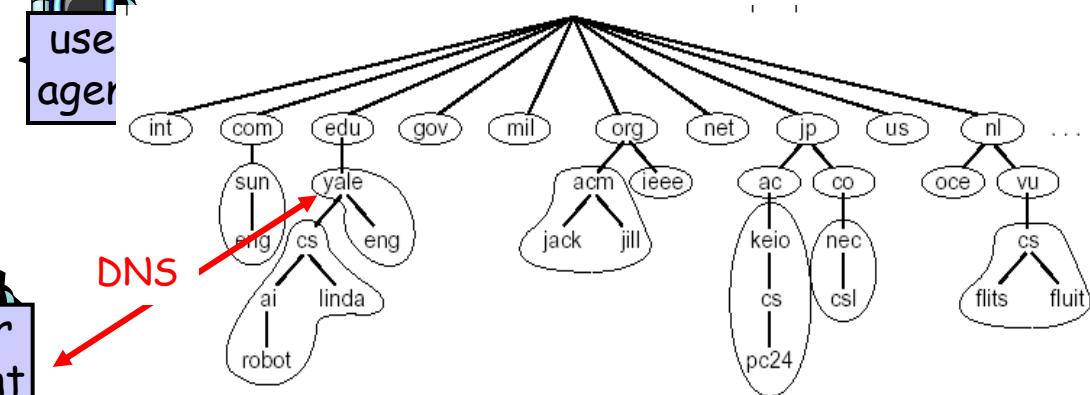
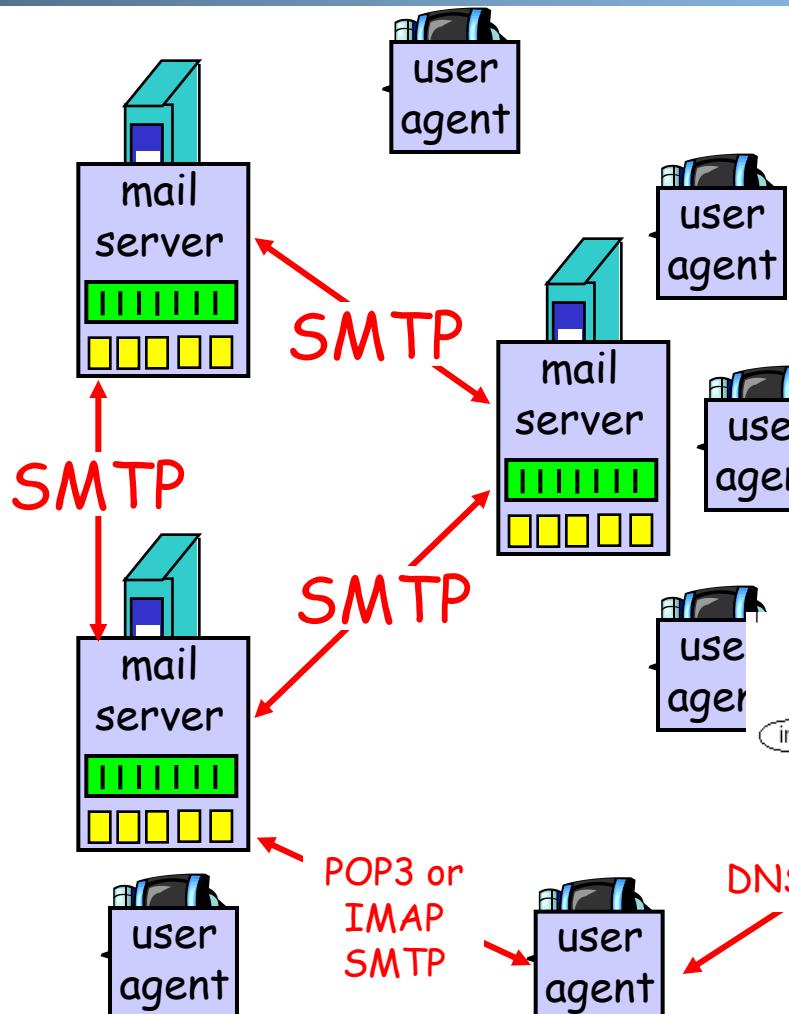
- Scalability and robustness bottleneck
- Administrative bottleneck

DNS: Distributed Management of the Domain Name Space

- ❑ A distributed database managed by authoritative name servers
 - divided into zones, where each zone is a sub-tree of the global tree
 - each zone has its own **authoritative name servers**
 - an authoritative name server of a zone may **delegate** a subset (i.e. a sub-tree) of its zone to another name server



Email Architecture + DNS



Root Zone and Root Servers

- The root zone is managed by the root name servers

- 13 root name servers worldwide

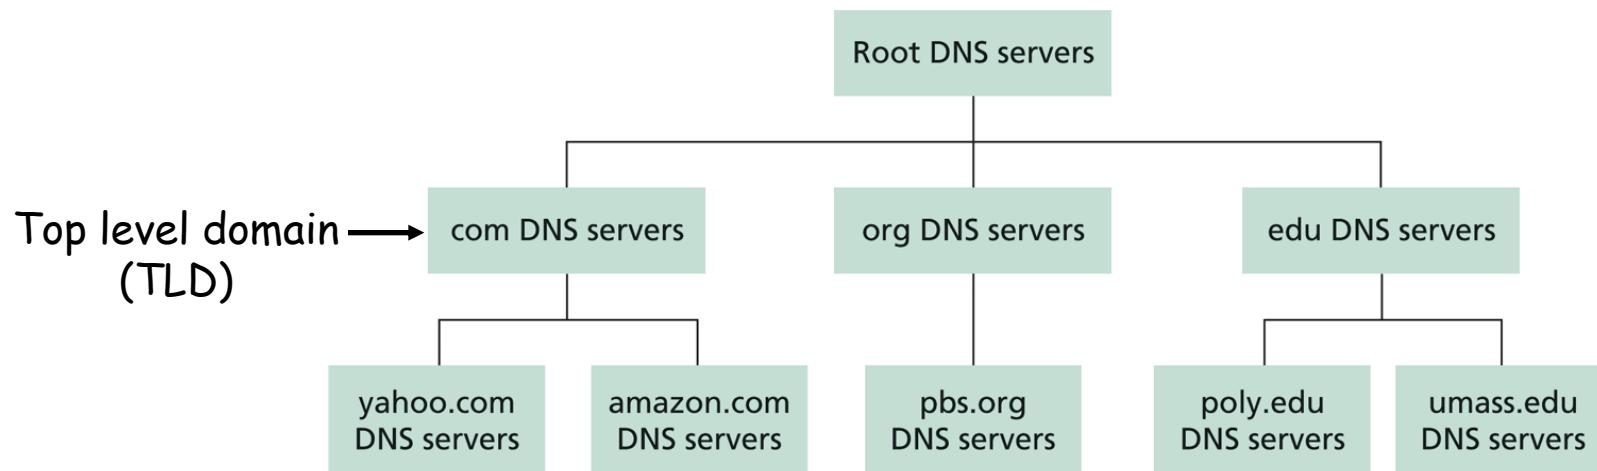
- a. Verisign, Dulles, VA
 - c. Cogent, Herndon, VA (also Los Angeles)
 - d. U Maryland College Park, MD
 - g. US DoD Vienna, VA
 - h. ARL Aberdeen, MD
 - j. Verisign, (11 locations)
 - i. Autonomica, Stockholm (plus 3 other locations)
 - k. RIPE London (also Amsterdam, Frankfurt)
 - m. WIDE Tokyo
 - e. NASA Mt View, CA
 - f. Internet Software C.
Palo Alto, CA
(and 17 other locations)
 - b. USC-ISI Marina del Rey, CA
 - l. ICANN Los Angeles, CA



See <http://root-servers.org/> for more details

Linking the Name Servers

- Each name server knows the addresses of the root servers
- Each name server knows the addresses of its immediate children (i.e., those it delegates)



Q: how to query a hierarchy?

DNS Message Flow:

Two Types of Queries

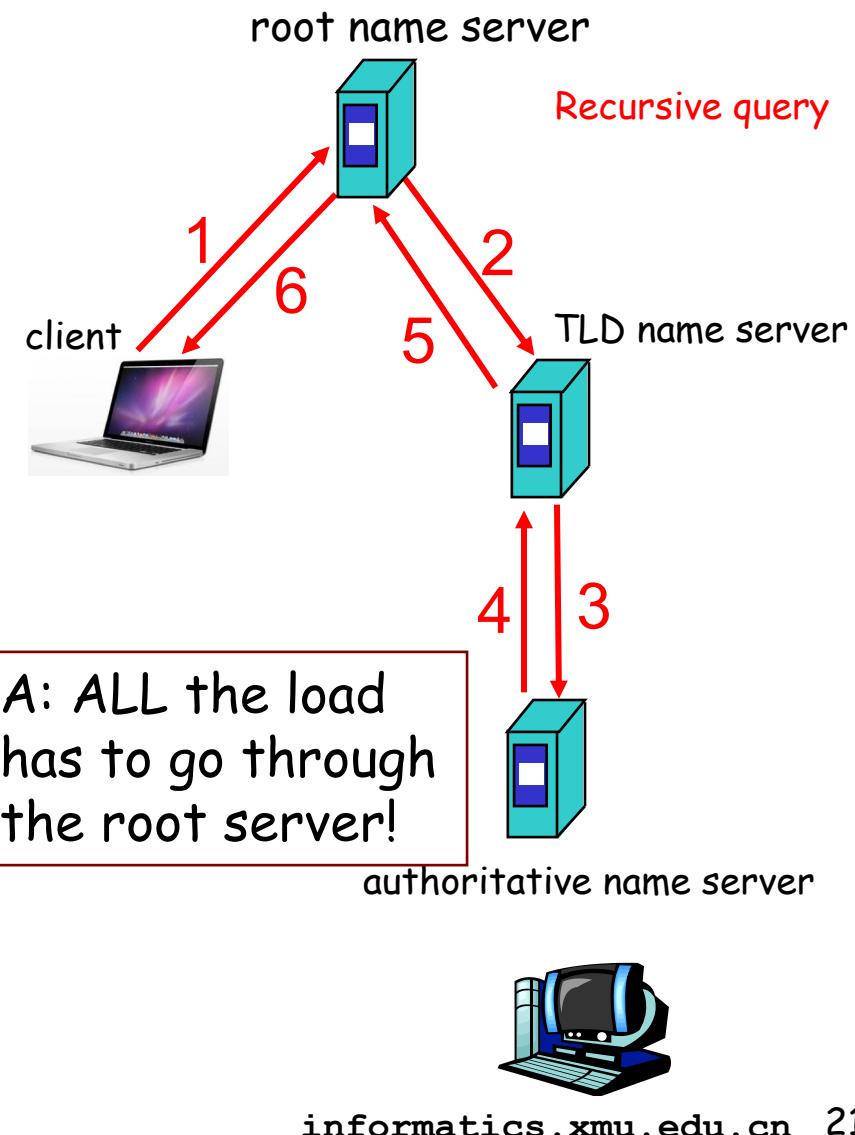
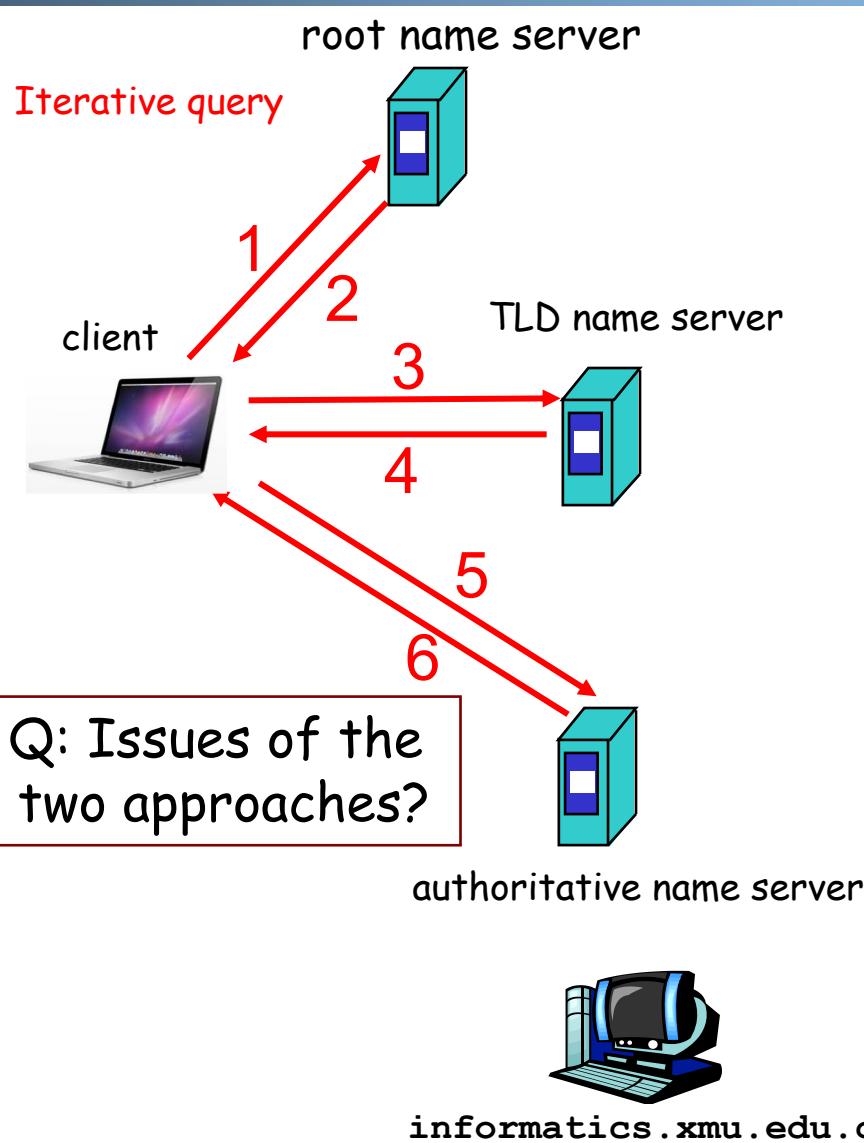
Recursive query:

- ❑ The contacted name server resolves the name completely

Iterated query:

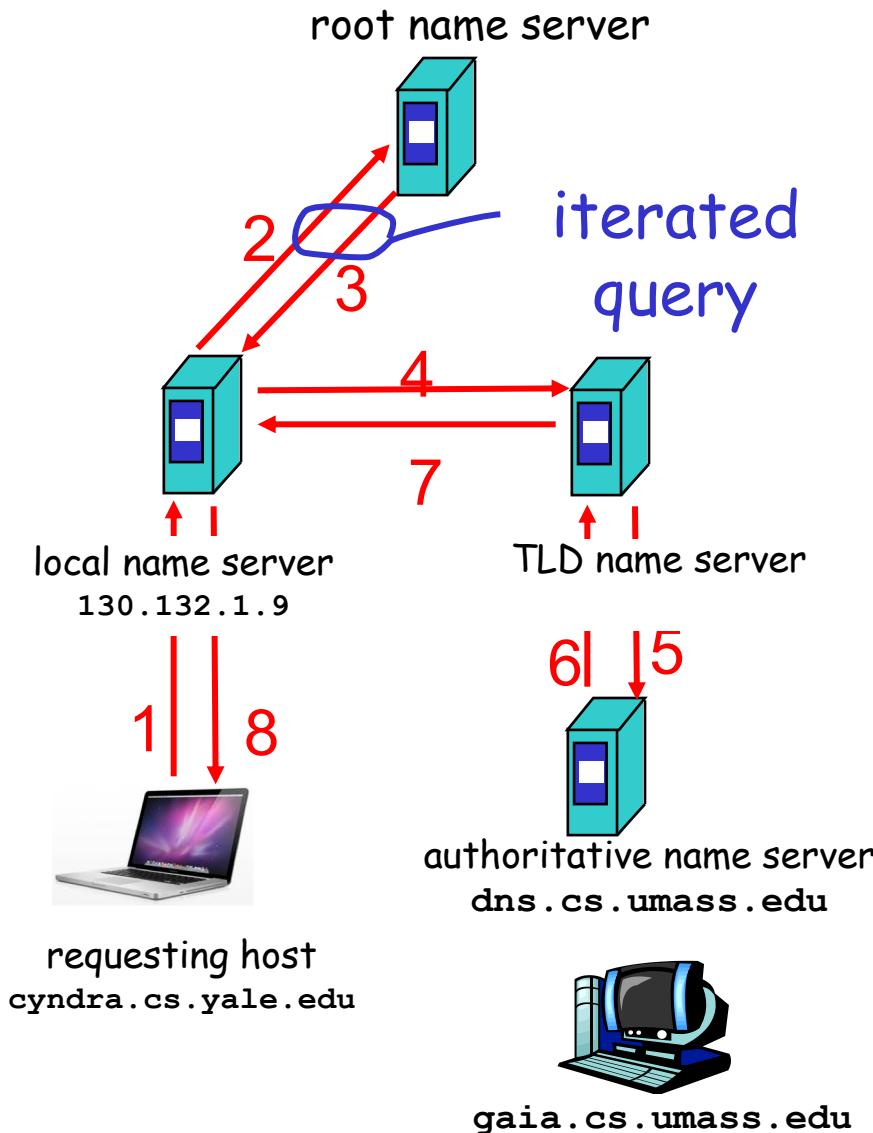
- ❑ Contacted server replies with name of server to contact
 - “I don’t know this name, but ask this server”

Two Extreme DNS Message Flows



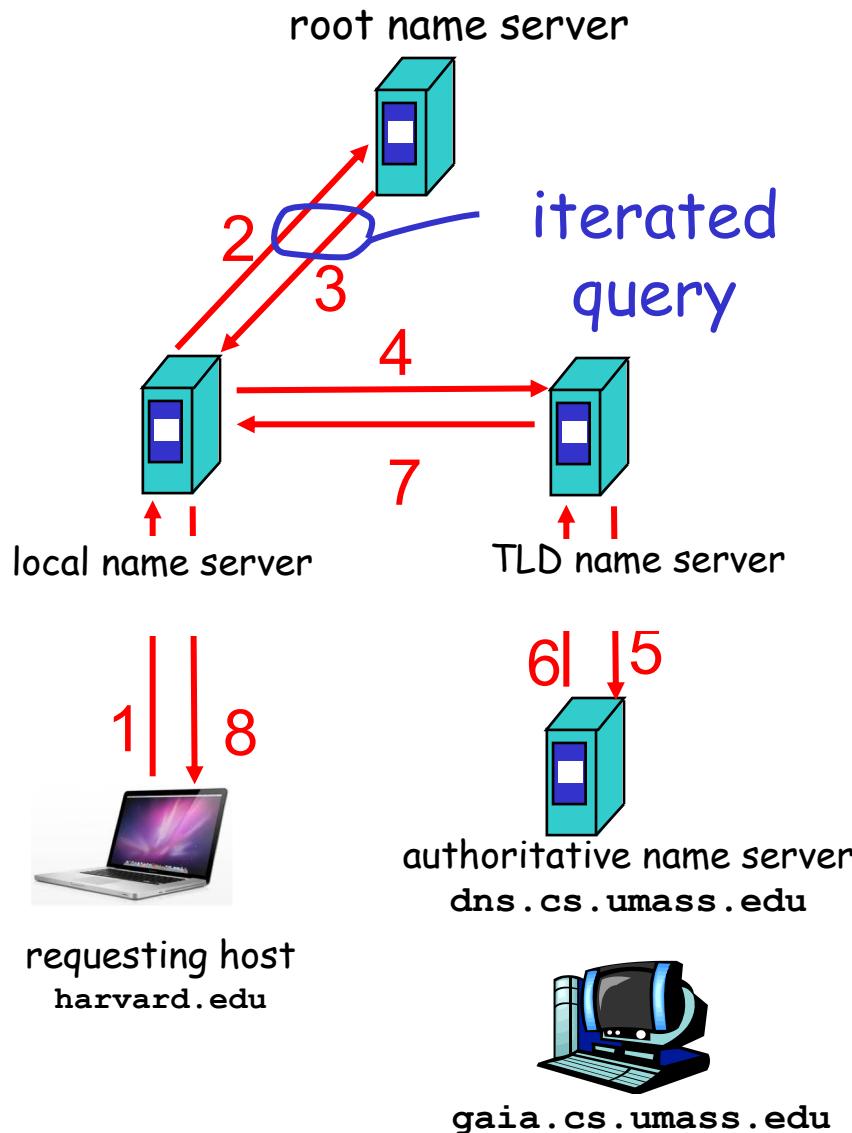
Typical DNS Message Flow: The Hybrid Case

- Host knows only local name server
- Local name server is learned from DHCP, or configured, e.g. /etc/resolv.conf
- Local DNS server helps clients resolve DNS names



Typical DNS Message Flow: The Hybrid Case

- Host knows only local name server
- Local name server is learned from DHCP, or configured, e.g. /etc/resolv.conf
- Local DNS server helps clients resolve DNS names
- Benefits of local name servers (often called **resolvers**)
 - simplifies client
 - caches/reuses results

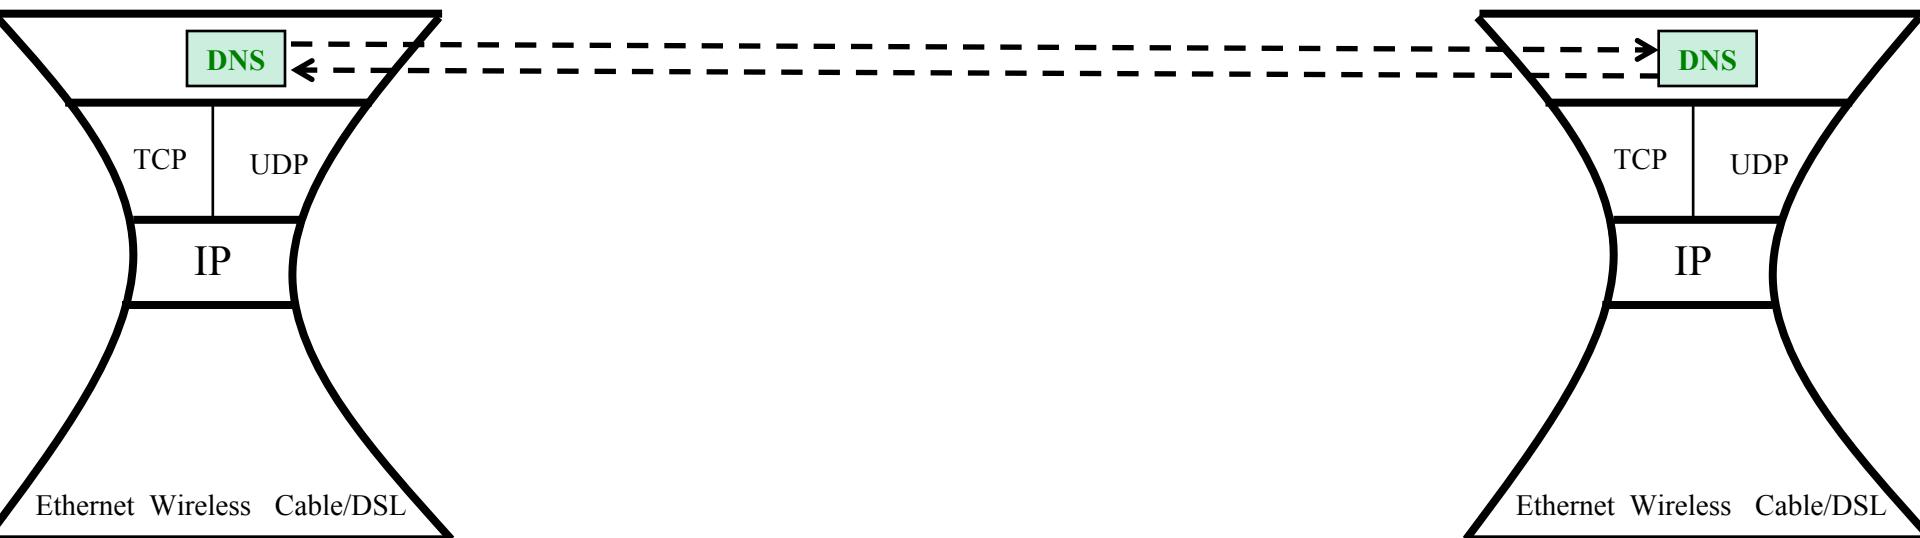


Outline

- Admin. and recap
- DNS
 - High-level design
 - *Details*

DNS Message Format?

Basic encoding decisions: UDP/TCP,
how to encode domain name, how to
encode answers...



Observing DNS Messages

- Capture the messages
 - DNS server is at port 53
 - Display and clear DNS cache
 - MacOS: <https://support.apple.com/en-us/HT202516>
sudo killall -HUP mDNSResponder
 - Ubuntu:
sudo systemd-resolve --flush-caches
sudo systemd-resolve --statistics
 - Try to load the dns-capture file from class Schedule page, if you do not want live capture

DNS Protocol, Messages

DNS protocol : typically over UDP (can use TCP);
query and *reply* messages, both with the **same** message format

Identification	Flags	12 bytes
Number of questions	Number of answer RRs	
Number of authority RRs	Number of additional RRs	
Questions (variable number of questions)		Name, type fields for a query
Answers (variable number of resource records)		RRs in response to query
Authority (variable number of resource records)		Records for authoritative servers
Additional information (variable number of resource records)		Additional “helpful” info that may be used

DNS Details

- Header (Sec. 4.1.1 of
<https://www.ietf.org/rfc/rfc1035.txt>)
- Encoding of questions (Sec. 4.1.2):
 - [Label-length label-chars]
- Encoding of answers (Sec. 4.1.3)
 - Pointer format
(<http://www.iana.org/assignments/dns-parameters/dns-parameters.xhtml>)
- See example DNS packets

Name Encoding

▼ Queries

▼ xmuxmu.edu.cn: type A, class IN

Name: xmuxmu.edu.cn

[Name Length: 10]

[Label Count: 3]

Type: A (Host Address) (1)

Class: IN (0x0001)

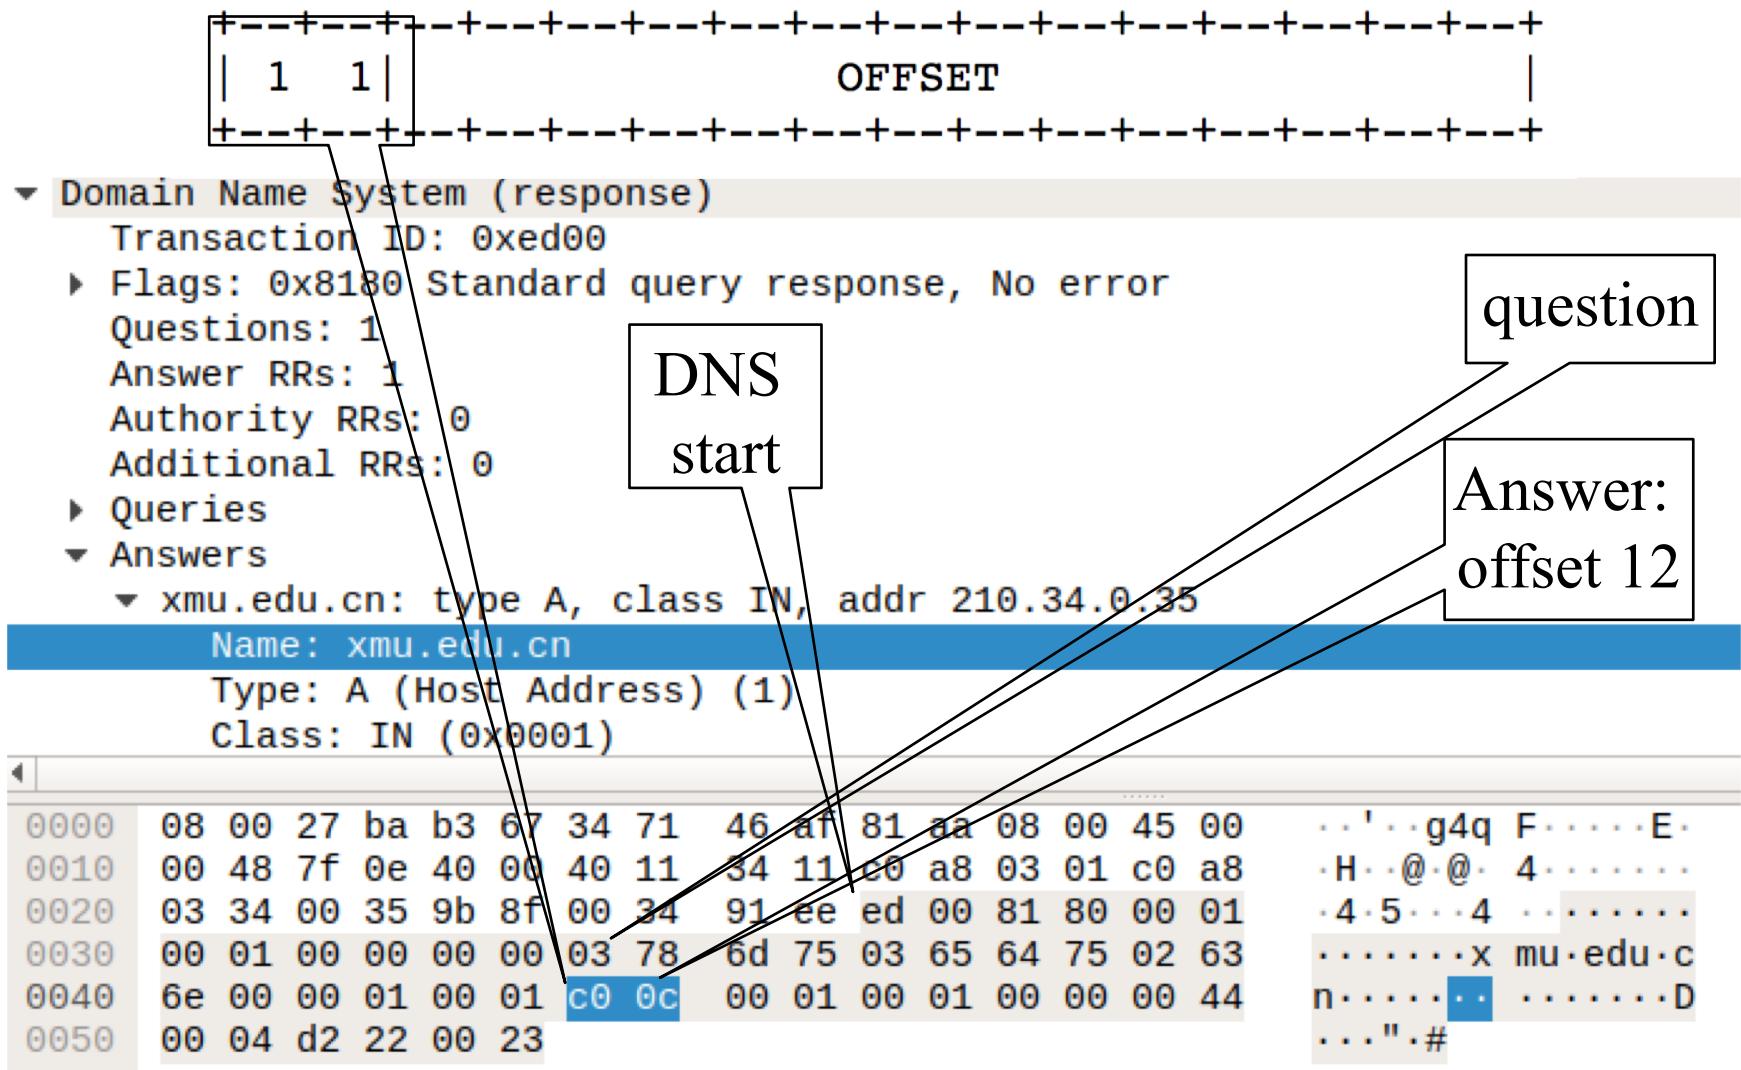
[Response In: 3]

	Hex															ASCII														
0000	34	71	46	af	81	aa	08	00	27	ba	b3	67	08	00	45	00	4qF	'	...	g	..	E	
0010	00	38	02	d4	40	00	40	11	b0	5b	c0	a8	03	34	c0	a8	·8	··@·@·	·[....	4	
0020	03	01	9b	8f	00	35	00	24	87	bb	ed	00	01	00	00	00	01	5	\$
0030	00	00	00	00	00	00	03	78	6d	75	03	65	64	75	02	63	x	mu	·edu	·c
0040	6e	00	00	01	00	01												n	·

Diagram illustrating the name encoding:

- The name "xmuxmu.edu.cn" is broken down into its labels: "cn", "xmu", "edu", and "cn".
- The "length" of each label is indicated by the number of bytes it occupies in the hex dump.
- Red arrows point from the labels to their corresponding byte ranges in the hex dump.

Message Compression (Label Pointer)



Recap: DNS Protocol, Messages

Many features: typically over UDP (can use TCP); *query* and *reply* messages with the same message format; *length/content encoding of names*; *simple compression*; *additional info as server push*

Identification	Flags	12 bytes
Number of questions	Number of answer RRs	
Number of authority RRs	Number of additional RRs	
Questions (variable number of questions)		Name, type fields for a query
Answers (variable number of resource records)		RRs in response to query
Authority (variable number of resource records)		Records for authoritative servers
Additional information (variable number of resource records)		Additional “helpful” info that may be used

What DNS did Right?

- Hierarchical delegation avoids central control, improving manageability and scalability
- Redundant servers improve robustness
 - see <http://www.internetnews.com/dev-news/article.php/1486981> for DDoS attack on root servers in Oct. 2002 (9 of the 13 root servers were crippled, but only slowed the network)
- Caching reduces workload and improves robustness
- Proactive answers reduce # queries on server and latency on client

Problems of DNS

- Simple query model, relatively static resource values and types make it harder to implement generic service discovery
 - e.g., service discovery of all printers
 - Although theoretically you can update the values of the records, it is rarely enabled
- Early binding (separation of DNS query from application query) does not work well in mobile, dynamic environments
 - e.g., load balancing, locate the nearest printer
- Each local domain needs servers, but an ad hoc domain may not have a DNS server

Outline

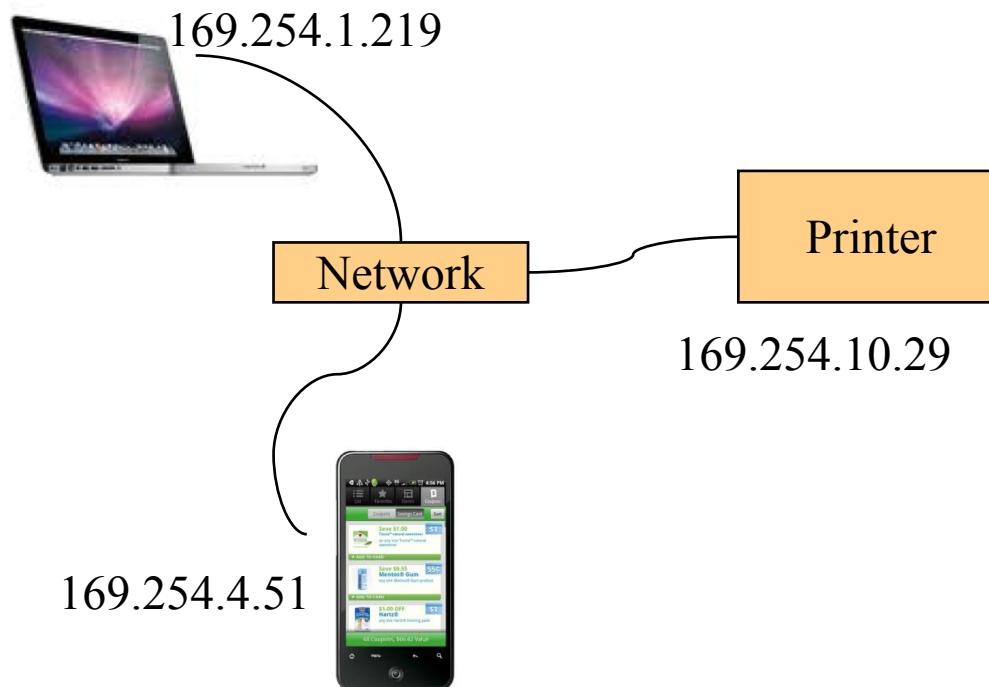
- Admin. and recap
- DNS
 - High-level design
 - Details
 - *Extensions/alternatives*

Discussions

- What extension(s) to standard DNS operations do we need to allow service discovery, say to implement Bonjour (discover all local printers)?
 - each printer needs to provide the following info: host, port, printer info (e.g., support postscript)

DNS-Service Discovery

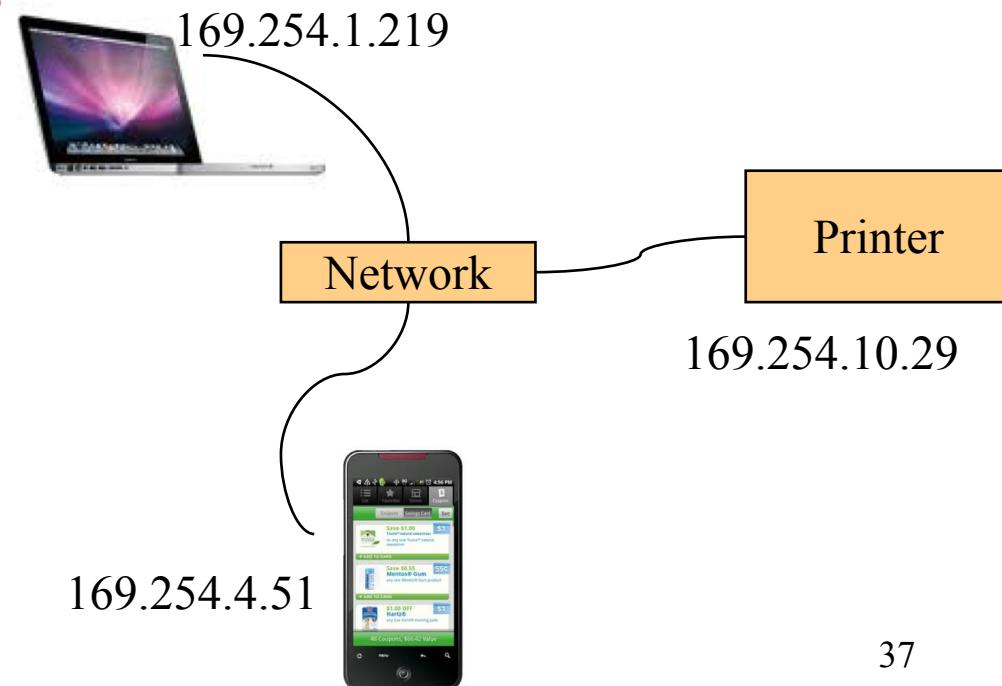
- ❑ Leverage DNS message format, but each node can announce its own services



Realizing DNS-SD without Central DNS Server: mDNS

□ Multicast in a small world

- no central address server
 - each node is a responder
- link-local addressing
 - send to **multicast** address: **224.0.0.251**



Example

- Use the avahi-publish-service command on Ubuntu as example
 - Advertise (register) an LPR printer on port 515

```
avahi-publish-service test _printer._tcp . 515  
pdl=application/postscript
```

txt for additional data

Name of instance providing the service

<type_service>.
<transport>

port

Example

- Use the dns-sd command on Mac as example
 - Advertise (register) an LPR printer on port 515

```
dns-sd -R "test" _printer._tcp . 515  
pdl=application/postscript
```

Name of instance providing the service

<type_service>. <transport>

domain (. means default, which is local)

port

Txt for additional data

1	0.0.0... fe80::5052:5af:92fa...	ff02::fb	MDNS	179 Standard query 0x0000 ANY test._printer._tcp.local, "QM" question SRV 0 0 515 qiao-VirtualBox.l...
2	0.0.0... 192.168.3.52	224.0.0.251	MDNS	159 Standard query 0x0000 ANY test._printer._tcp.local, "QM" question SRV 0 0 515 qiao-VirtualBox.l...
3	0.255... fe80::5052:5af:92fa...	ff02::fb	MDNS	179 Standard query 0x0000 ANY test._printer._tcp.local, "QM" question SRV 0 0 515 qiao-VirtualBox.l...
4	0.255... 192.168.3.52	224.0.0.251	MDNS	159 Standard query 0x0000 ANY test._printer._tcp.local, "QM" question SRV 0 0 515 qiao-VirtualBox.l...
5	0.501... fe80::5052:5af:92fa...	ff02::fb	MDNS	179 Standard query 0x0000 ANY test._printer._tcp.local, "QM" question SRV 0 0 515 qiao-VirtualBox.l...
6	0.501... 192.168.3.52	224.0.0.251	MDNS	159 Standard query 0x0000 ANY test._printer._tcp.local, "QM" question SRV 0 0 515 qiao-VirtualBox.l...
7	0.701... 192.168.3.52	224.0.0.251	MDNS	248 Standard query response 0x0000 TXT, cache flush PTR test._printer._tcp.local SRV, cache flush 0...
8	0.701... fe80::5052:5af:92fa...	ff02::fb	MDNS	252 Standard query response 0x0000 TXT, cache flush PTR test._printer._tcp.local SRV, cache flush 0...
9	1.919... 192.168.3.52	224.0.0.251	MDNS	248 Standard query response 0x0000 TXT, cache flush PTR test._printer._tcp.local SRV, cache flush 0...
10	1.919... fe80::5052:5af:92fa...	ff02::fb	MDNS	252 Standard query response 0x0000 TXT, cache flush PTR test._printer._tcp.local SRV, cache flush 0...
11	4.127... 192.168.3.52	224.0.0.251	MDNS	248 Standard query response 0x0000 TXT, cache flush PTR test._printer._tcp.local SRV, cache flush 0...
12	4.127... fe80::5052:5af:92fa...	ff02::fb	MDNS	252 Standard query response 0x0000 TXT, cache flush PTR test._printer._tcp.local SRV, cache flush 0...
13	11.73... fe80::5052:5af:92fa...	ff02::fb	MDNS	183 Standard query 0x0000 PTR _services._dns-sd._udp.local, "QM" question PTR _companion-link._tcp...
14	11.73... 192.168.3.52	224.0.0.251	MDNS	305 Standard query 0x0000 PTR _services._dns-sd._udp.local, "QM" question TXT Qiao\342\200\231s Mac...
15	11.83... 192.168.3.38	224.0.0.251	MDNS	230 Standard query response 0x0000 PTR _companion-link._tcp.local SRV, cache flush 0 0 54580 Qiaos...
16	12.71... fe80::5052:5af:92fa...	ff02::fb	MDNS	183 Standard query 0x0000 PTR _services._dns-sd._udp.local, "QM" question PTR _companion-link._tcp...
17	12.71... 192.168.3.52	224.0.0.251	MDNS	212 Standard query 0x0000 PTR _services._dns-sd._udp.local, "QM" question PTR _companion-link._tcp...
18	14.72... fe80::5052:5af:92fa...	ff02::fb	MDNS	183 Standard query 0x0000 PTR _services._dns-sd._udp.local, "QM" question PTR _companion-link._tcp...
19	14.72... 192.168.3.52	224.0.0.251	MDNS	212 Standard query 0x0000 PTR _services._dns-sd._udp.local, "QM" question PTR _companion-link._tcp...

Offline Exercise

- Use the dns-sd /avahi-publish-service command as example
 - Advertise (register) a web page on local machine

```
dns-sd -R "My Test" _http._tcp . 80  
path=/path-to-page.html
```

Issue: How to Query

- Query needs a back pointer, PTR records
- Exercise: Use the dns-sd / avahi-service-publish command as example
 - Browse web pages on local machines

```
dns-sd -B _http._tcp
avahi-browse -rt _http._tcp
```

Network Service Discovery in Android

- Based on DNS-SD/mDNS
- Foundation for peer-to-peer/Wi-Fi Direct in Android
- See <https://developer.android.com/training/connect-devices-wirelessly/nsd.html> for programming using nsd

Backup Slides

(General Service/Naming Discovery)

General Service/Naming

Discovery Paradigm: Linda

- “Distributed workspace” by David Gelernter in the 80’s at Yale
- Very influential in naming and resource discovery
- Key issues
 - How to name services/resources
 - How to resolve names

The Linda Paradigm

- Naming scheme:
 - arbitrary tuples (heterogeneous-type vectors)

- Name resolution:
 - Nodes write into shared memory
 - Nodes read matching tuples from shared memory
 - exact matching is required for extraction

Linda: Core API

- `out()`: writes tuples to shared space
 - example: `out("abc", 1.5, 12)`.
 - result: insert ("abc", 1.5, 12) into space
- `read()`: retrieves tuple copy matching arg list (blocking)
 - example: `read("abc", ?A, ?B)`
 - result: finds ("abc", 1.5, 12) and sets local variables
 $A = 1.5, B = 12$. Tuple ("abc", 1.5, 12) is still resident in space.
- `in()`: retrieves and deletes matching tuple from space (blocking)
 - example: same as above except ("abc", 1.5, 12) is deleted
- `eval(expression)`: similar to `out` except that the tuple argument to `eval` is evaluated
 - example: `eval("ab",-6,abs(-6))` creates tuple ("ab", -6, 6)

Linda Extension: JavaSpaces

- Industry took Linda principles and made modifications
 - add transactions, leases, events
 - store Java objects instead of tuples
 - a very comprehensive service discovery system

- Definitive book, “JavaSpaces Principles, Patterns, and Practice”