

---

Network Applications:  
HTTP/1.0/1.1/2;  
High-Performance Server Design (Per Thread)

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

10/19/2021

# Outline

---

- Admin and recap
- HTTP/1.0
  - Basic design
  - Dynamic content
- HTTP “acceleration”
- Network server design

# Admin

---

- HTTP server assignment (part 1) to be posted later today
  
- Plagiarism is NOT allowed in any assignment
  - Discussion and reading are encouraged, but copying is NOT!
  - If you reuse others' (e.g., classmates, friends, Stackoverflow, CSDN) code as part of your solution, please mark it in comments.

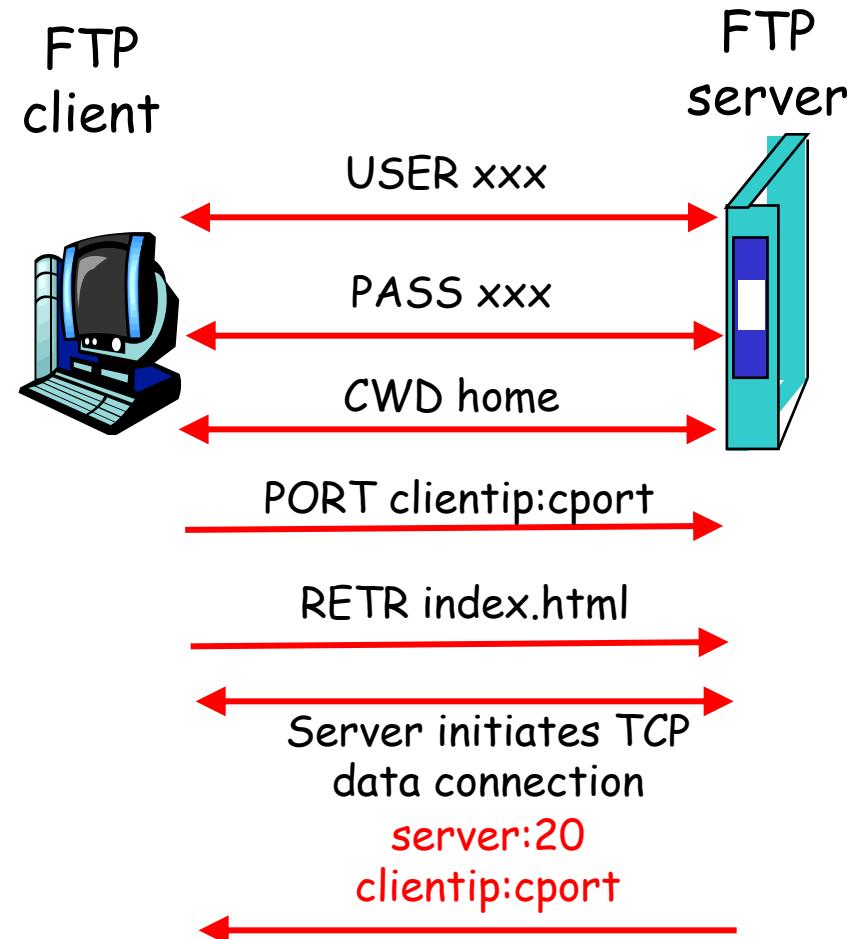
# Recap: FTP

## □ A stateful protocol

- state established by commands such as
  - USER/PASS, CWD, TYPE

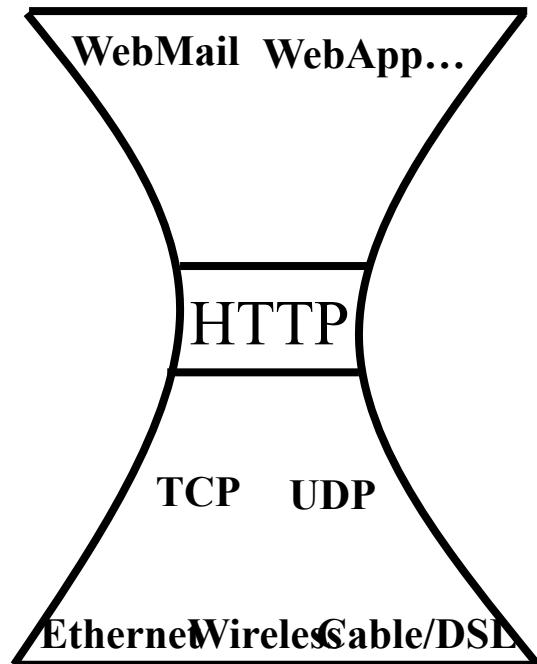
## □ Multiple TCP connections

- A control connection
- Data connections
  - Two approaches: PORT vs PASV
  - GridFTP: concurrent data connections; block data transfer mode



# Recap: HTTP

- Wide use of HTTP for Web applications
- Example: RESTful API
  - RESTful design
    - [http://www.ics.uci.edu/~fielding/pubs/dissertation/rest\\_arch\\_style.htm](http://www.ics.uci.edu/~fielding/pubs/dissertation/rest_arch_style.htm)
    - <http://docs.oracle.com/javaee/6/tutorial/doc/giepu.html>



# Recap: HTTP

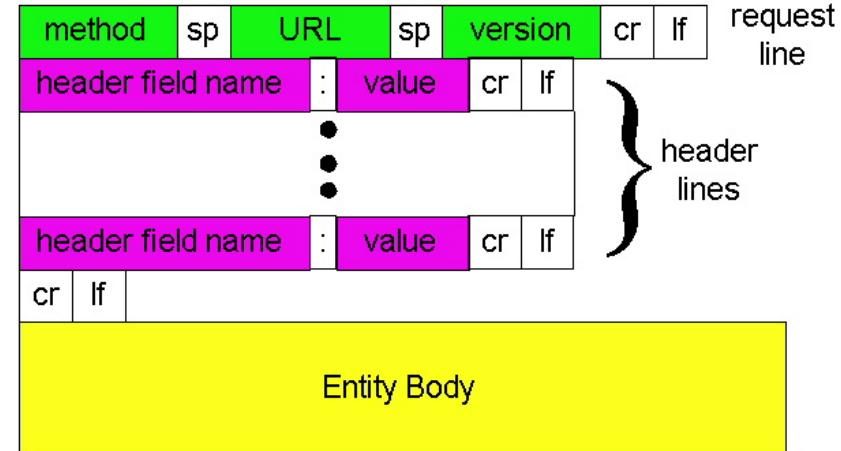
## ❑ C-S app serving Web pages

### ○ message format

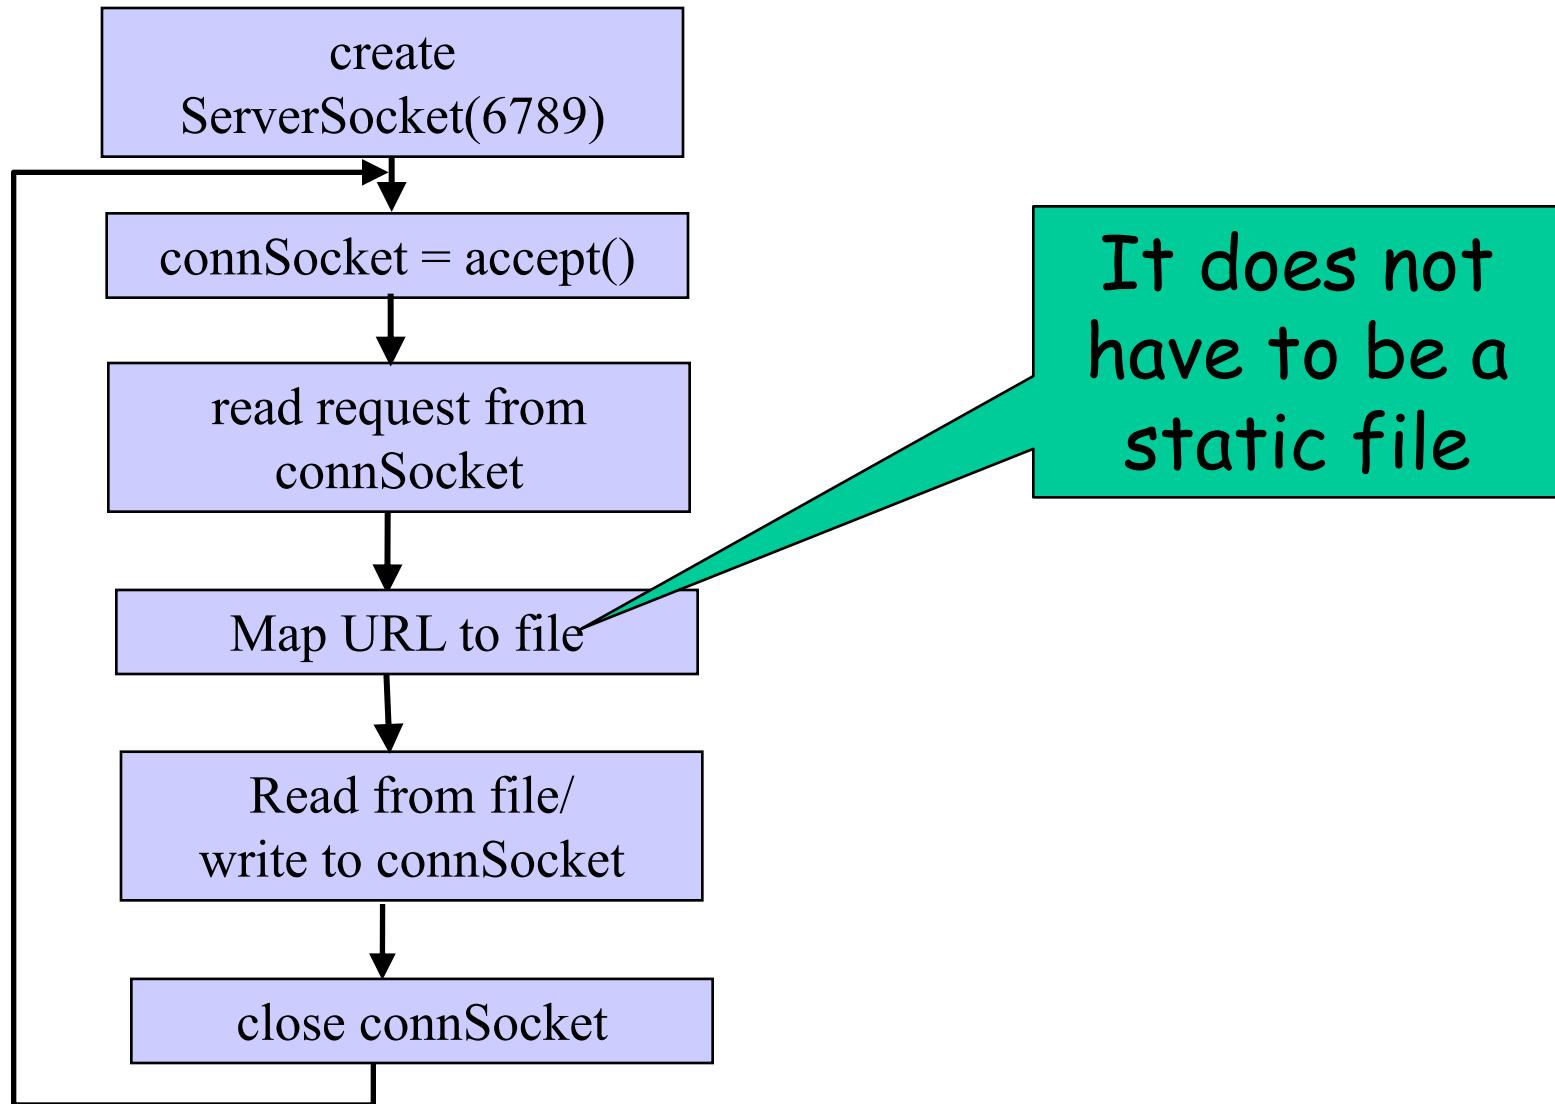
- request/response line, header lines, entity body
- simple methods, rich headers

### ○ message flow

- stateless server, thus states such as cookie and authentication are needed in each message



# Recap: Basic HTTP/1.0 Server



# Outline

---

- Admin and recap
- HTTP/1.0
  - Basic design
  - *Dynamic content*

# Dynamic Content Pages

---

- There are multiple approaches to make dynamic web pages:
  - Embed code into pages (**server side include**)
    - http server includes an interpreter for the type of pages
  - Invoke external programs (http server is agnostic to the external program execution)
    - E.g., **Common Gateway Interface (CGI)**

`http://www.cs.yale.edu/index.shtml`

`http://www.cs.yale.edu/cgi-bin/ureserve.pl`

`http://www.google.com/search?q=Yale&sourceid=chrome`

## Example SSI

---

- See programming/examples-java-socket/BasicWebServer/ssi/index.shtml, header.shtml, ...

## Example SSI

---

- See programming/examples-java-socket/BasicWebServer/ssi/index.shtml, header.shtml, ...
  
- To enable ssi, need configuration to tell the web server (see conf/apache-htaccess)
  - <https://httpd.apache.org/docs/2.2/howto/htaccess.html> (Server Side Includes example)

# CGI: Invoking External Programs

---

## □ Two issues

- Input: Pass HTTP request parameters to the external program
- Output: Redirect external program output to socket

## Example: Typical CGI Implementation

---

- Starts the executable as a child process
  - Passes HTTP request as environment variables
    - <http://httpd.apache.org/docs/2.2/env.html>
    - CGI standard: <http://www.ietf.org/rfc/rfc3875>
  - Redirects input/output of the child process to the socket

# Example: CGI

## □ Example:

- GET /search?q=Yale&sourceid=chrome HTTP/1.0
- **setup environment variables, in particular**  
**`$QUERY_STRING=q=Yale&sourceid=chrome`**
- **start search and redirect its input/output**

<https://docs.oracle.com/javase/7/docs/api/java/lang/ProcessBuilder.html>

# Example

- <http://172.28.229.215/BasicWebServer/cgi/price.cgi?appl>

```
#!/usr/bin/perl -w

$company = $ENV{'QUERY_STRING'};
print "Content-Type: text/html\r\n";
print "\r\n";

print "<html>";
print "<h1>Hello! The price is ";

if ($company =~ /appl/) {
    my $var_rand = rand();
    print 450 + 10 * $var_rand;
} else {
    print "150";
}

print "</h1>";
print "</html>";
```

## Client Using Dynamic Pages

---

- See ajax.html and wireshark for client code example

<http://172.28.229.215/BasicWebServer/cgi/ajax.html>

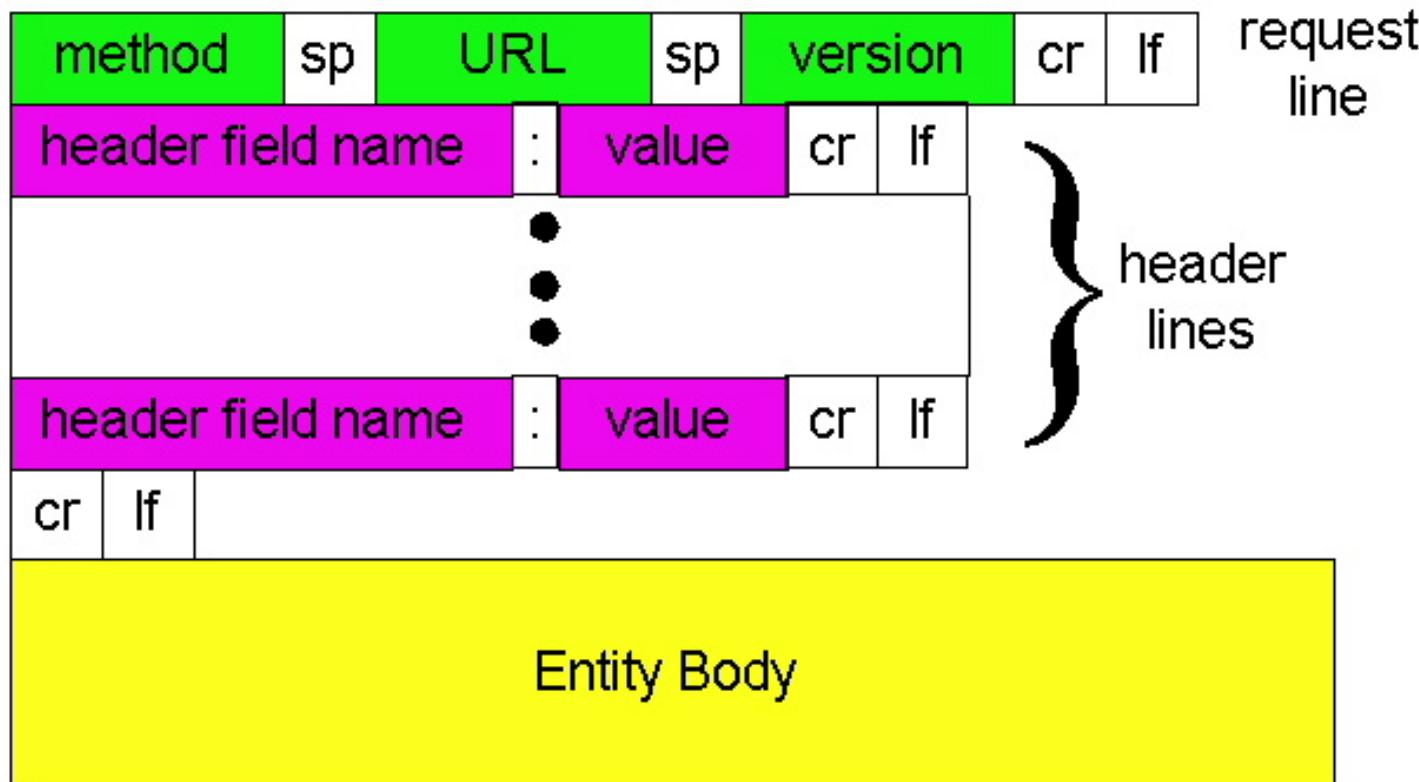
## Discussions

---

- What features are missing in HTTP that we have covered so far?

# HTTP: POST

- If an HTML page contains forms or parameter too large, they are sent using POST and encoded in message body



# HTTP: POST Example

---

POST /path/script.cgi HTTP/1.0

User-Agent: MyAgent

Content-Type: application/x-www-form-urlencoded

Content-Length: 15

item1=A&item2=B

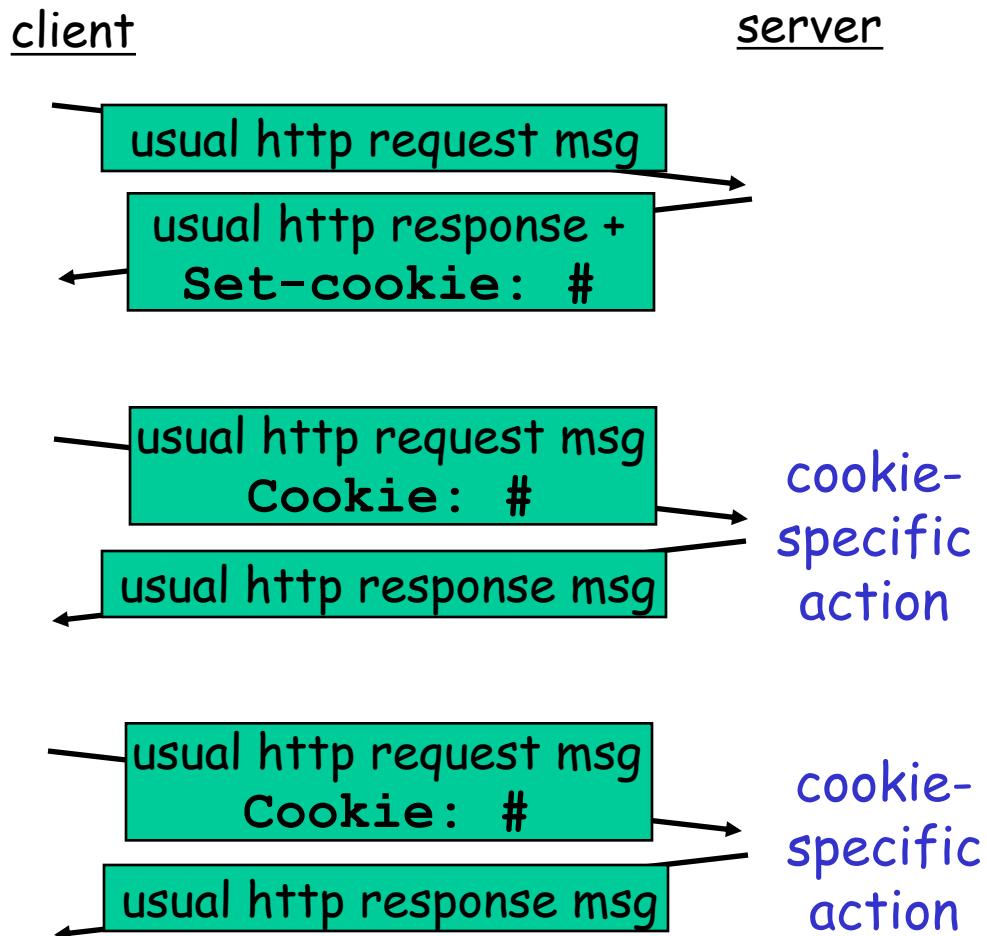
Example using nc:

<https://github.com/programming-examples/java-socket/tree/BasicWebServer/nc/>

# Stateful User-server Interaction: Cookies

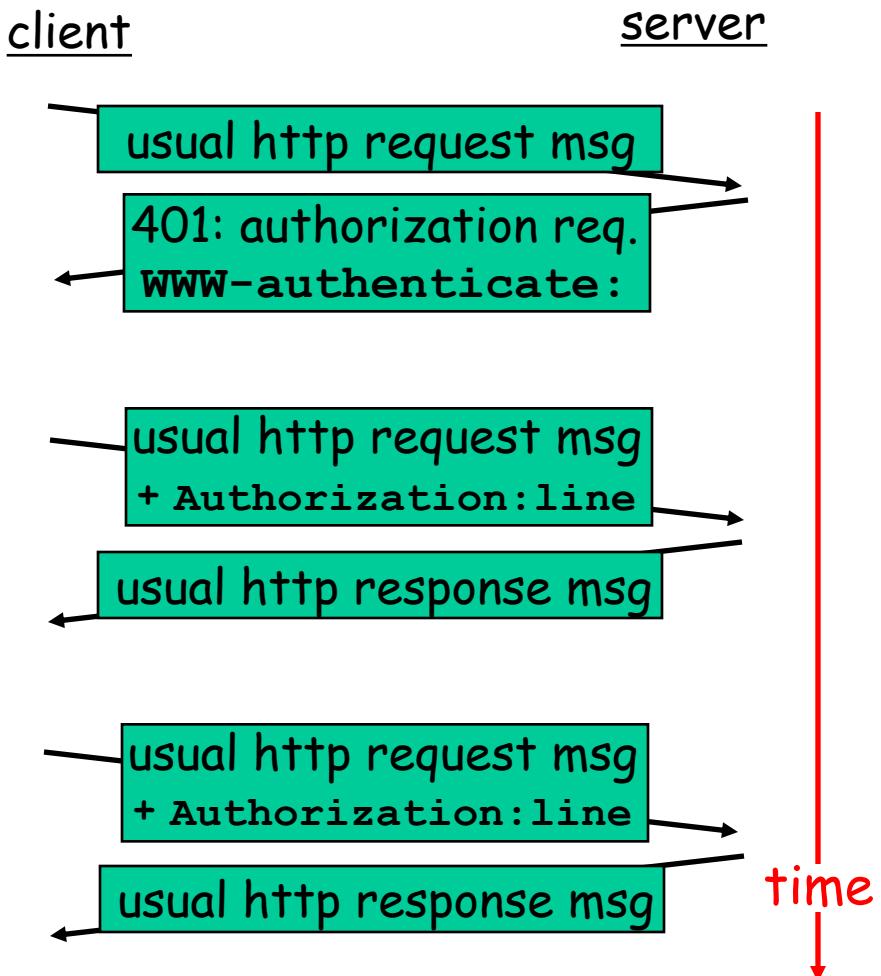
Goal: no explicit application level session

- ❑ Server sends “cookie” to client in response msg
  - **Set-cookie:** 1678453
- ❑ Client presents cookie in later requests
  - **Cookie:** 1678453
- ❑ Server matches presented-cookie with server-stored info
  - authentication
  - remembering user preferences, previous choices



# Authentication of Client Request

- Authentication goal:** control access to server documents
- **stateless:** client must present authorization in each request
  - **authorization:** typically name, password
    - Authorization: header line in request
    - if no authorization presented, server refuses access, sends **www-authenticate:** header line in response



Browser caches name & password so that user does not have to repeatedly enter it.

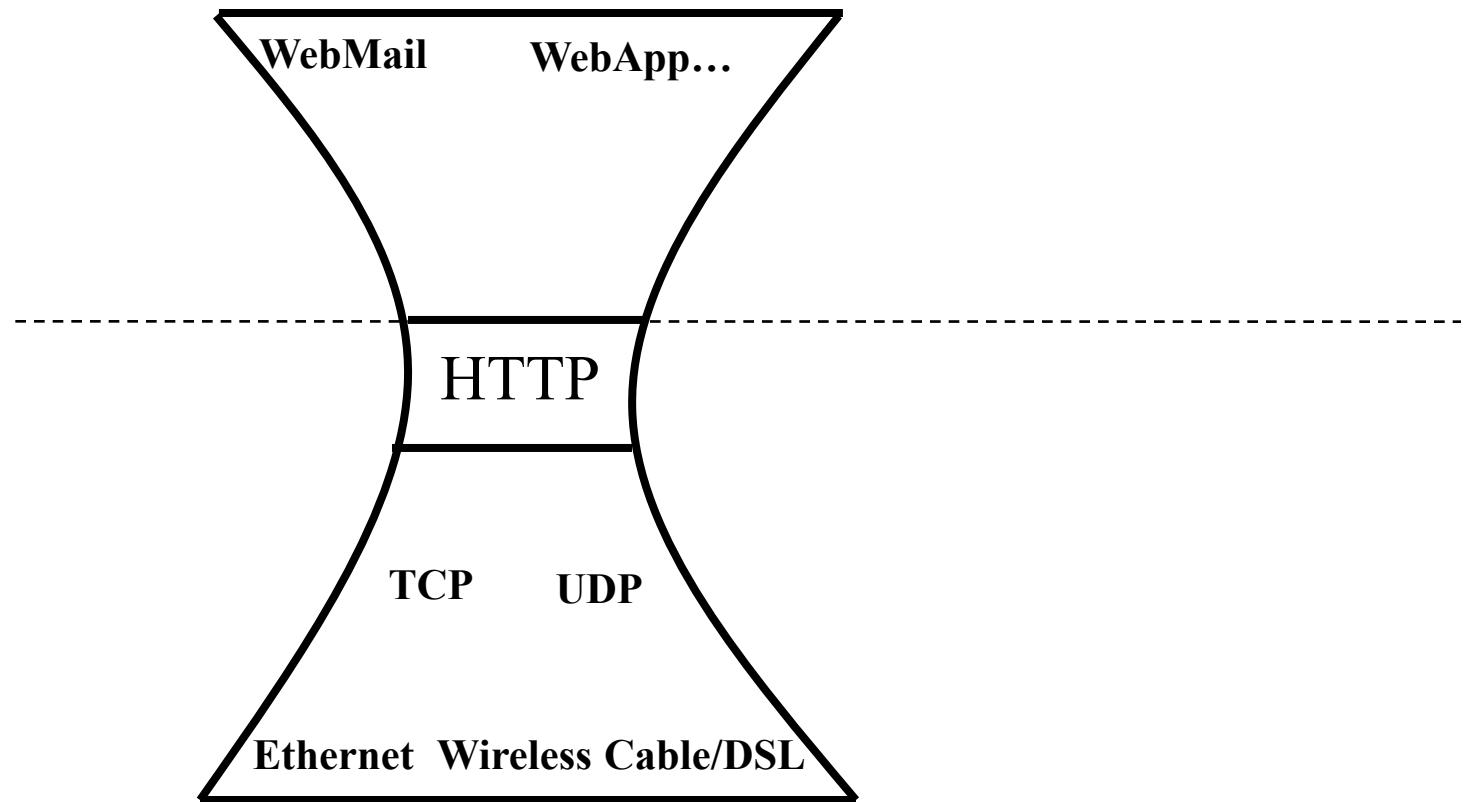
## Example: Amazon S3

---

### □ Amazon S3 API

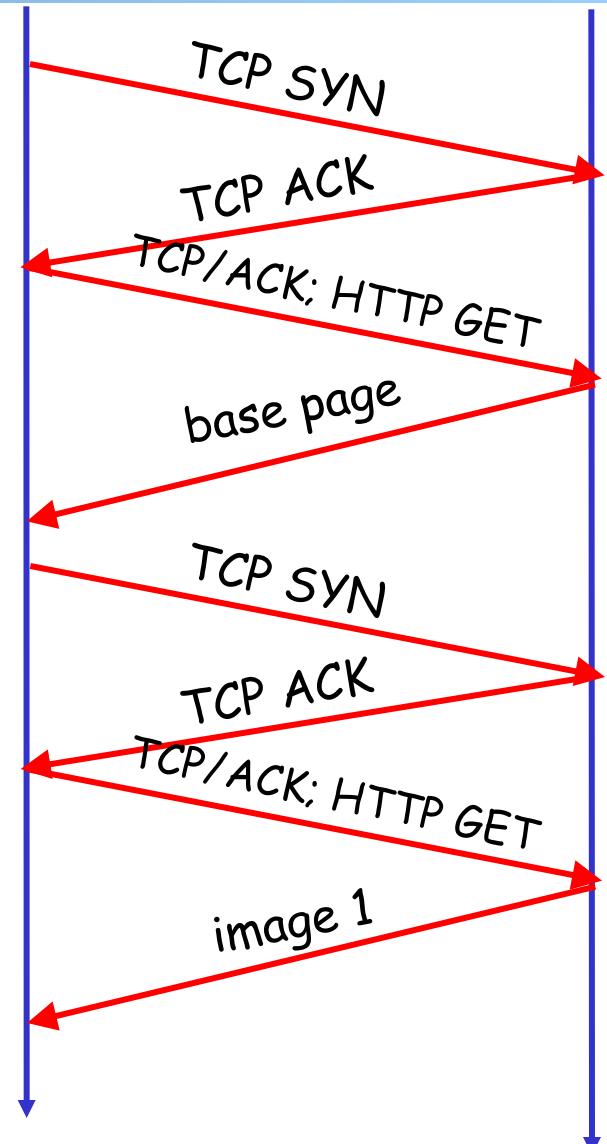
- <http://docs.aws.amazon.com/AmazonS3/latest/API/APIRest.html>

# HTTP as the Thin Waist



# Protocol Flow of Basic HTTP/1.0

- $\geq 2$  RTTs per object:
  - TCP handshake --- 1 RTT
  - client request and server responds --- at least 1 RTT  
(if object can be contained in one packet)



# Outline

---

- Admin and recap
- HTTP/1.0
- *HTTP "acceleration"*

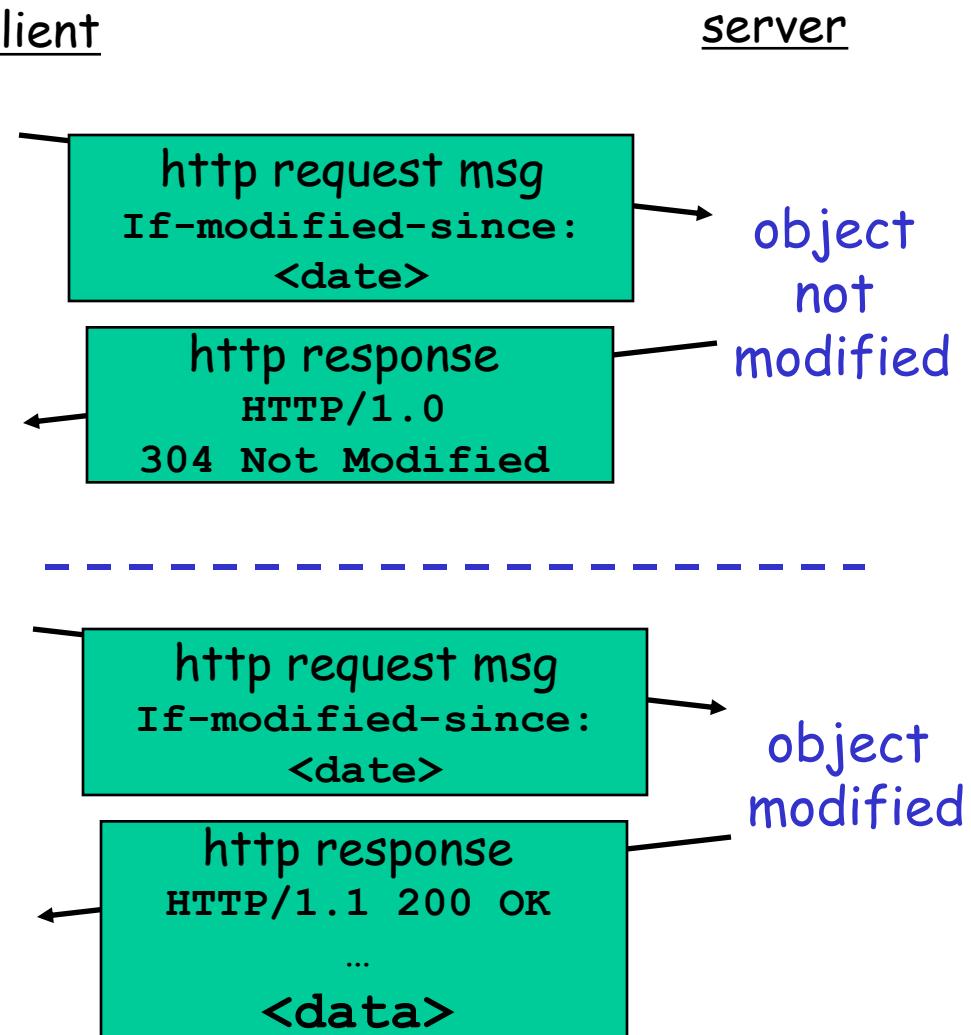
# Substantial Efforts to Speedup Basic HTTP/1.0

- Reduce the number of objects fetched [Browser cache]
- Reduce data volume [Compression of data]
- Header compression [HTTP/2]
- Reduce the latency to the server to fetch the content [Proxy cache]
- Remove the extra RTTs to fetch an object [Persistent HTTP, aka HTTP/1.1]
- Increase concurrency [Multiple TCP connections]
- Asynchronous fetch (multiple streams) using a single TCP [HTTP/2]
- Server push [HTTP/2]



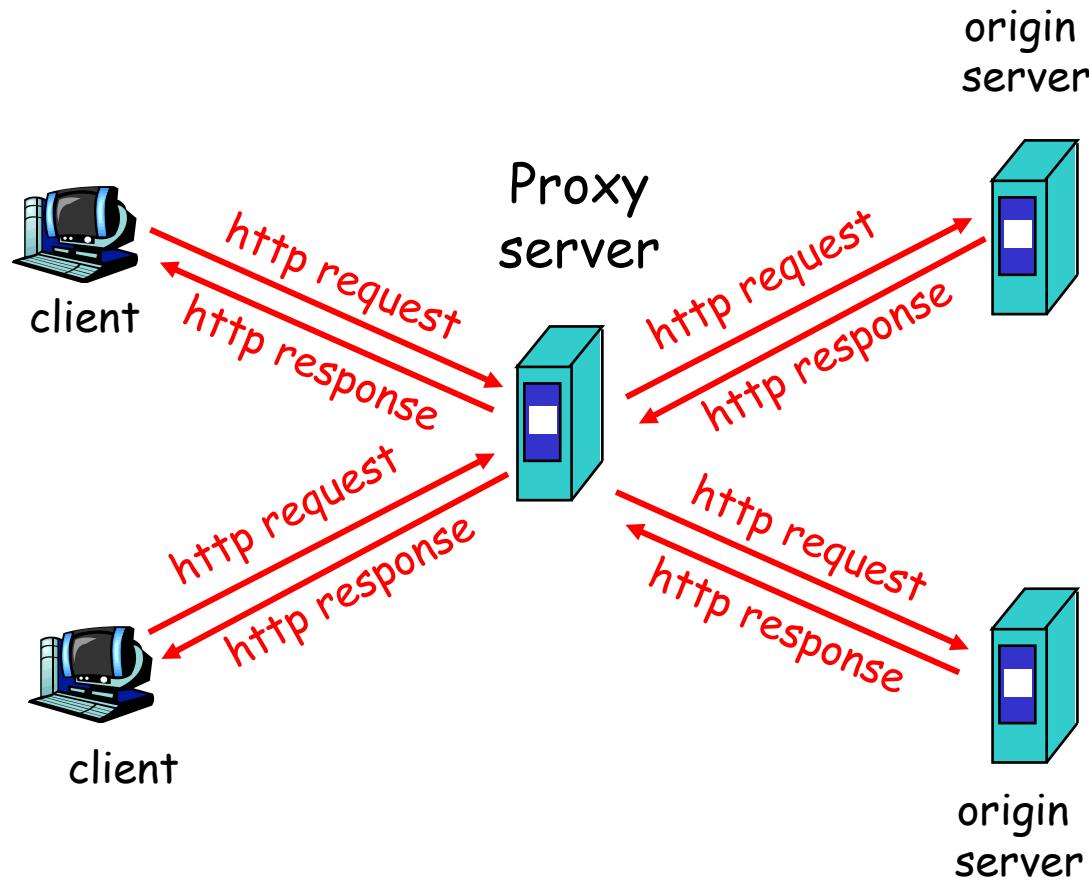
# Browser Cache and Conditional GET

- **Goal:** don't send object if client has up-to-date stored (cached) version
- client: specify date of cached copy in http request  
**If-modified-since:**  
    <date>
- server: response contains no object if cached copy up-to-date:  
**HTTP/1.0 304 Not Modified**



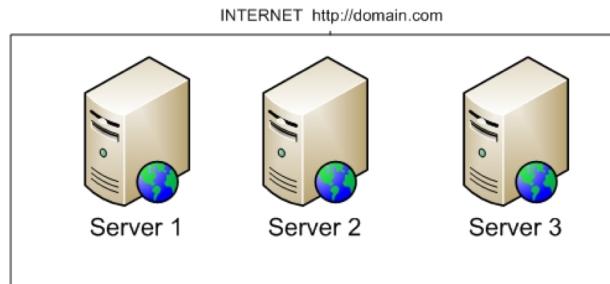
# Web Caches (Proxy)

**Goal:** satisfy client request without involving origin server

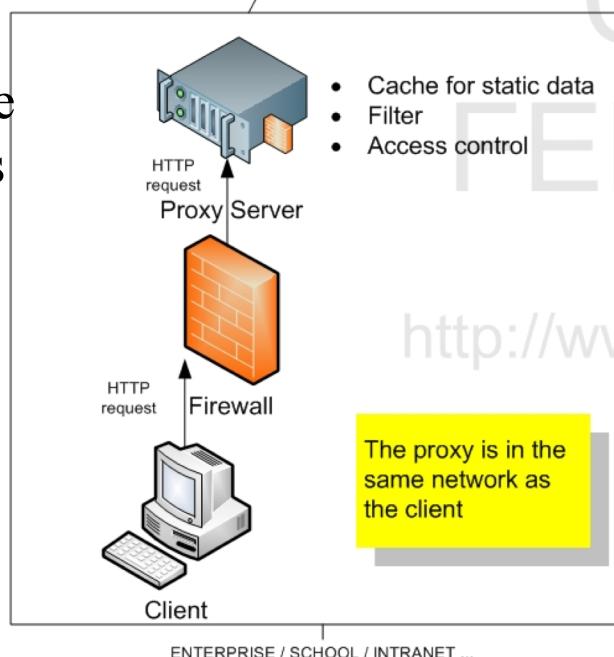


## Two Types of Proxies

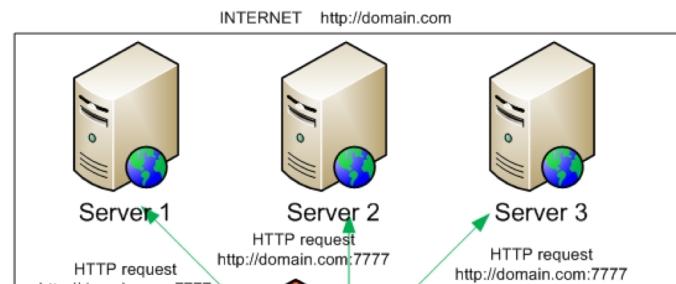
(FORWARD) PROXY server



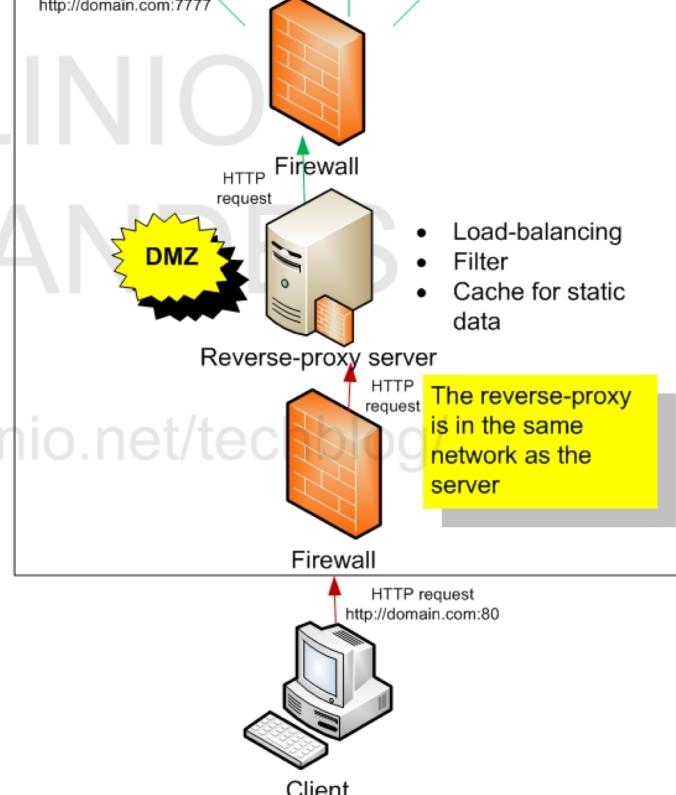
Typically  
in the same  
network as  
the client



REVERSE-PROXY server



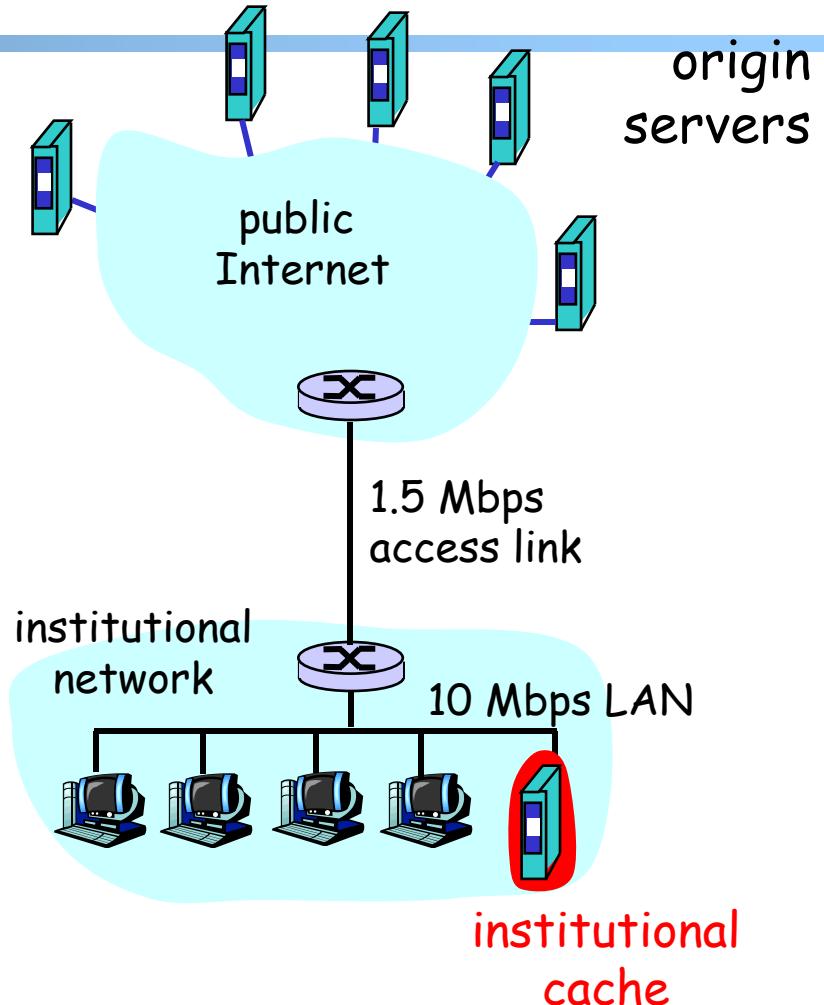
Typically in  
the same  
network as  
the server



# Benefits of Forward Proxy

**Assume:** cache is “close” to client (e.g., in same network)

- smaller response time: cache “closer” to client
- decrease traffic to distant servers
  - o link out of institutional/local ISP network often is bottleneck



# No Free Lunch: Problems of Web Caching

---

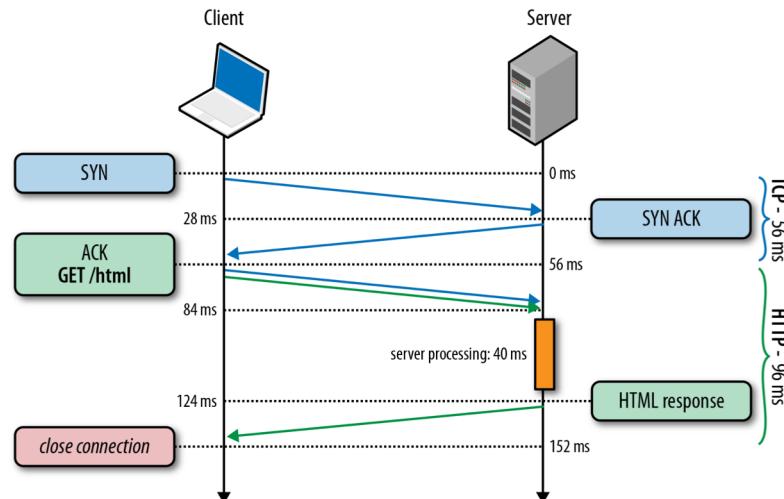
- The major issue of web caching is how to maintain consistency
- Two ways
  - pull
    - Web caches periodically pull the web server to see if a document is modified
  - push
    - whenever a server gives a copy of a web page to a web cache, they sign a lease with an expiration time; if the web page is modified before the lease, the server notifies the cache

## HTTP/1.1: Persistent (keepalive/pipelining) HTTP

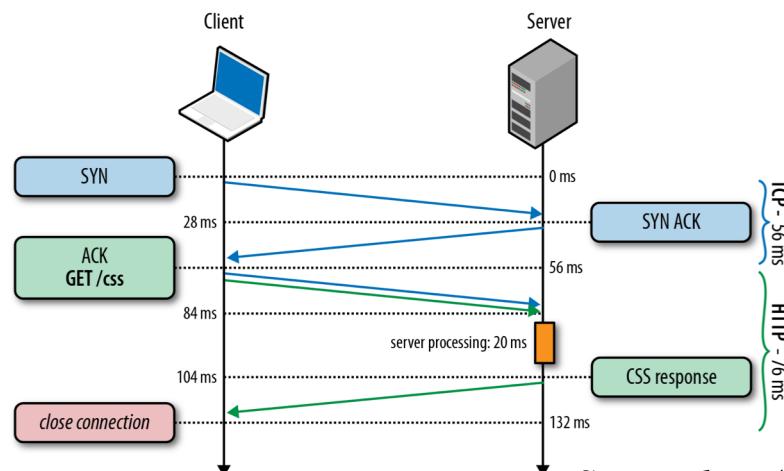
- HTTP/1.0 allows a single request outstanding, while HTTP/1.1 allows request pipelining
  - On same TCP connection: server parses request, responds, parses new request, ...
  - Client sends requests for all referenced objects as soon as it receives base HTML
- Benefit
  - Fewer RTTs
  - See Joshua Graessley WWDC 2012 talk: 3x within iTunes

# HTTP/1.0, Keep-Alive, Pipelining

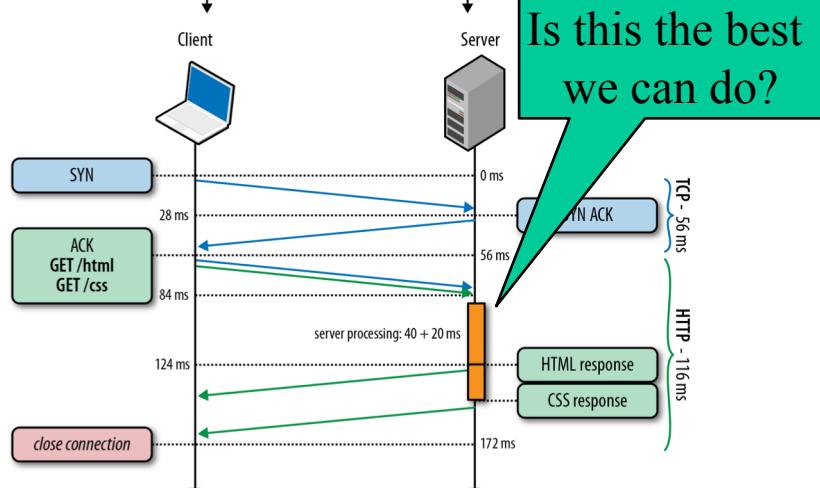
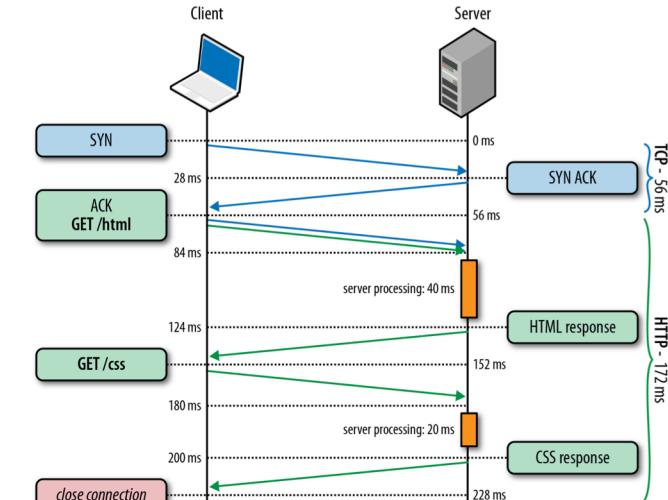
TCP connection #1, Request #1: HTML request



TCP connection #2, Request #2: CSS request

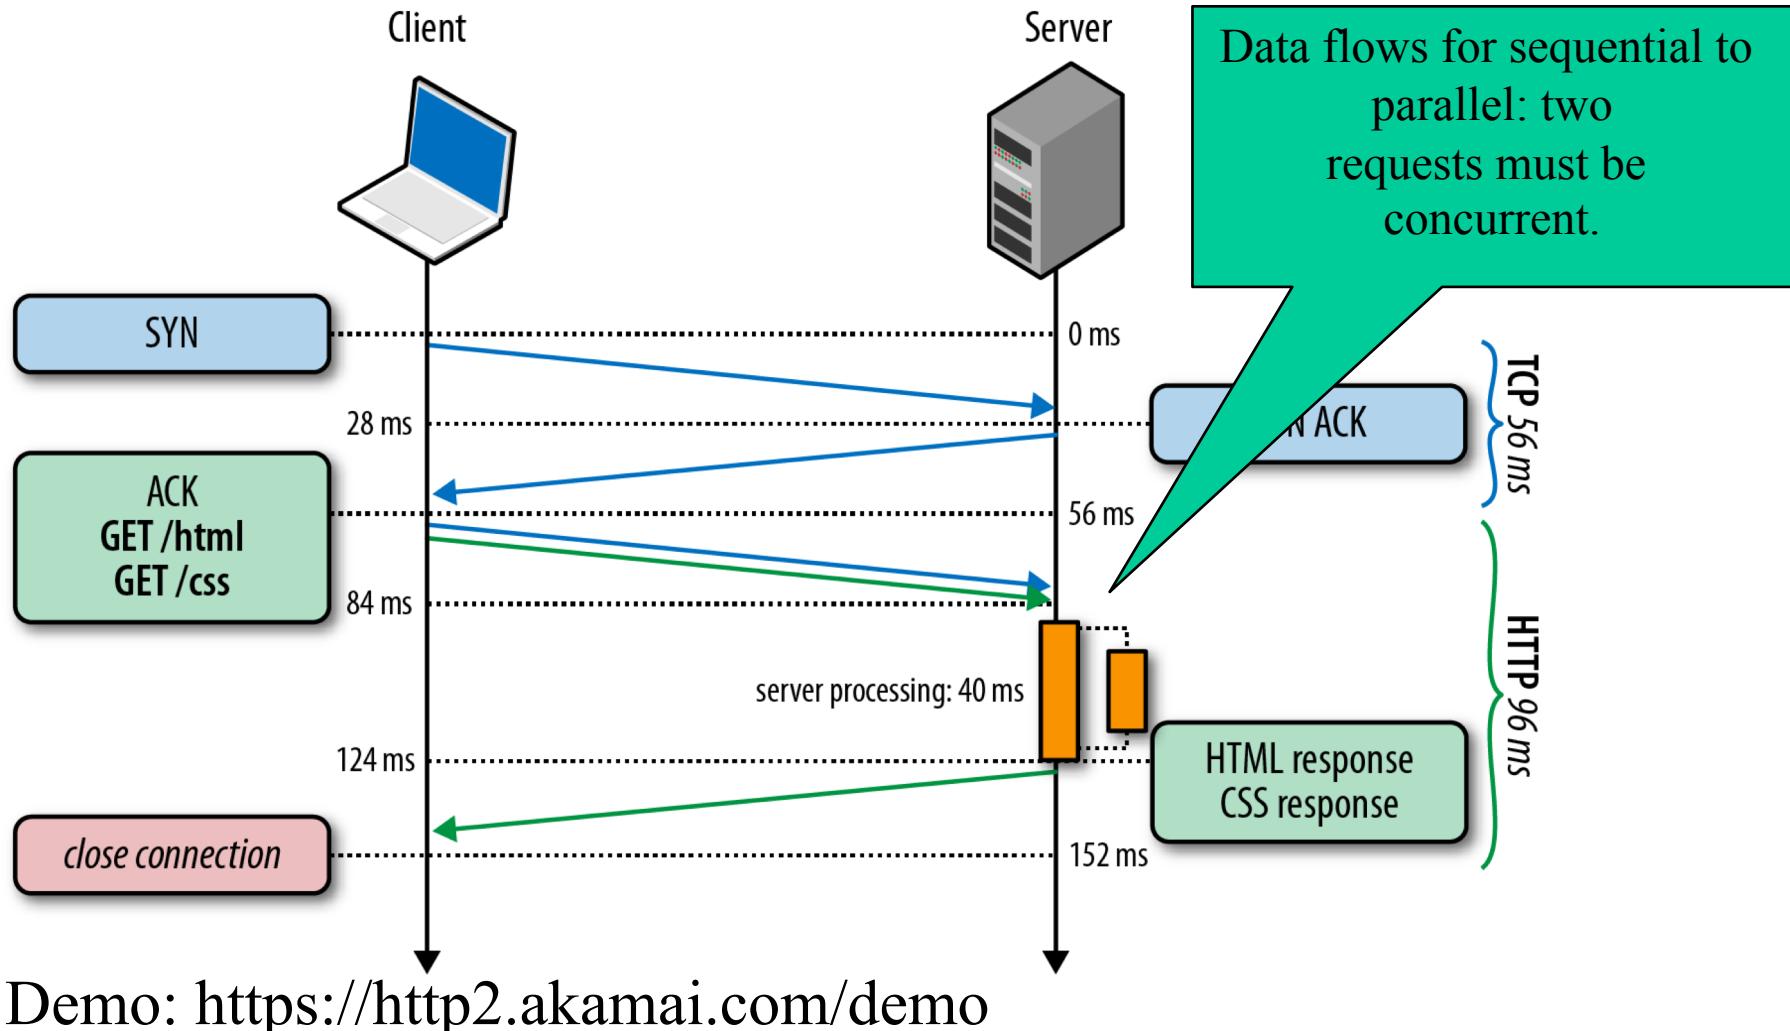


TCP connection #1, Request #1-2: HTML + CSS



Source: <http://chimera.labs.oreilly.com/books/123000000545/ch11.html>

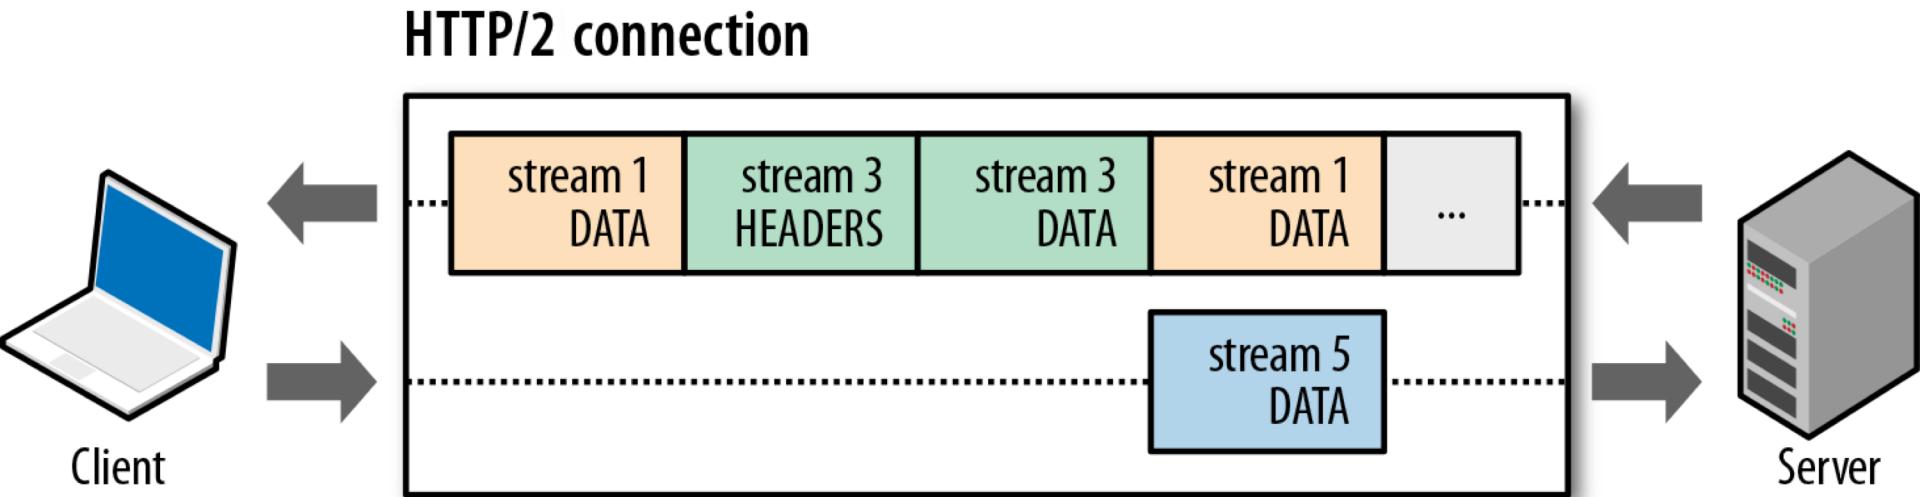
# HTTP/2 Basic Idea: Remove Head-of-Line Blocking in HTTP/1.1



# Observing HTTP/2

- `export SSLKEYLOGFILE=/tmp/keylog.txt`
- Start Chrome, e.g.,
  - Mac: /Applications/Google Chrome.app/Contents/MacOS/Google Chrome
  - Ubuntu: firefox
- Visit HTTP/2 pages, such as  
<https://http2.akamai.com>
- Wireshark:
  - Mac: Wireshark -> preferences -> protocol -> TSL (pre)-master-secret log file name
  - Ubuntu: edit -> preferences -> protocol -> SSL (pre)-master-secret log file name

# HTTP/2 Design: Multi-Streams



Bit	+0..7	+8..15	+16..23	+24..31
0		Length		Type
32	Flags			
40	R		Stream Identifier	
...			<i>Frame Payload</i>	

**HTTP/2 Binary Framing**

# HTTP/2 Header Compression

Request headers

:method	GET
:scheme	https
:host	example.com
:path	/resource
user-agent	Mozilla/5.0 ...
custom-hdr	some-value

Static table

1	:authority	
2	:method	GET
...	...	...
51	referer	
...	...	...
62	user-agent	Mozilla/5.0 ...
63	:host	example.com
...	...	...



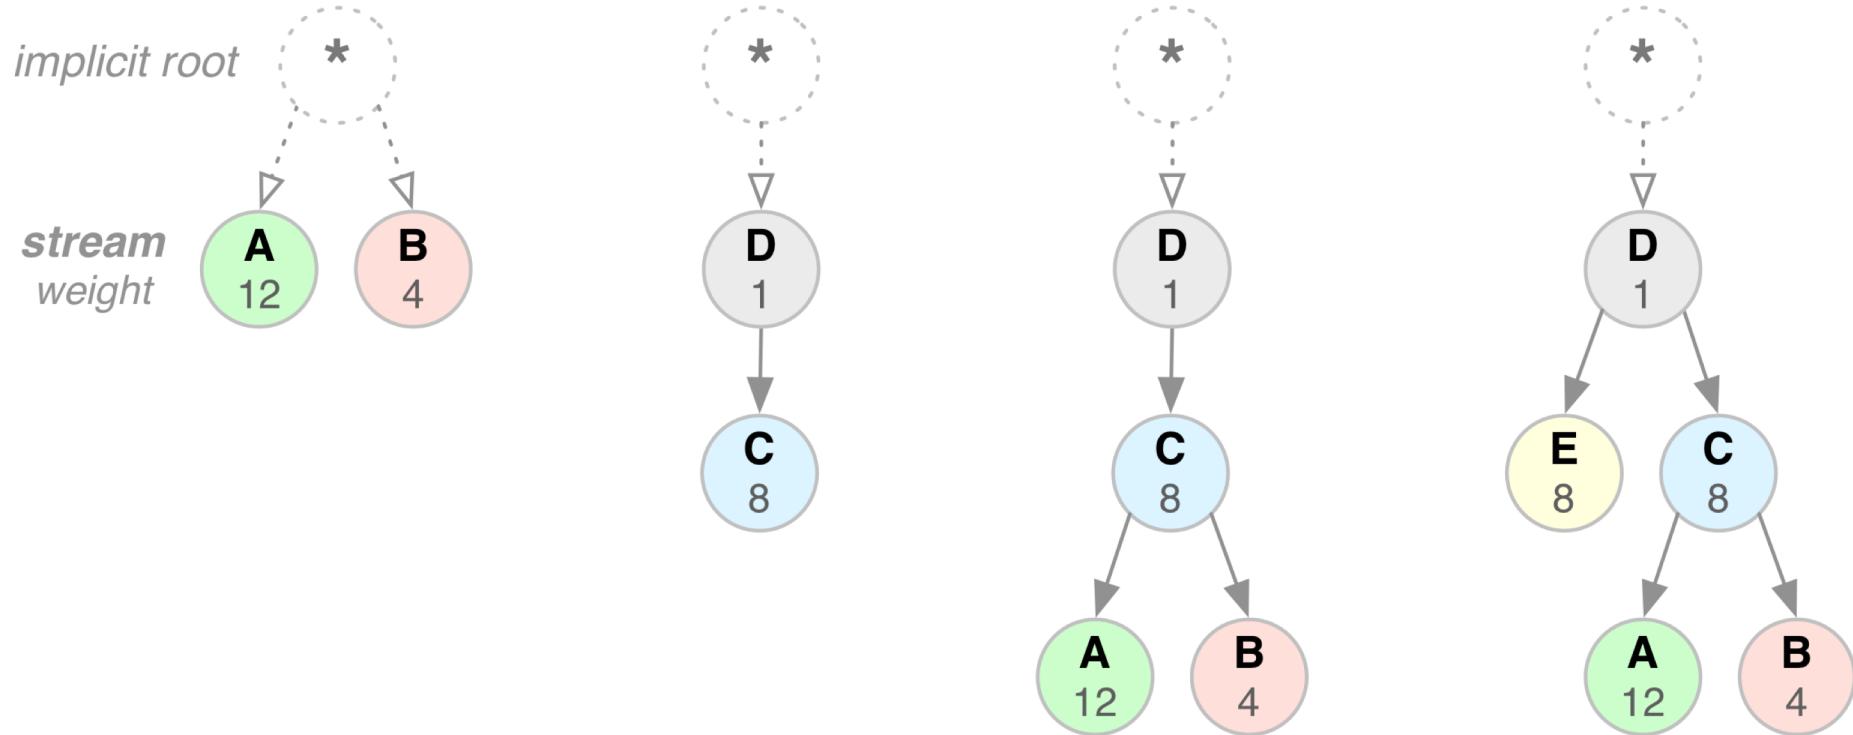
Encoded headers

2
7
63
19
Huffman("/resource")
62
Huffman("custom-hdr")
Huffman("some-value")

Dynamic table

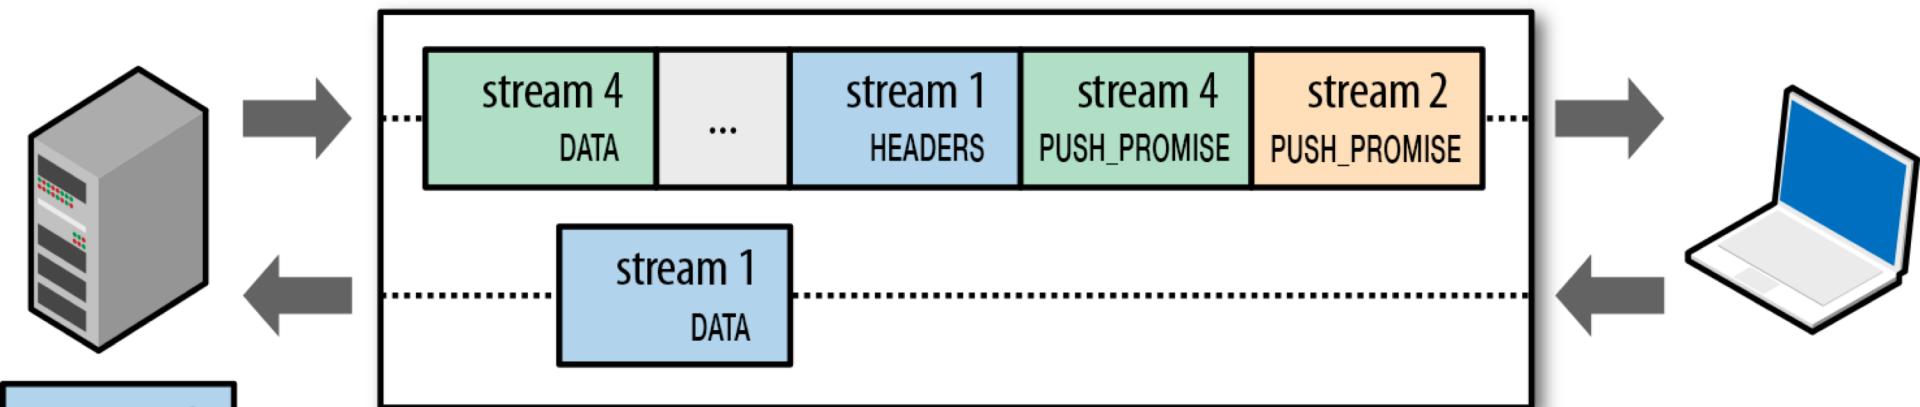


# HTTP/2 Stream Dependency and Weights



# HTTP/2 Server Push

## HTTP/2 connection



**stream 1: /page.html** (client request)

**stream 2: /script.js** (push promise)

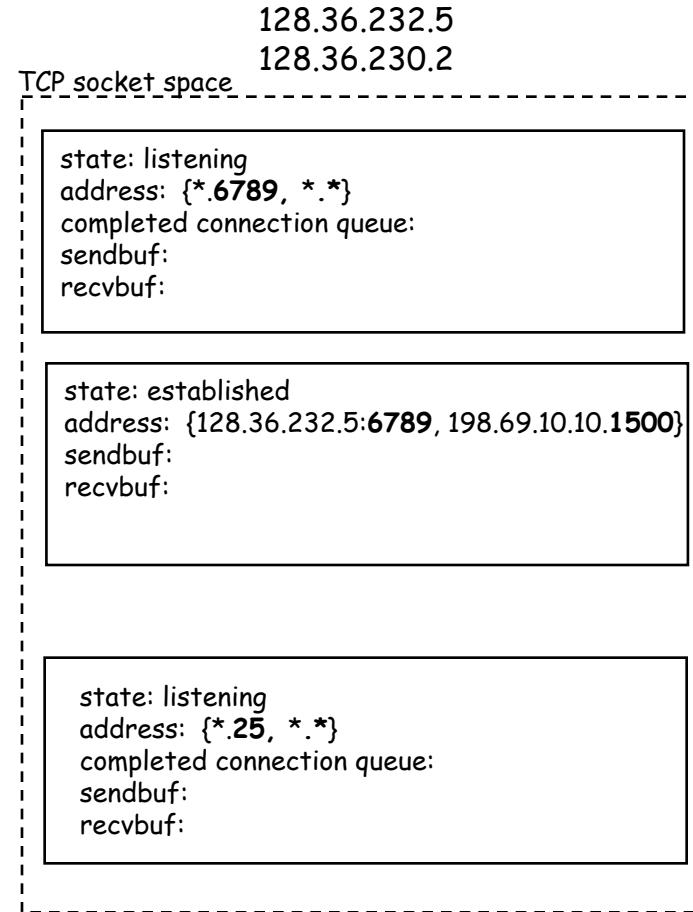
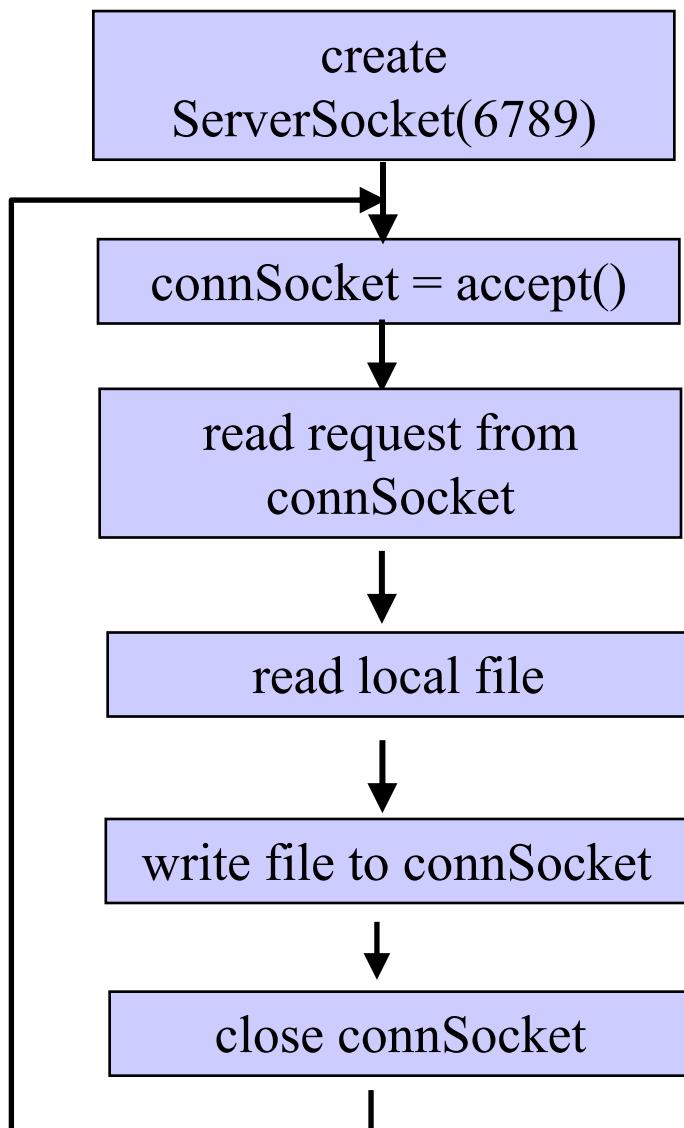
**stream 4: /style.css** (push promise)

# Outline

---

- Admin and recap
- HTTP/1.0
- HTTP "acceleration"
- *Network server design*

# WebServer Implementation



Discussion: what does each step do and how long does it take?

# Demo

---

- Try TCPServer
- Start two TCPClient
  - Client 1 starts early but stops
  - Client 2 starts later but inputs first