

# Network Layer:

## Distance Vector Protocols Variations

## Link-State Protocol

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf21/index.shtml>

12/02/2021

# Outline

---

- Admin and recap
- Network overview
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
      - Link state protocols (distributed state synchronization)

# Admin

---

- Lab assignment six posted
  - Due Jan. 4, 2022

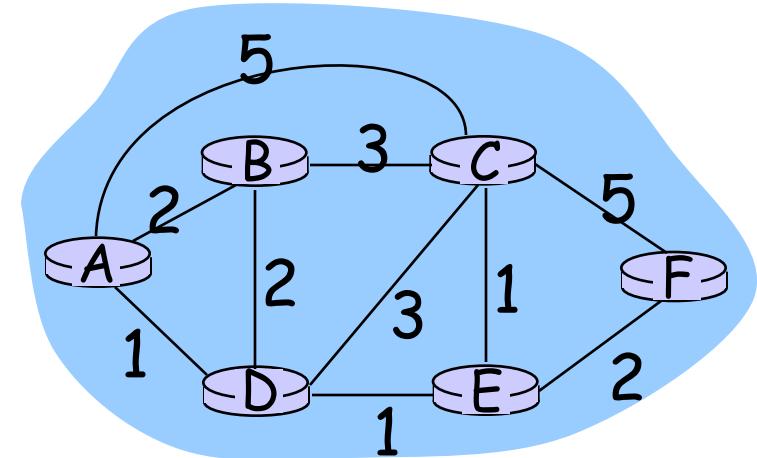
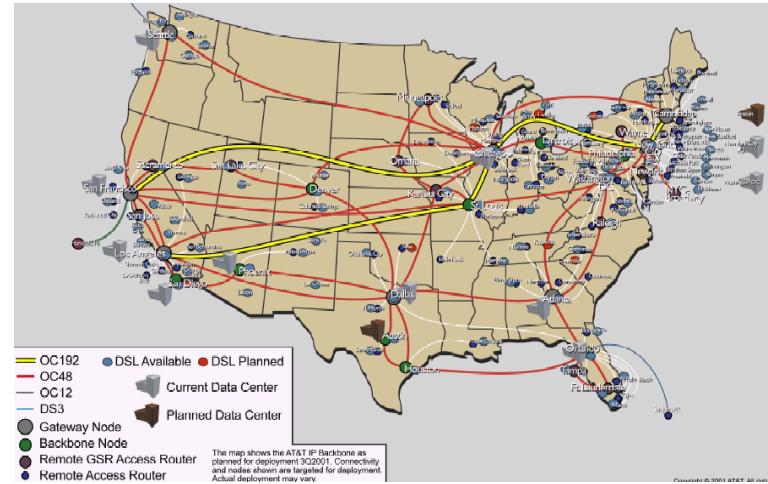
# Recap: Routing Context

## Routing

**Goal:** determine “good” paths (sequences of routers) thru networks from source to dest.

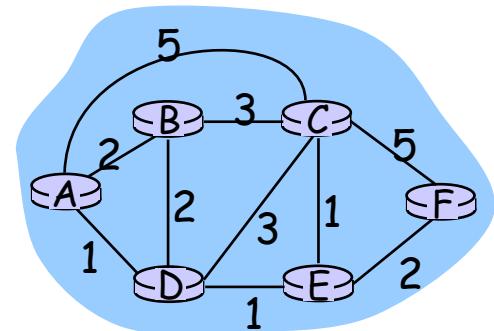
Often depends on a graph abstraction:

- graph nodes are routers
- graph edges are physical links
  - links have properties: delay, capacity, \$ cost, **policy**



# Recap: Routing Design Space

- Routing has a large design space
  - who decides routing?
    - source routing: end hosts make decision
    - network routing: networks make decision
  - how many paths from source s to destination d?
    - multi-path routing
    - single path routing
  - what does routing compute?
    - network cost minimization (shortest path routing)
    - QoS aware
  - will routing adapt to network traffic demand?
    - adaptive routing
    - static routing
  - ...



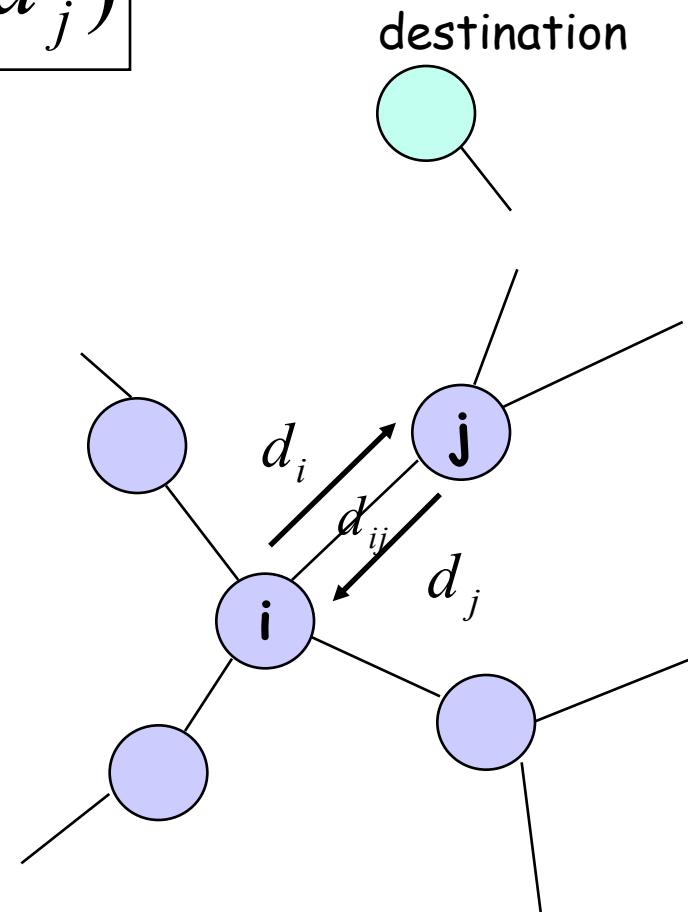
# Recap: Distance Vector Routing: Basic Idea (Bellman-Ford Alg)

- r At node  $i$ , the basic update rule

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

where

- $d_i$  denotes the distance estimation from  $i$  to the destination,
- $N(i)$  is set of neighbors of node  $i$ , and
- $d_{ij}$  is the distance of the direct link from  $i$  to  $j$

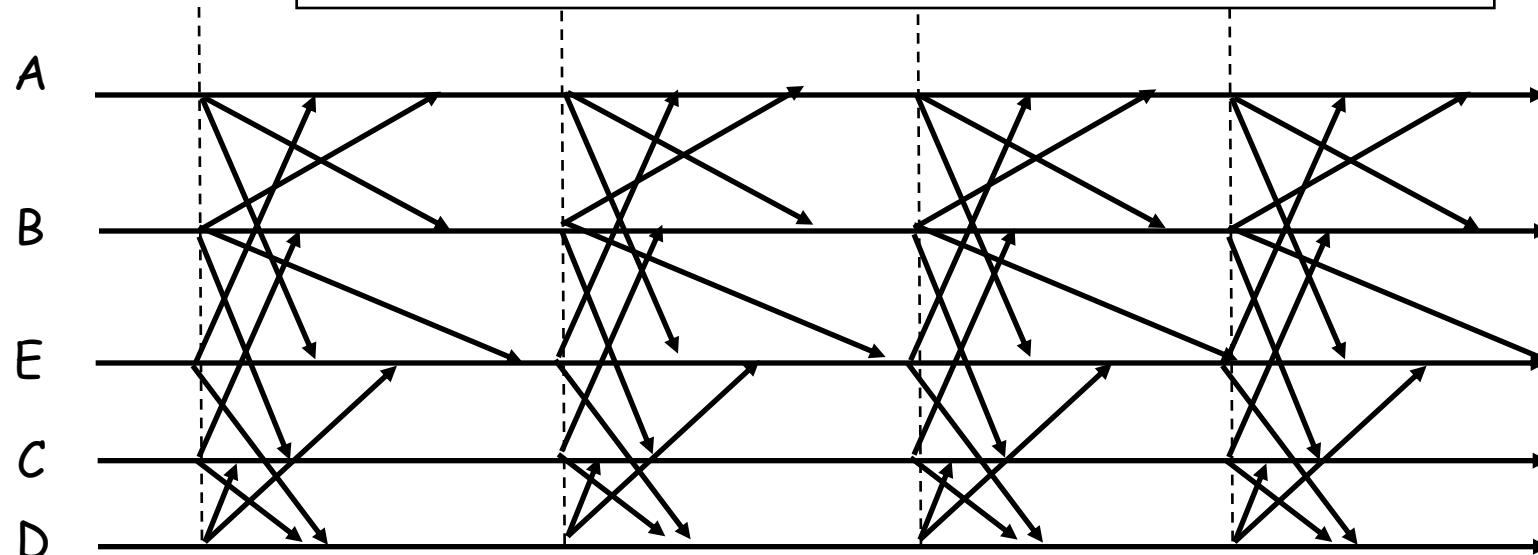


# Recap: Synchronous Bellman-Ford (SBF)

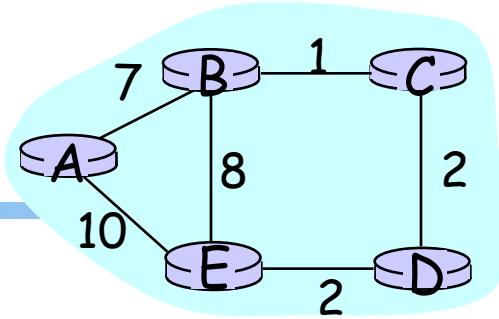
## □ Nodes update in rounds:

- there is a global clock;
- at the beginning of each round, each node sends its estimate to all of its neighbors;
- at the end of the round, updates its estimation

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$



## Recap: SBF/∞



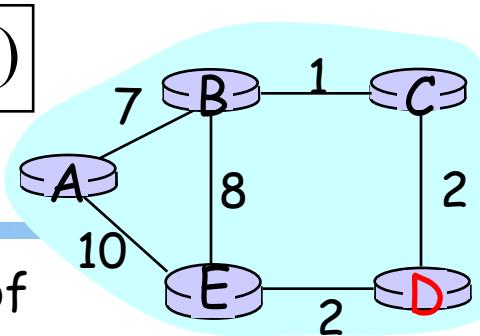
- Initialization (time 0):

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ \infty & \text{otherwise} \end{cases}$$

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

## Example

Consider D as destination;  $d(t)$  is a vector consisting of estimation of each node at round  $t$



|        | A        | B        | C        | E        | D |
|--------|----------|----------|----------|----------|---|
| $d(0)$ | $\infty$ | $\infty$ | $\infty$ | $\infty$ | 0 |
| $d(1)$ | $\infty$ | $\infty$ | 2        | 2        | 0 |
| $d(2)$ | 12       | 3        | 2        | 2        | 0 |
| $d(3)$ | 10       | 3        | 2        | 2        | 0 |
| $d(4)$ | 10       | 3        | 2        | 2        | 0 |

Observation:  $d(0) \geq d(1) \geq d(2) \geq d(3) \geq d(4) = d^*$

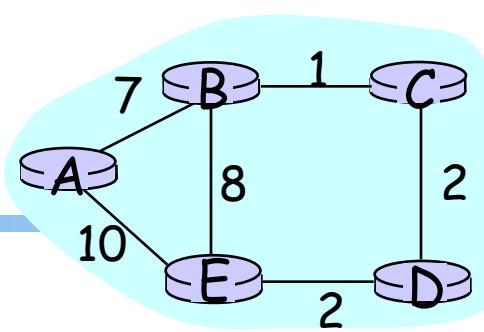
$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

## A Nice Property of SBF: Monotonicity

- Consider two configurations  $d(t)$  and  $d'(t)$
- If  $d(t) \geq d'(t)$ 
  - i.e., each node has a higher estimate in one scenario ( $d$ ) than in another scenario ( $d'$ ),
- then  $d(t+1) \geq d'(t+1)$ 
  - i.e., each node has a higher estimate in  $d$  than in  $d'$  after one round of synchronous update.

Recap: SBF at another

Initial Configuration: SBF/-1



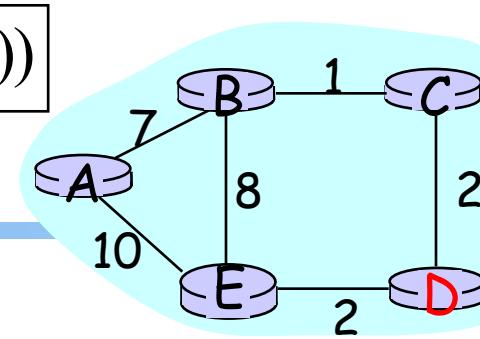
- Initialization (time 0):

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ -1 & \text{otherwise} \end{cases}$$

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

## Example

Consider D as destination



|      | A  | B  | C  | E  | D |
|------|----|----|----|----|---|
| d(0) | -1 | -1 | -1 | -1 | 0 |
| d(1) | 6  | 0  | 0  | 2  | 0 |
| d(2) | 7  | 1  | 1  | 2  | 0 |
| d(3) | 8  | 2  | 2  | 2  | 0 |
| d(4) | 9  | 3  | 3  | 2  | 0 |
| d(5) | 10 | 3  | 3  | 2  | 0 |
| d(6) | 10 | 3  | 3  | 2  | 0 |

Observation:  $d(0) \leq d(1) \leq d(2) \leq d(3) \leq d(4) \leq d(5) = d(6) = d^*$

## Correctness of SBF/-1

---

- SBF/-1 converges due to monotonicity
  
- Remaining question:
  - Can we guarantee that SBF/-1 converges to shortest path?

# Correctness of SBF/-1

- Common between SBF/ $\infty$  and SBF/-1: they solve the Bellman equation

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

where  $d_D = 0$ .

- We have proven SBF/ $\infty$  is the shortest path solution.
- SBF/-1 computes shortest path if Bellman equation has a unique solution.

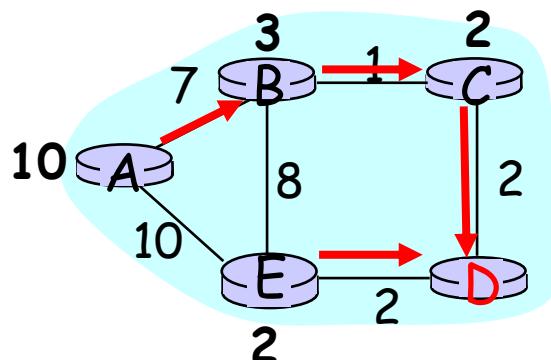
$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

## Uniqueness of Solution to BE

- Assume another solution  $d$ , we will show that  $d = d^*$

case 1: we show  $d \geq d^*$

Since  $d$  is a solution to BE, we can construct paths as follows: for each  $i$ , pick a  $j$  which satisfies the equation; since  $d^*$  is shortest,  $d \geq d^*$



$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

## Uniqueness of Solution to BE

Case 2: we show  $d \leq d^*$

assume we run SBF with two initial configurations:

- one is  $d$
- another is  $\text{SBF}/\infty (d^\infty)$ ,

-> monotonicity and convergence of  $\text{SBF}/\infty$  imply that  $d \leq d^*$

## Discussion

---

- Will SBF converge under other non-negative initial conditions?
- Problems of running *synchronous* BF?

# Outline

---

- Admin and recap
- Network overview
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - *Distributed distance vector protocols*
        - synchronous Bellman-Ford (SBF)
      - *asynchronous Bellman-Ford (ABF)*

# Asynchronous Bellman-Ford (ABF)

- No notion of global iterations
  - each node updates at its own pace
- Asynchronously each node  $i$  computes

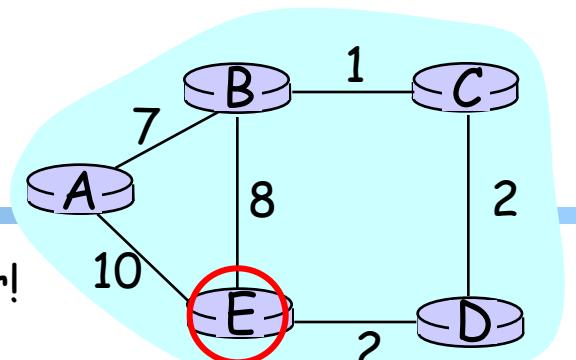
$$d_i = \min_{j \in N(i)} (d_{ij} + d_j^i)$$

using last received value  $d_j^i$  from neighbor  $j$ .

- Asynchronously node  $j$  sends its estimate to its neighbor  $i$ :
  - We assume that there is an upper bound on the delay of estimate packet

# ABF: Example

Below is just one step! The protocol repeats forever!



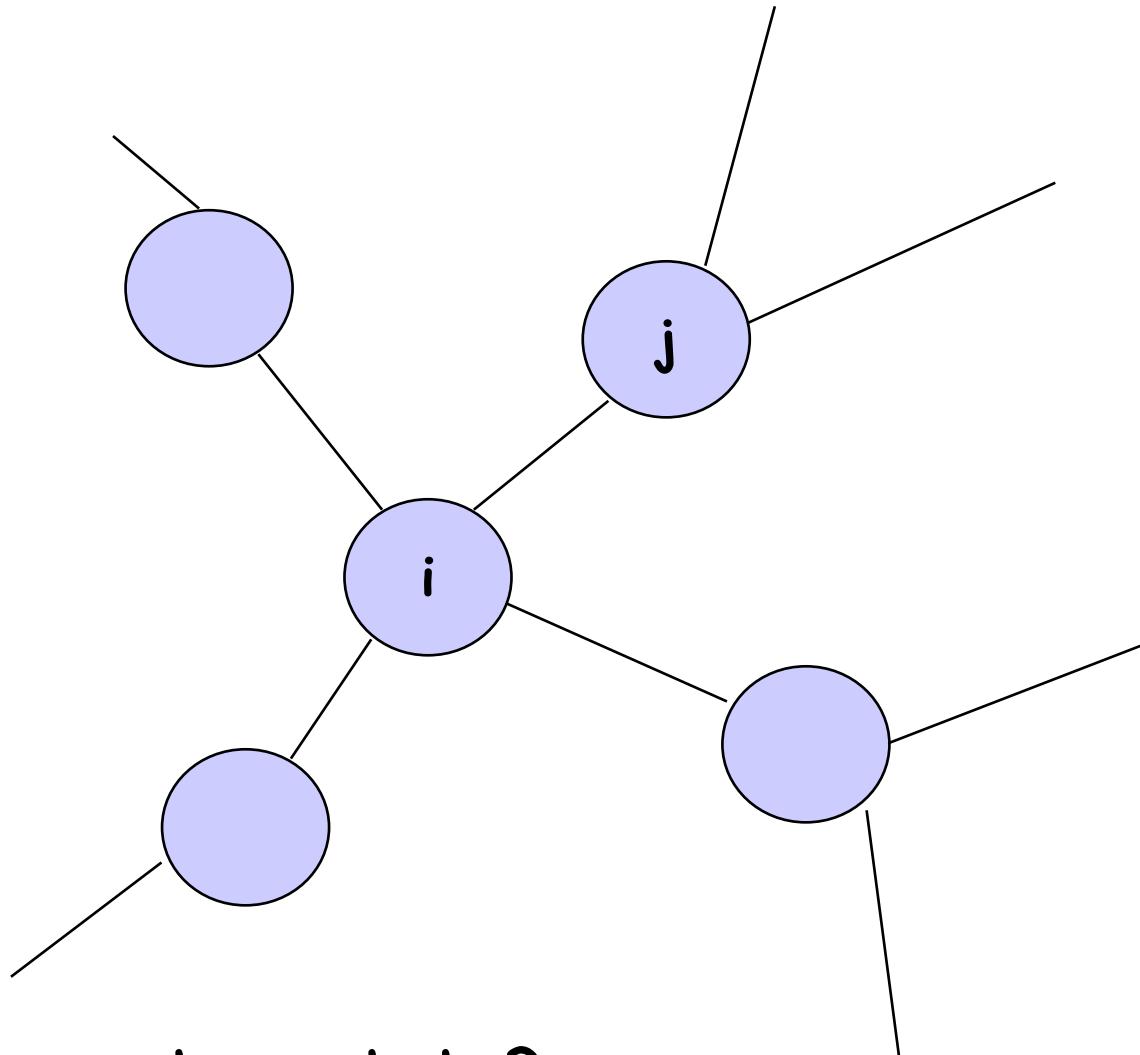
| distance tables from neighbors |          |          |          | computation |          |          | E's distance table | distance table E sends to its neighbors |
|--------------------------------|----------|----------|----------|-------------|----------|----------|--------------------|---|
| d <sub>E</sub> ( )             | A        | B        | D        | A           | B        | D        |                    |   |
| A                              | 0        | 7        | $\infty$ | 10          | 15       | $\infty$ | A: 10              | A: 10                                   |
| B                              | 7        | 0        | $\infty$ | 17          | 8        | $\infty$ | B: 8               | B: 8                                    |
| C                              | $\infty$ | 1        | 2        | $\infty$    | 9        | 4        | D: 4               | C: 4                                    |
| D                              | $\infty$ | $\infty$ | 0        | $\infty$    | $\infty$ | 2        | D: 2               | D: 2                                    |
|                                | 10       | 8        | 2        |             |          |          |                    | E: 0                                    |

## Asynchronous Bellman-Ford (ABF)

---

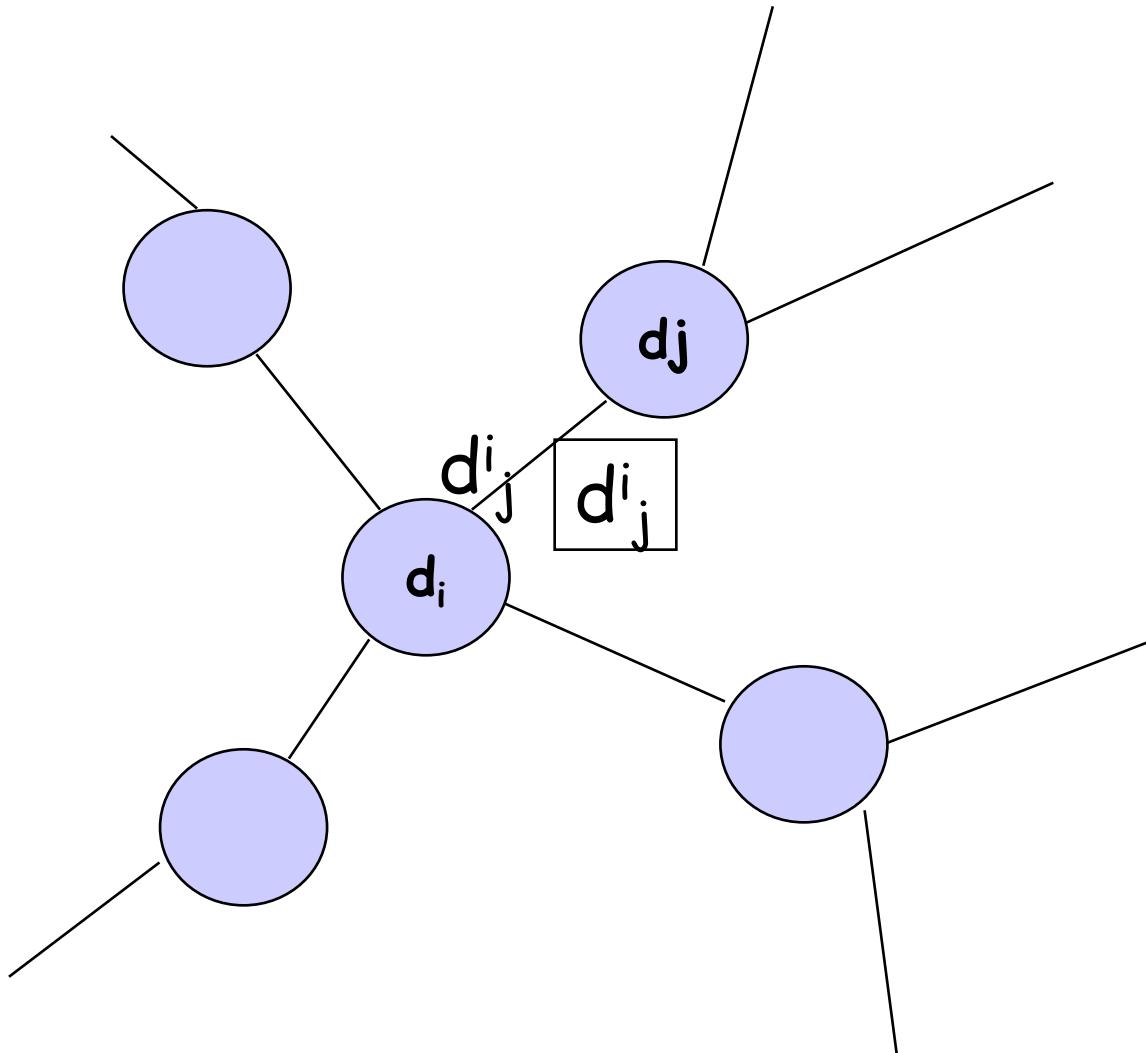
- ABF will eventually converge to the shortest path
  - links can go down and come up - but if topology is stabilized after some time  $t$  and connected, ABF will eventually converge to the shortest path !

# ABF Convergence Proof Complexity: Complex System State



What is system state?

# System State



three types of distance state from node j:

- $d_j$ : current distance estimate state at node j
- $d_{i,j}^i$ : last  $d_j$  that neighbor i received
- $d_{i,j}^t$ : those  $d_j$  that are still in transit to neighbor i

# ABF Convergence Proof: The Sandwich Technique

## □ Basic idea:

- bound system state using extreme states

## □ Extreme states:

- SBF/ $\infty$ ; call the sequence U()
- SBF/-1; call the sequence L()

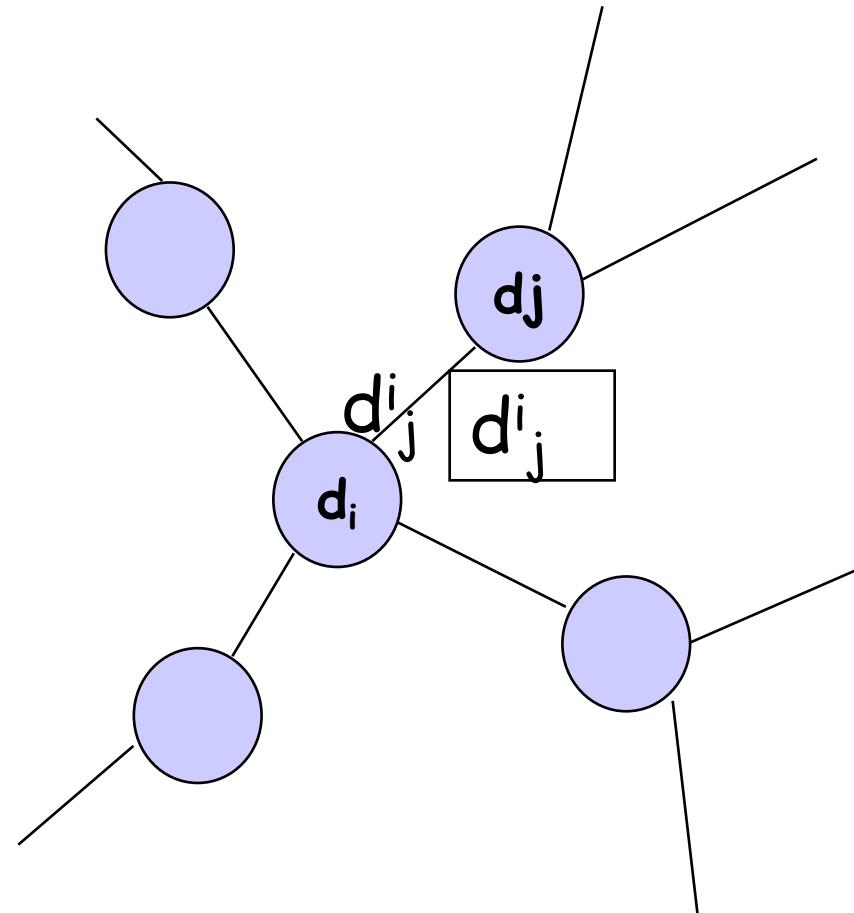
# ABF Convergence

---

- Consider the time when the topology is stabilized as time 0
- $U(0)$  and  $L(0)$  provide upper and lower bounds at time 0 on all corresponding elements of states
  - $L_j(0) \leq d_j \leq U_j(0)$  for all  $d_j$  state at node j
  - $L_j(0) \leq d^i_j \leq U_j(0)$
  - $L_j(0) \leq$  update messages  $d^i_j \leq U_j(0)$

# ABF Convergence

- $d_j$ 
  - after at least one update at node  $j$ :  
 $d_j$  falls between  
 $L_j(1) \leq d_j \leq U_j(1)$
  
- $d^i_j$ :
  - eventually all  $d^i_j$  that are only bounded by  $L_j(0)$  and  $U_j(0)$  are replaced with in  $L_j(1)$  and  $U_j(1)$



# Asynchronous Bellman-Ford: Summary

---

## ❑ Distributed

- each node communicates its routing table to its directly-attached neighbors

## ❑ Iterative

- continues periodically or when link changes, e.g. detects a link failure

## ❑ Asynchronous

- nodes need *not* exchange info/iterate in lock step!

## ❑ Convergence

- in finite steps, independent of initial condition if network is connected

## Summary: Distributed Distance-Vector

---

- Tool box: a key technique for proving convergence (**liveness**) of distributed protocols: **monotonicity** and **bounding-box (sandwich)** design
  - Consider two configurations  $d(t)$  and  $d'(t)$ :
    - if  $d(t) \leq d'(t)$ , then  $d(t+1) \leq d'(t+1)$
  - Identify two extreme configurations to sandwich any real configurations

# Outline

---

- Admin and recap
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
        - synchronous Bellman-Ford (SBF)
        - asynchronous Bellman-Ford (ABF)
  - *properties of DV*

# Properties of Distance-Vector Algorithms

- Good news propagate fast

| A | B        | C        | D        | E        |                   |
|---|----------|----------|----------|----------|-------------------|
|   |          |          |          |          | Initially         |
| 1 | $\infty$ | $\infty$ | $\infty$ | $\infty$ | After 1 exchange  |
| 1 | 2        | $\infty$ | $\infty$ |          | After 2 exchanges |
| 1 | 2        | 3        | $\infty$ |          | After 3 exchanges |
| 1 | 2        | 3        | 4        |          | After 4 exchanges |

# Properties of Distance-Vector Algorithms

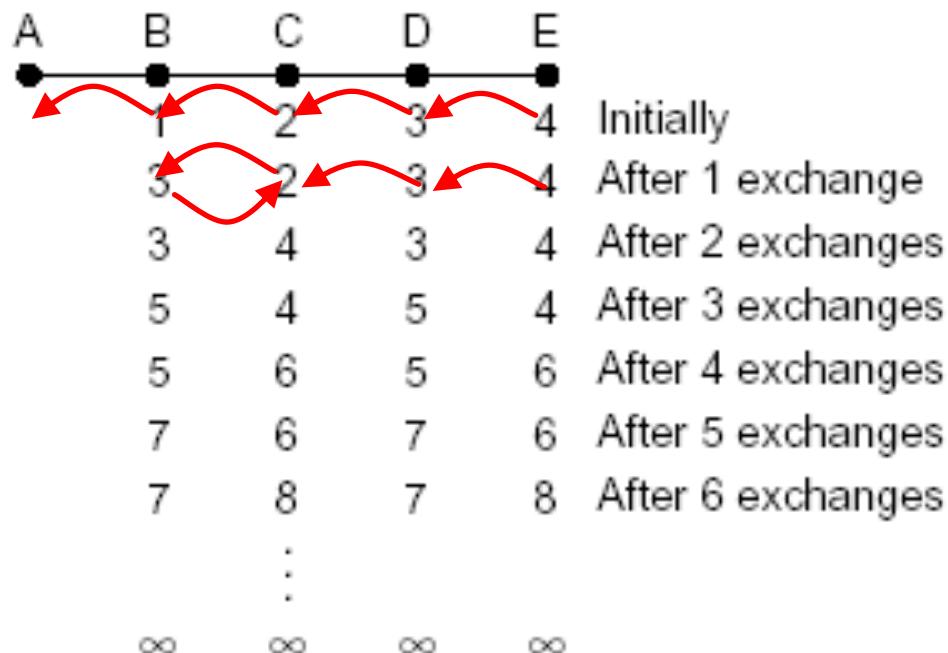
- Bad news propagate slowly

| A-B link down | A | B | C | D | E |                   |
|---------------|---|---|---|---|---|-------------------|
|               | 1 | 2 | 3 | 4 |   | Initially         |
| A-B link down | 3 | 2 | 3 | 4 |   | After 1 exchange  |
|               | 3 | 4 | 3 | 4 |   | After 2 exchanges |
|               | 5 | 4 | 5 | 4 |   | After 3 exchanges |
|               | 5 | 6 | 5 | 6 |   | After 4 exchanges |
|               | 7 | 6 | 7 | 6 |   | After 5 exchanges |
|               | 7 | 8 | 7 | 8 |   | After 6 exchanges |
|               |   |   |   |   |   |                   |
|               | ⋮ |   |   |   |   |                   |
|               | ∞ | ∞ | ∞ | ∞ |   |                   |

- This is called the *counting-to-infinity* problem
- Q: what causes counting-to-infinity?

# Counting-To-Infinity is Because of Routing Loop

- Counting-to-infinity is caused by a routing loop, which is a **global state** (consisting of the nodes' local states) at a global moment (observed by an oracle) such that there exist nodes A, B, C, ... E such that A (locally) thinks B as next hop, B thinks C as next hop, ... E thinks A as next hop



# Discussion

---

- Why avoid routing loops is hard?
  
- Any proposals to avoid routing loops?

# Outline

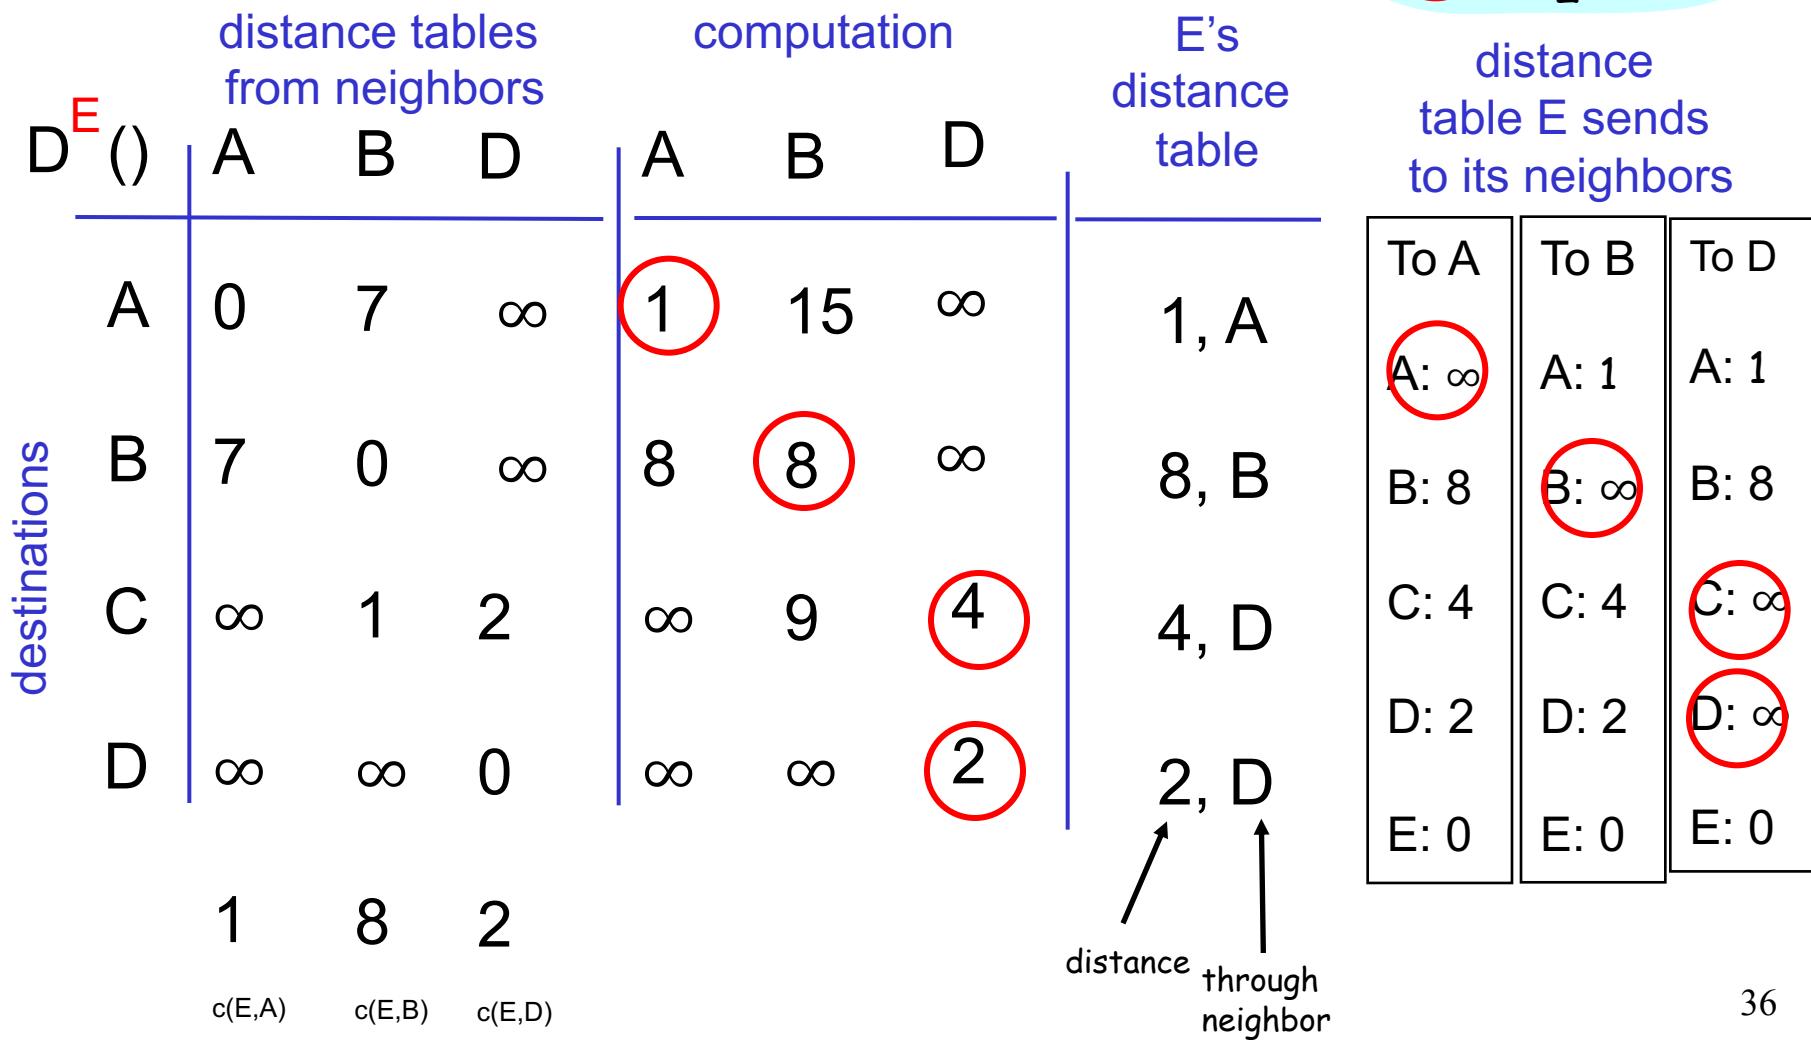
---

- Admin and recap
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
        - synchronous Bellman-Ford (SBF)
        - asynchronous Bellman-Ford (ABF)
        - properties of DV
          - DV w/ loop prevention

➤ *reverse poison*

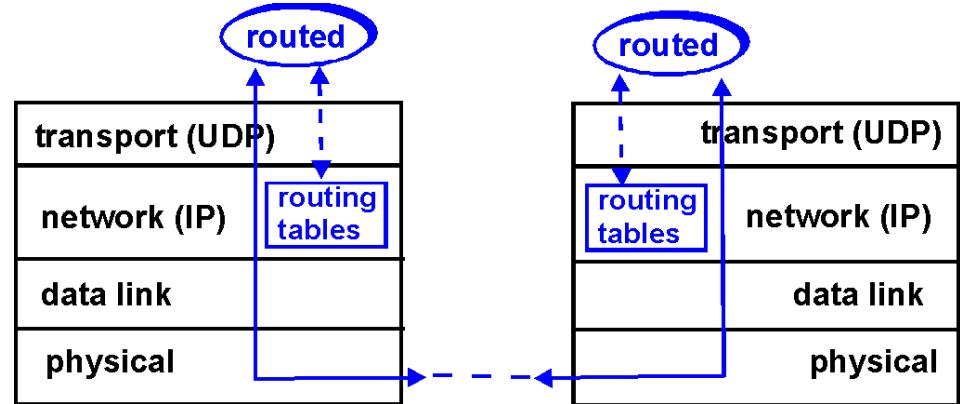
# The Reverse-Poison (Split-horizon) Hack

If the path to dest is through neighbor h, report  $\infty$  to neighbor h for dest.



# DV+RP => RIP ( Routing Information Protocol)

- Included in BSD-UNIX Distribution in 1982
- Link cost: 1
- Distance metric: # of hops
- Distance vectors
  - exchanged every 30 sec via Response Message (also called **advertisement**) using UDP
  - each advertisement: route to up to 25 destination nets



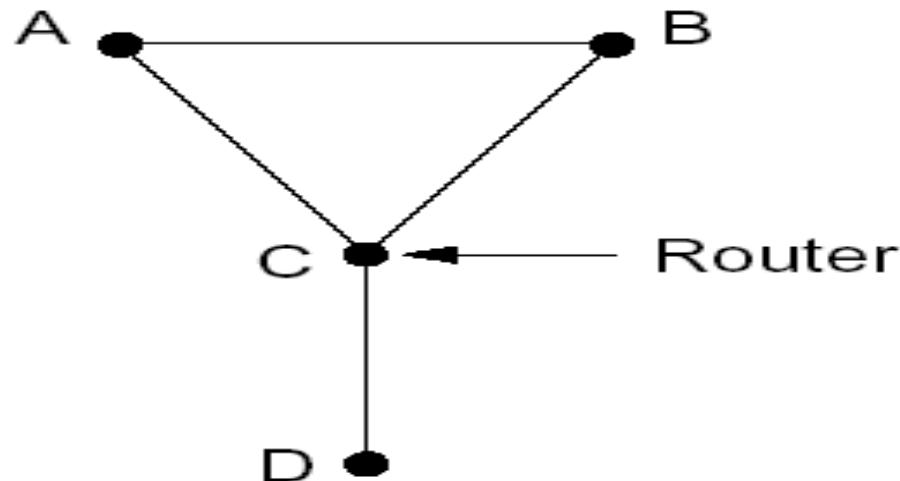
## RIP: Link Failure and Recovery

If no advertisement heard after 180 sec -->  
neighbor/link declared dead

- routes via neighbor invalidated
- new advertisements sent to neighbors
- neighbors in turn send out new advertisements  
(if tables changed)
- link failure info quickly propagates to entire net
- reverse-poison used to prevent **ping-pong** loops
- set infinite distance = 16 hops (why?)

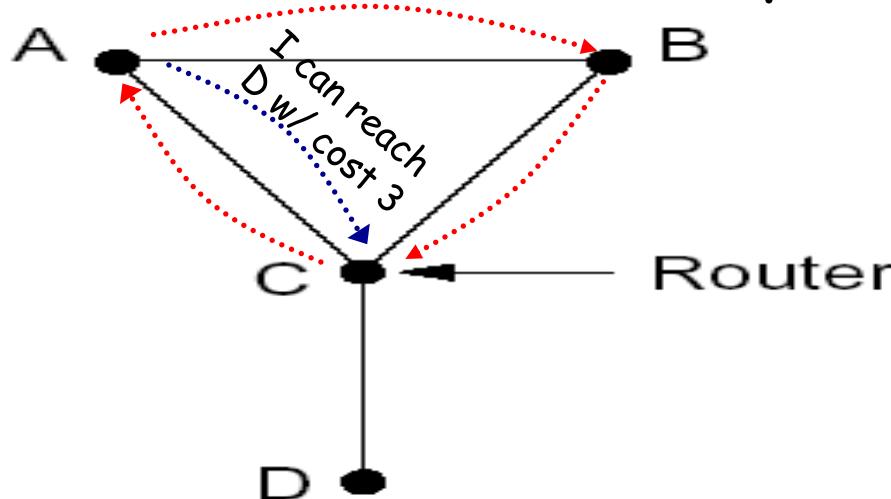
## General Routing Loops and Reverse-poison

- Exercise: Can Reverse-poison guarantee no loop for this network?



## General Routing Loops and Reverse-poison

- Reverse-poison removes two-node loops but may not remove more-node loops



- Unfortunate timing can lead to a loop
  - When the link between C and D fails, C will set its distance to D as  $\infty$
  - A receives the bad news ( $\infty$ ) from C, A will use B to go to D
  - A sends the news to C
  - C sends the news to B

# Outline

---

- Admin and recap
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
        - synchronous Bellman-Ford (SBF)
        - asynchronous Bellman-Ford (ABF)
        - properties of DV
          - DV w/ loop prevention
            - reverse poison
            - *destination-sequenced DV (DSDV)*

# Destination-Sequenced Distance Vector protocol (DSDV)

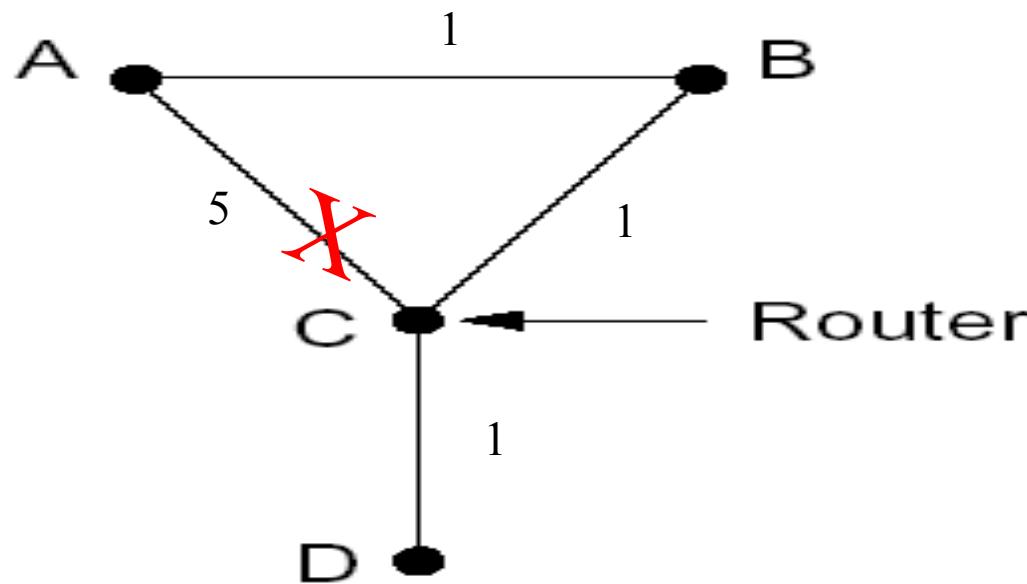
- Basic idea: use sequence numbers to partition computation
  - tags each route with a sequence number
  - each destination node D periodically advertises monotonically increasing even-numbered sequence numbers
  - when a node realizes that **the link it uses to reach destination D is broken**, it advertises an **infinite metric** and a **sequence number which is one greater** than the previous route (i.e., an odd seq. number)
    - the route is repaired by a later even-number advertisement from the destination

## DSDV: More Detail



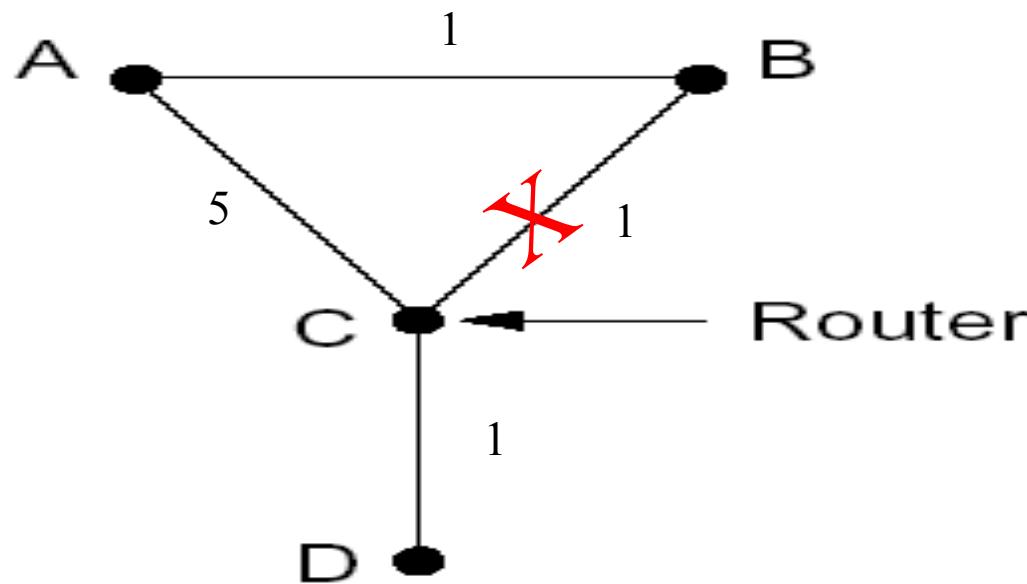
- Let's assume the destination node is D
- There are optimizations but we present a simple version:
  - each node B maintains  $(S^B, d^B)$ , where  $S^B$  is the sequence number at B for destination D and  $d^B$  is the best distance using a neighbor from B to D
- Both periodical and triggered updates
  - periodically: D increases its seq. by 2 and broadcasts with  $(S^D, 0)$
  - if B is using C as next hop to D and B discovers that C is no longer reachable
    - B increases its sequence number  $S^B$  by 1, sets  $d^B$  to  $\infty$ , and sends  $(S^B, d^B)$  to all neighbors

# Example



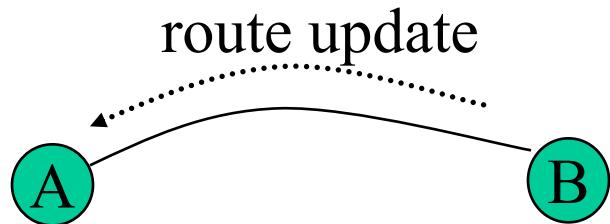
Will this trigger an update?

# Example



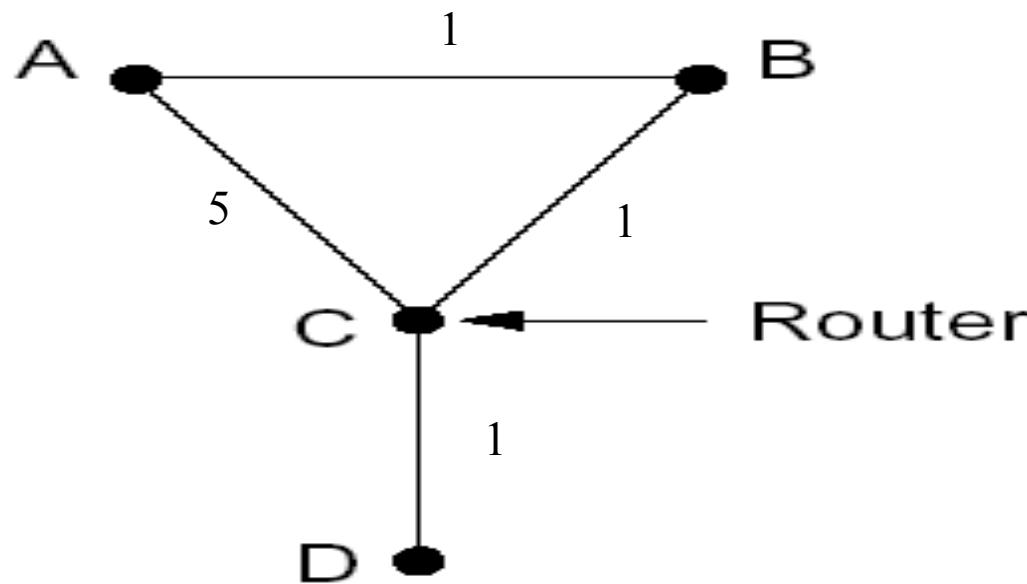
Will this trigger an update?

# DSDV: Update Alg.



- Consider simple version, no optimization
- Update after receiving a message
  - assume B sends to A its current state  $(S^B, d^B)$
  - when A receives  $(S^B, d^B)$ 
    - if  $S^B > S^A$ , then  
**// always update if a higher seq#**
      - »  $S^A = S^B$
      - » if  $(d^B == \infty)$   $d^A = \infty$ ; else  $d^A = d^B + d(A, B)$
    - else if  $S^A == S^B$ , then
      - » if  $d^A > d^B + d(A, B)$   
**// update for the same seq# only if better route**
        - $d^A = d^B + d(A, B)$  and uses B as next hop

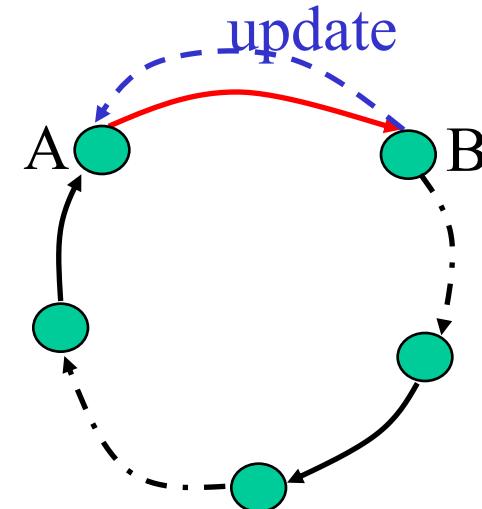
# Example



Exercise: update process after D increases its seq# to next even number.

# Claim: DSDV will NEVER Form a Loop

- Initially no loop (no one has next hop so no loop)
- Derive contradiction if a loop forms after a node processes an update,
  - e.g., when A receives the update from B, A decides to use B as next hop and forms a loop

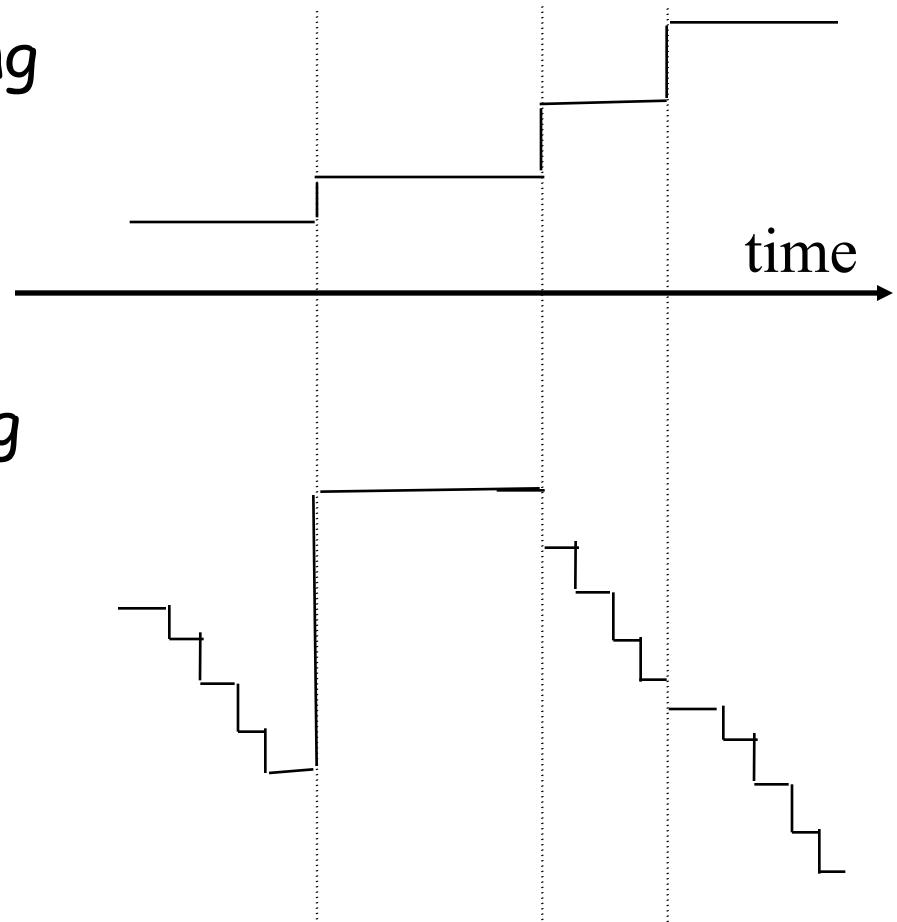


# Technique: Global Invariants

- **Global Invariant** is a very effective method in understanding **safety** of distributed asynchronous protocols
- Invariants are defined over the states of the distributed nodes
  
- Consider any node B.
- Let's identify some invariants over the state of node B, i.e.,  $(S^B, d^B)$ .

# Invariants of a Single Node B

- Some invariants about the state of a node B
  - [I1]  $S^B$  is non-decreasing



- [I2]  $d^B$  is non-increasing  
for the same  
sequence number

# Invariants of if A Considers B as Next Hop

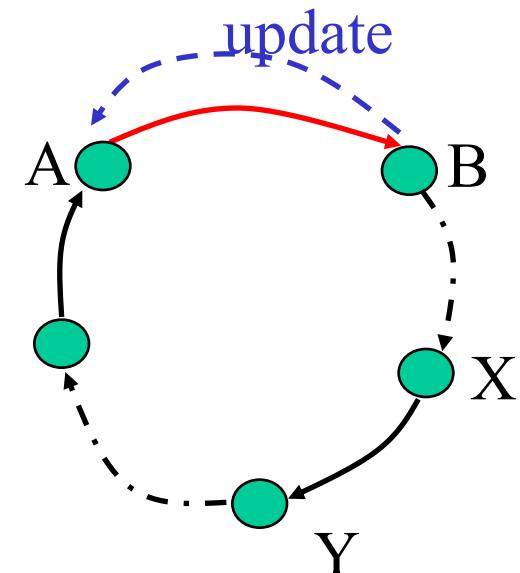
- Some invariants if A considers B as next hop



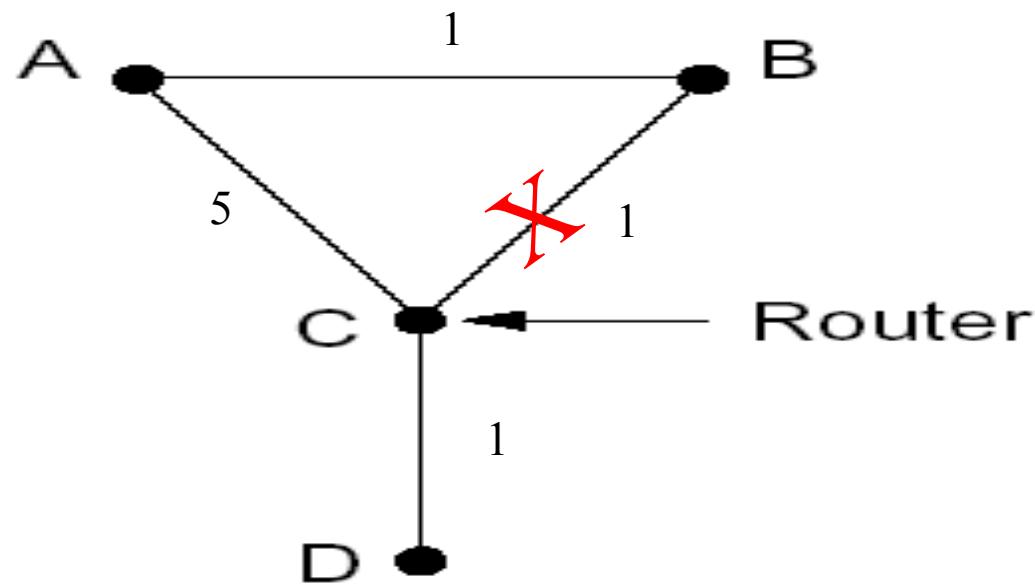
- [I3]  $d^A$  is not  $\infty$
- [I4]  $S^B \geq S^A$   
because A is having the seq# which B last sent to A; B's seq# might be increased after B sent its state
  - [I5] if  $S^B == S^A$
  - then  $d^B < d^A$  because  $d^A$  is based on  $d^B$  which B sent to A some time ago,  $d^B < d^A$  since all link costs are positive;  $d^B$  might be decreased after B sent its state

# Loop Freedom of DSDV

- Consider a critical moment
  - A starts to consider B as next hop, and we have a loop
- According to invariant I4 for each link in the loop (X considers Y as next hop) :  
 $S^Y \geq S^X$
- Two cases:
  - exists  $S^Y > S^X$ 
    - by transition along the loop  $S^B > S^B$
  - all nodes along the loop have the same sequence number
    - apply I5, by transition along the loop  $d^B > d^B$



# Issue of DSDV



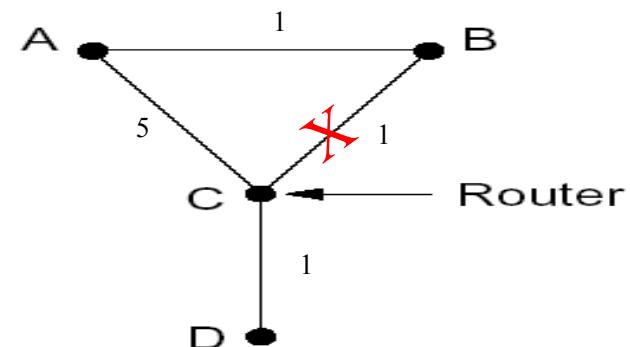
# Outline

---

- Admin and recap
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
        - synchronous Bellman-Ford (SBF)
        - asynchronous Bellman-Ford (ABF)
        - properties of DV
          - DV w/ loop prevention
            - reverse poison
            - destination-sequenced DV (DSDV)
            - *diffusive update algorithm (DUAL) and EIGRP*

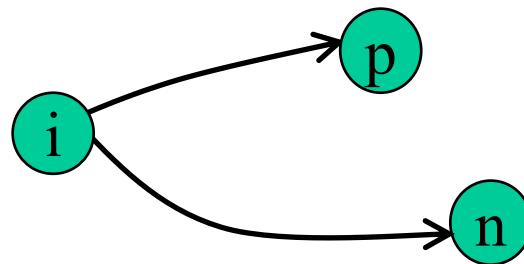
# Basic Idea

- DSDV guarantees no loop, but at the price of not using any backup path before destination re-announces reachability.
- Basic idea: Sufficient condition to guarantee no loop using backup paths (called switching)?



## Key Idea: Feasible Successors

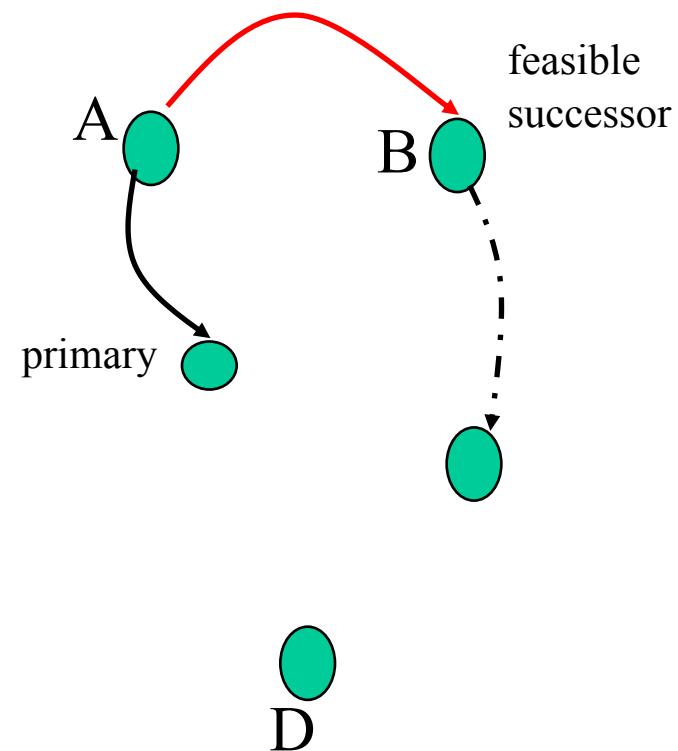
- If the reported distance of a neighbor n is lower than the total distance using primary (current shortest), the neighbor n is a feasible successor



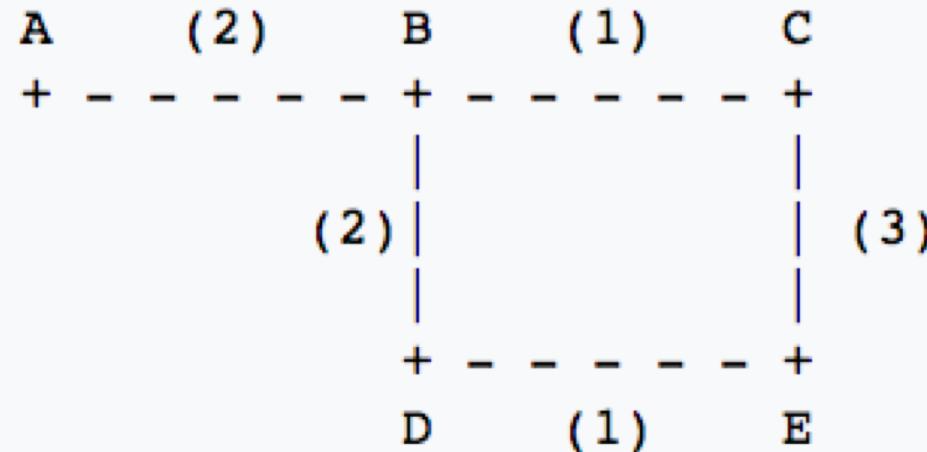
$$d_n + d_{i \rightarrow n} \geq d_{\text{primary}} + d_{i \rightarrow \text{primary}} > d_n$$

# Intuition

- Since the reported distance of B is lower than my total distance, B cannot be using me (along a path) to reach the destination



# Example



- Assume A is destination, consider E

|            | Reported Dist. | Total Dist. |
|------------|----------------|-------------|
| Neighbor C | 3              | 6           |
| Neighbor D | 4              | 5           |

# Summary: Distance Vector Routing

- Basic DV protocol
  - take away: use monotonicity as a technique to understand liveness/convergence
    - highly recommended reading of Bersekas/Gallager chapter
- Fix counting-to-infinity problem
  - DSDV
    - Idea: uses sequence number to avoid routing loops
      - seq# partitions routing updates from different outside events
      - within same event, no loop so long each node only decreases its distance
    - Analysis: use global invariants to understand/design safety/no routing loops
  - EIRGP (DUAL)
    - Idea: introduces a sufficient condition for local recovery

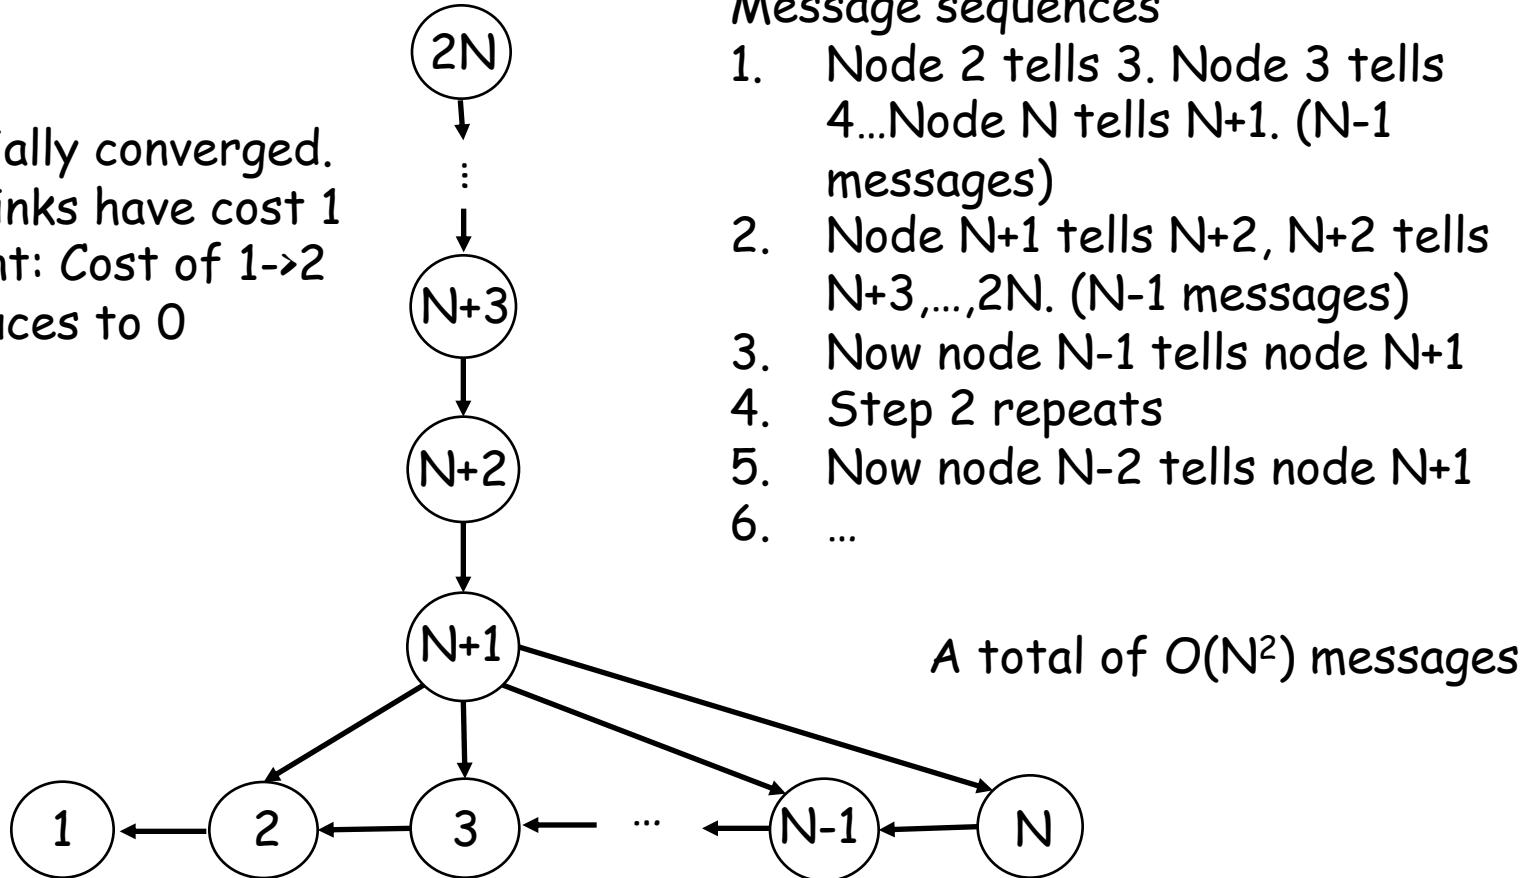
## Discussion: Distance Vector Routing

---

- What do you like about distributed, distance vector routing?
  
- What do you **not** like about distributed, distance vector routing?

# Churns of DV: One Example

Initially converged.  
All links have cost 1  
Event: Cost of  $1 \rightarrow 2$   
reduces to 0



# Outline

---

- Admin and recap
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
      - *Link state protocols (distributed state synchronization)*

# Link-State Routing

- Basic idea: Not distributed computing, only distributed state distribution
- Net topology, link costs are distributed to all nodes
  - all nodes have same info
  - Each node computes its shortest paths from itself to all other nodes
    - standard Dijkstra's algorithm as path compute alg
    - Allows multiple same-cost paths
    - Multiple cost metrics per link (for type of service routing)
- Most commonly used routing protocol (e.g., OSPF/ISIS) by most networks in Internet

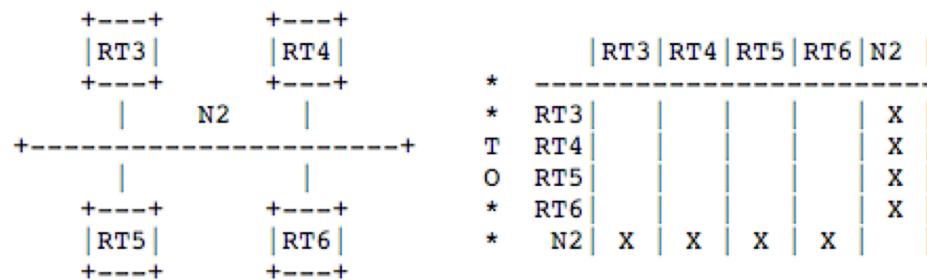
# Example: Link State and Directed Graph (OSPFv2)

\*\*FROM\*\*



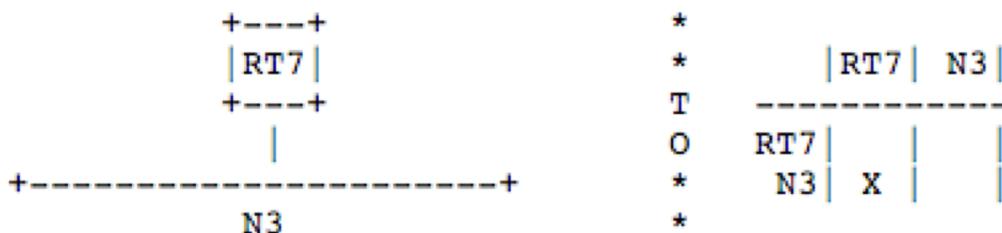
Physical point-to-point networks

\*\*FROM\*\*



Multi-access networks

\*\*FROM\*\*



Stub multi-access networks

# Example: Link State and Directed Graph (OSPFv2)

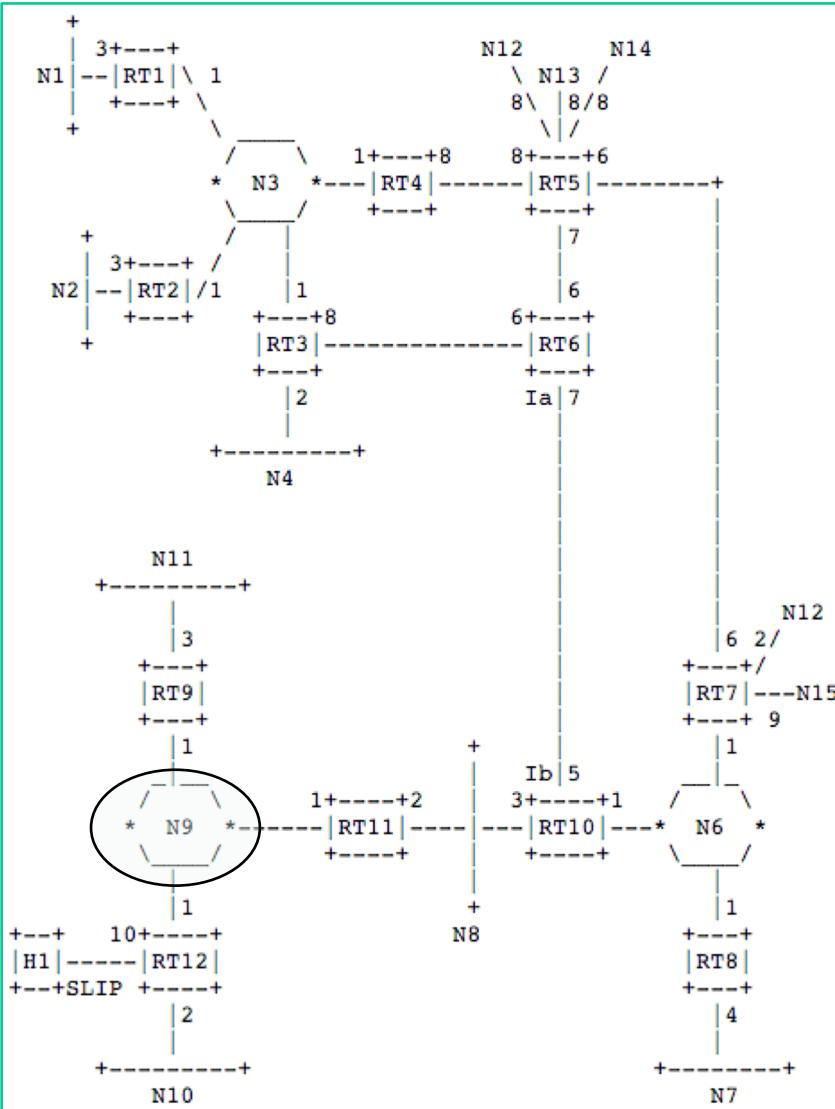


Figure 2: A sample Autonomous System

| **FROM** |      |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|----------|------|----|----|----|----|----|----|----|----|----|----|----|----|----|---|
| RT       | RT   | RT | RT | RT | RT | RT | RT | RT | RT | RT | N3 | N6 | N8 | N9 |   |
| RT1      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| RT2      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| RT3      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| RT4      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| RT5      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| RT6      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| RT7      |      |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| *        | RT8  |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| *        | RT9  |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| T        | RT10 |    |    |    |    |    |    |    |    |    |    |    | 0  | 0  | 0 |
| O        | RT11 |    |    |    |    |    |    |    |    |    |    |    | 0  | 0  | 0 |
| *        | RT12 |    |    |    |    |    |    |    |    |    |    |    | 0  |    |   |
| *        | N1   | 3  |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N2   |    | 3  |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N3   | 1  | 1  | 1  | 1  |    |    |    |    |    |    |    |    |    |   |
|          | N4   |    |    | 2  |    |    |    |    |    |    |    |    |    |    |   |
|          | N6   |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N7   |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N8   |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N9   |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N10  |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N11  |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N12  |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N13  |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N14  |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | N15  |    |    |    |    |    |    |    |    |    |    |    |    |    |   |
|          | H1   |    |    |    |    |    |    |    |    |    |    |    |    | 10 |   |

Figure 3: The resulting directed graph

# Outline

---

- Admin and recap
- Network control plane
  - Routing
    - Link weights assignment
    - Routing computation
      - Distance vector protocols (distributed computing)
      - *Link state protocols (distributed state synchronization)*
        - data structure to be distributed
        - *state distribution protocol*

# Basic Link State Broadcast Protocol

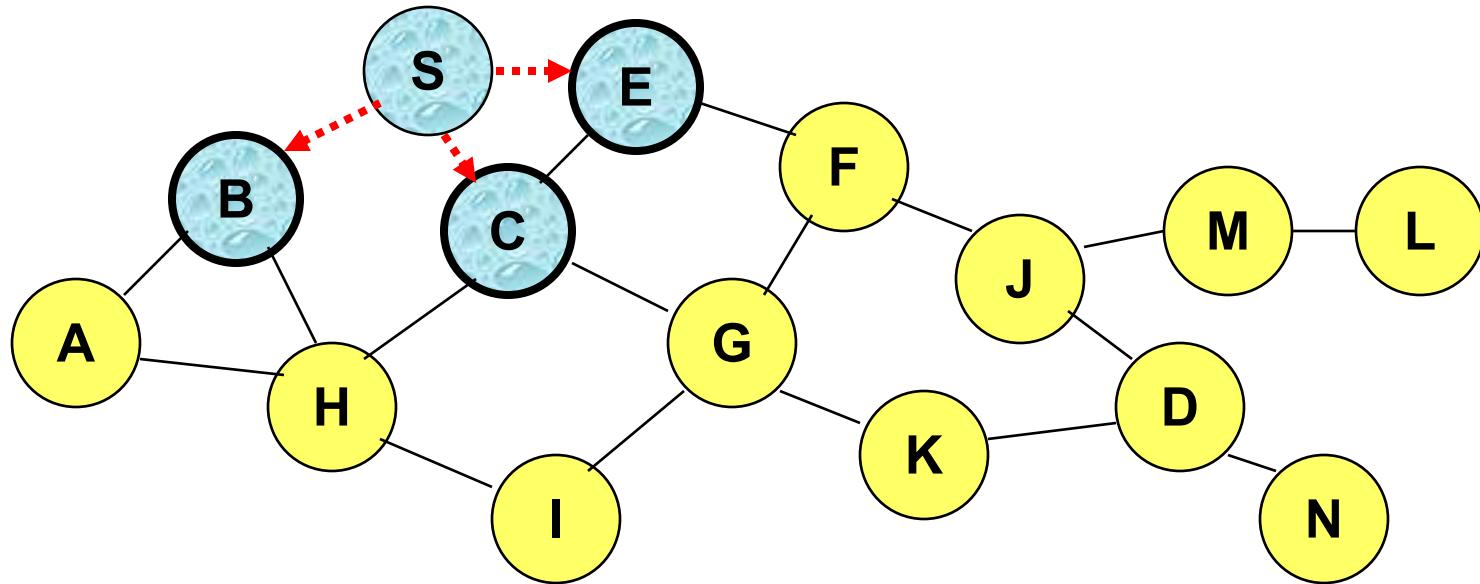
---

Basic event structure at node n

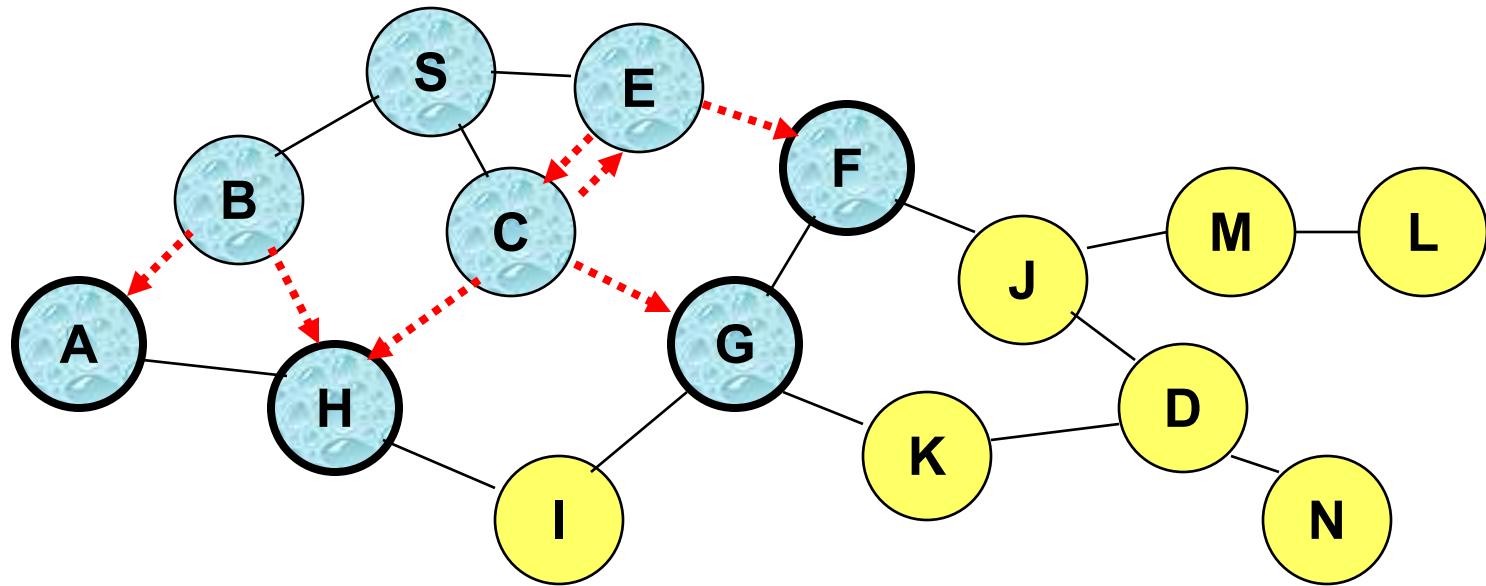
- on initialization:
  - broadcast LSA[e] for each link e connected to n
- on state change to a link e connected to n:
  - broadcast LSA[e] = new status
- on receiving an LSA[e]:
  - if (does not have LSA[e])  
forwards LSA[e] to all links except the incoming link

# Link State Broadcast

Node S updates link states connected to it.



# Link State Broadcast



To avoid forwarding the same link state announcement (LSA) multiple times (forming a loop), each node remembers the received LSAs.

- Second LSA[S] received by E from C is discarded
- Second LSA[S] received by C from E is discarded as well
- Node H receives LSA[S] from two neighbors, and will discard one of them

# Discussion

---

- Issues of the basic link state protocol?
  - Recall: goal is to efficiently distribute to each node to a correct, complete link state map

# Link State Broadcast: Issues

---

- Problem: Out of order delivery
  - link down and then up
  - A node may receive up first and then down
  
- Solution
  - Each link update is given a sequence number: (initiator, seq#, link, status)
    - the initiator should increase the seq# for each new update
  - If the seq# of an update of a link is not higher than the highest seq# a router has seen, drop the update
  - Otherwise, forward it to all links except the incoming link (real implementation using packet buffer)
  
  - Problem of solution: seq# corruption
  - Solution: age field (e.g.,  
<https://tools.ietf.org/html/rfc1583#page-102>)

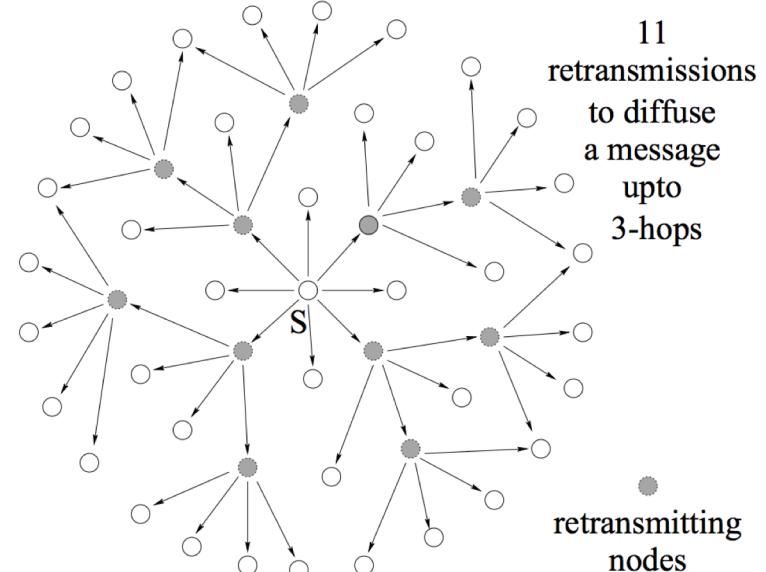
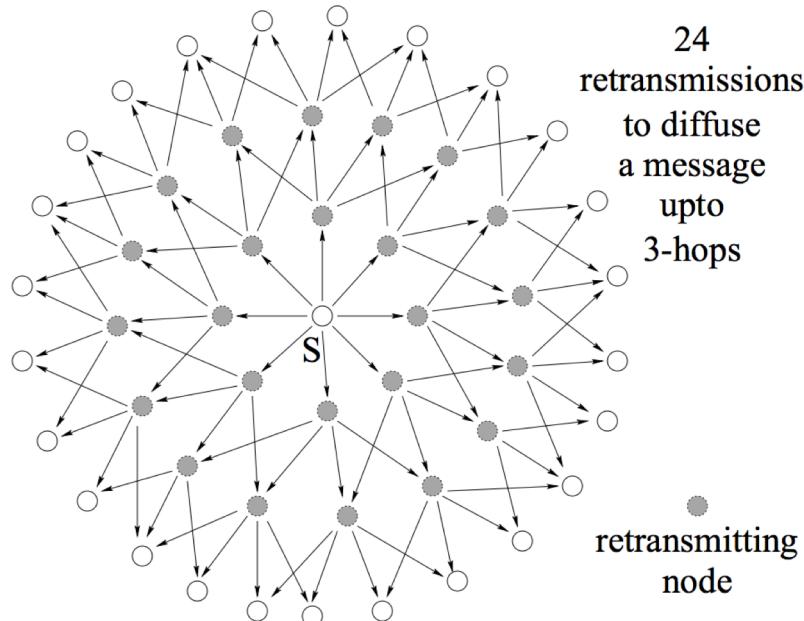
# Link State Broadcast: Issues

---

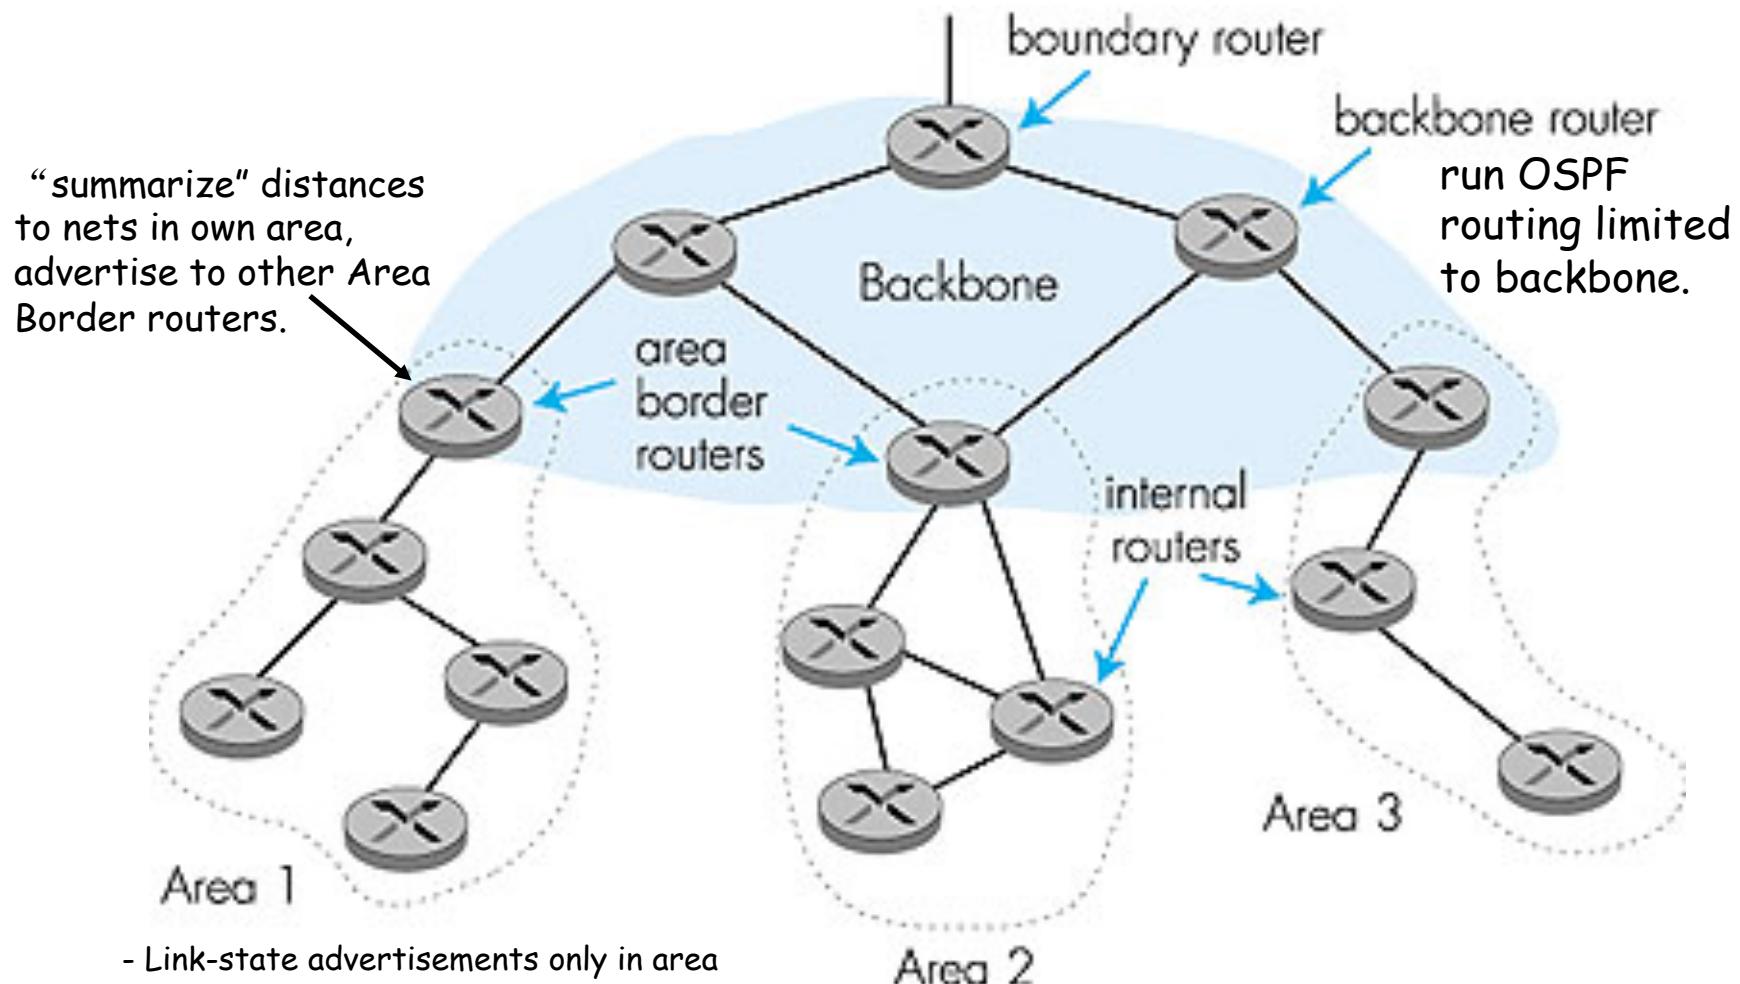
- Problem: network partition and then reconnect, how to sync across the reconnected components
  
- Solution: updates are sent periodically

# Link State Broadcast: Issues

## □ Problem: Broadcast redundancy



# Hierarchical OSPF



**Two-level hierarchy:** local area, backbone.

# Summary: Link State

## □ Basic LS protocol

- take away: instead of computing routing results using distributed computing, distributed computing is for only link state distribution (synchronization)

## □ Link state distribution can still have much complexity, e.g., out of order delivery, partition and reconnect, scalability

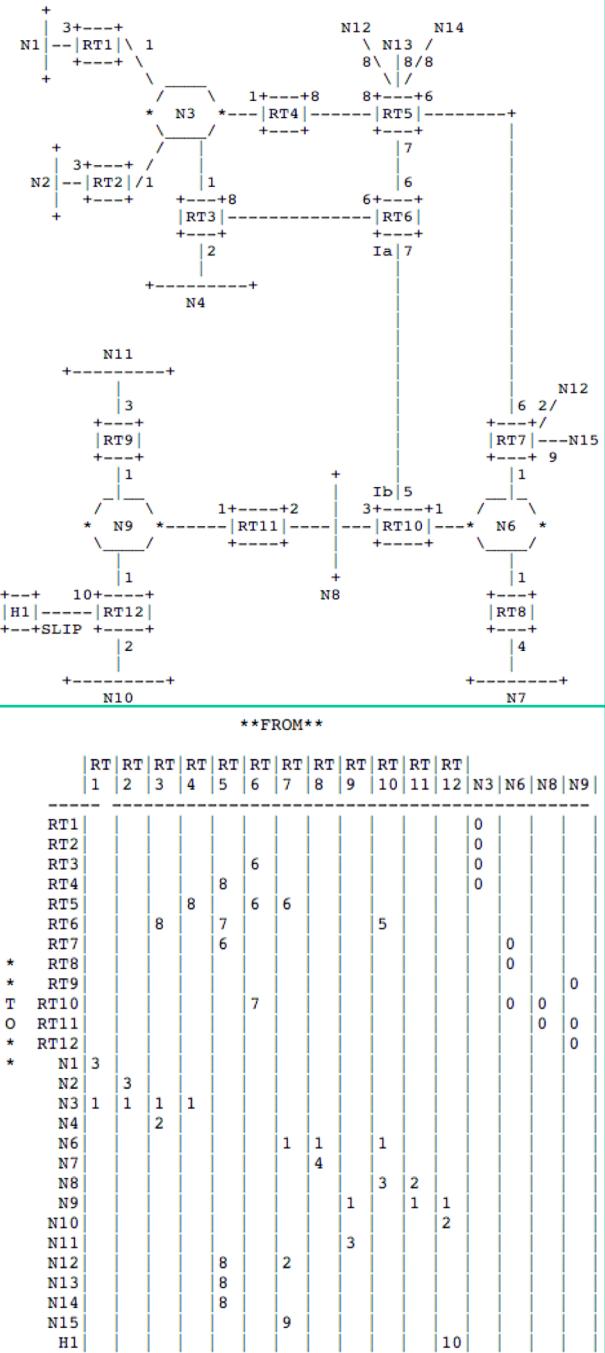


Figure 3: The resulting directed graph