
Network Transport Layer: Reliable Data Transfer; TCP

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf22/index.shtml>

10/25/2022

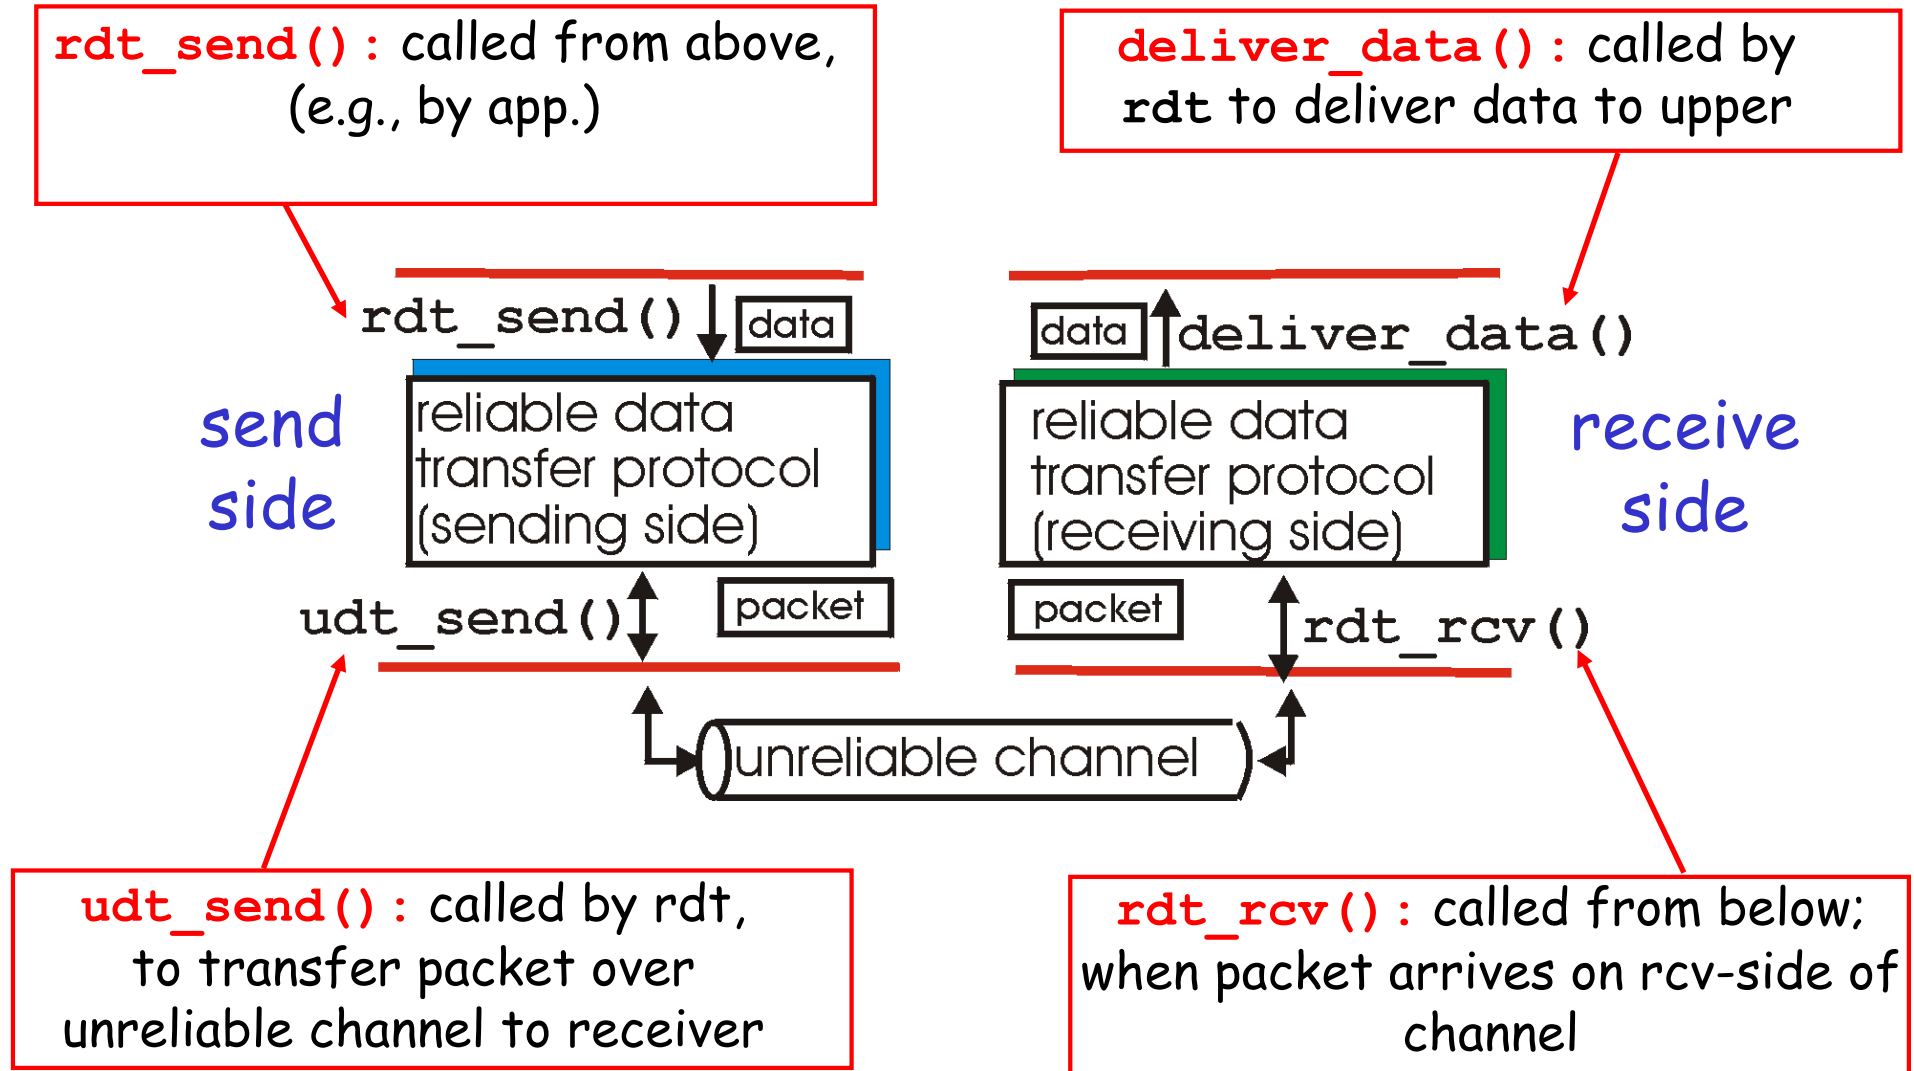
Outline

- ❑ Admin and recap
- ❑ Reliable data transfer

Admin

- ❑ Guest lecture on Oct. 27 class (morning)
 - ❑ Guest lecturer: Ennan Zhai (Alibaba)
 - ❑ Lecture title: Intent-Based Network in Alibaba

Recap: Reliable Data Transfer Context

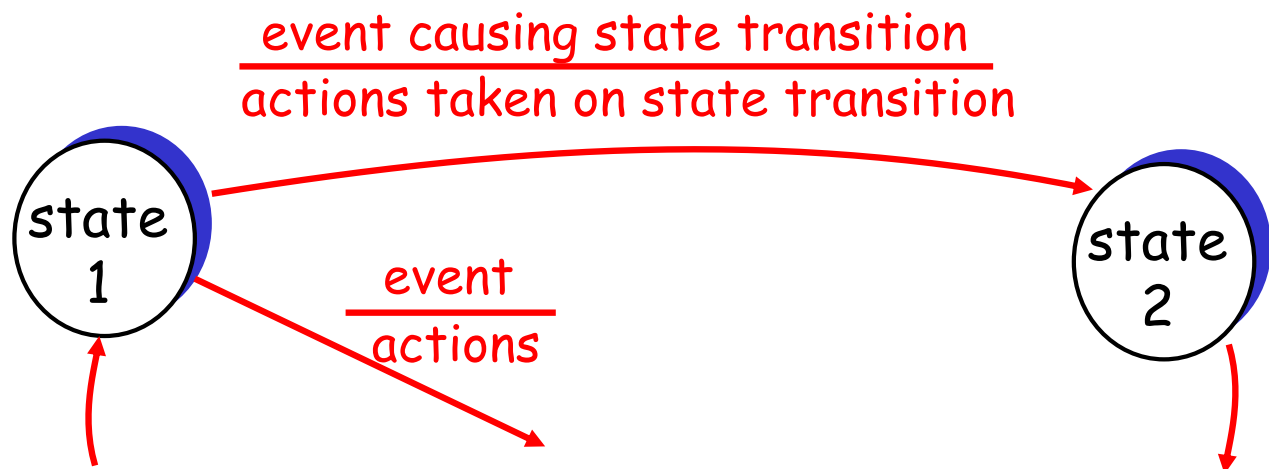


Reliable Data Transfer: Getting Started

We'll:


- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer
 - but control info will flow on both directions !
- use **finite state machines (FSM)** to specify sender, receiver

state: when in this "state" next state uniquely determined by next event




Rdt1.0: reliable transfer over a reliable channel

- separate FSMs for sender, receiver:
 - sender sends data into underlying channel
 - receiver reads data from underlying channel



sender


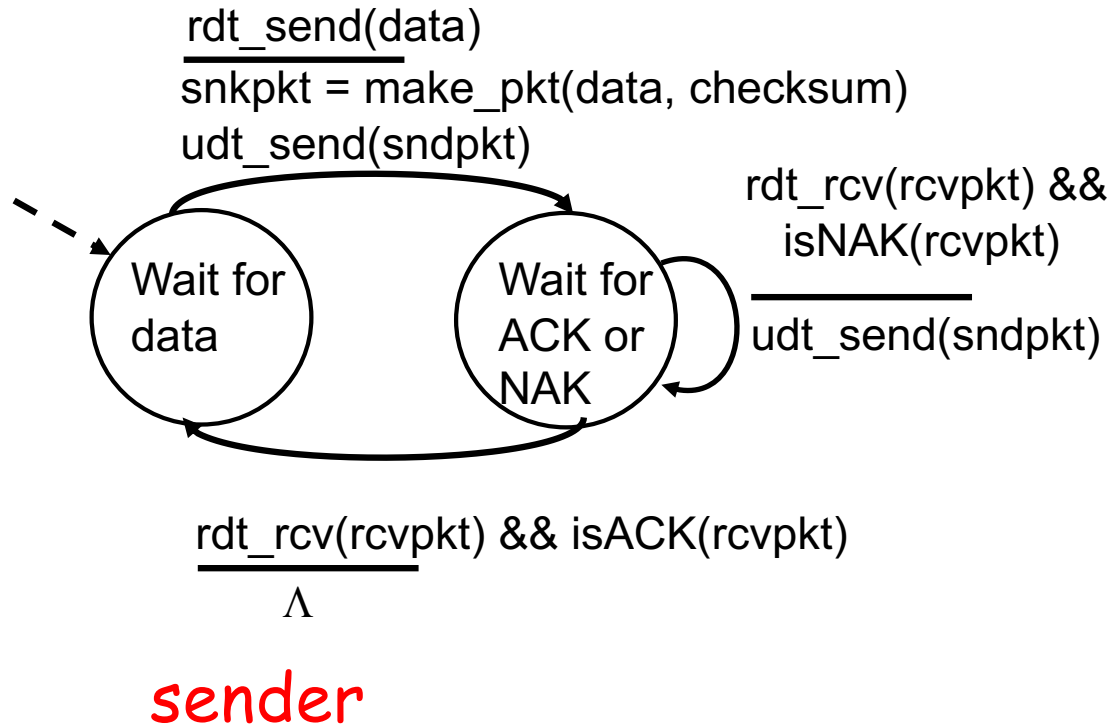


receiver

Exercise: Prove correctness of Rdt1.0.

Correctness: for every single packet, one and only one copy is received by receiver correctly (no error) and in-order


Recap: rdt2.0: Reliability allowing only Data Msg Corruption



rdt2.0 is Incomplete!

What happens if ACK/NAK corrupted?

- Although sender receives feedback, but doesn't know what happened at receiver!



rdt2.1b: Summary

Sender:

- ❑ seq # added to pkt
- ❑ must check if received ACK/NAK corrupted

Receiver:

- ❑ must check if received packet is duplicate
 - by checking if the packet has the expected pkt seq #

rdt2.1c: Summary

Sender:

- ❑ state must “remember” whether “current” pkt has 0 or 1 seq. #

Receiver:

- ❑ must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #

rdt2.2: a NAK-free protocol

- ❑ Same functionality as rdt2.1c, using ACKs only
- ❑ Instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed
- ❑ Duplicate ACK at sender results in same action as NAK: *retransmit current pkt*

rdt3.0: Channels with Errors and Loss

New assumption:

underlying channel can also lose packets (data or ACKs)

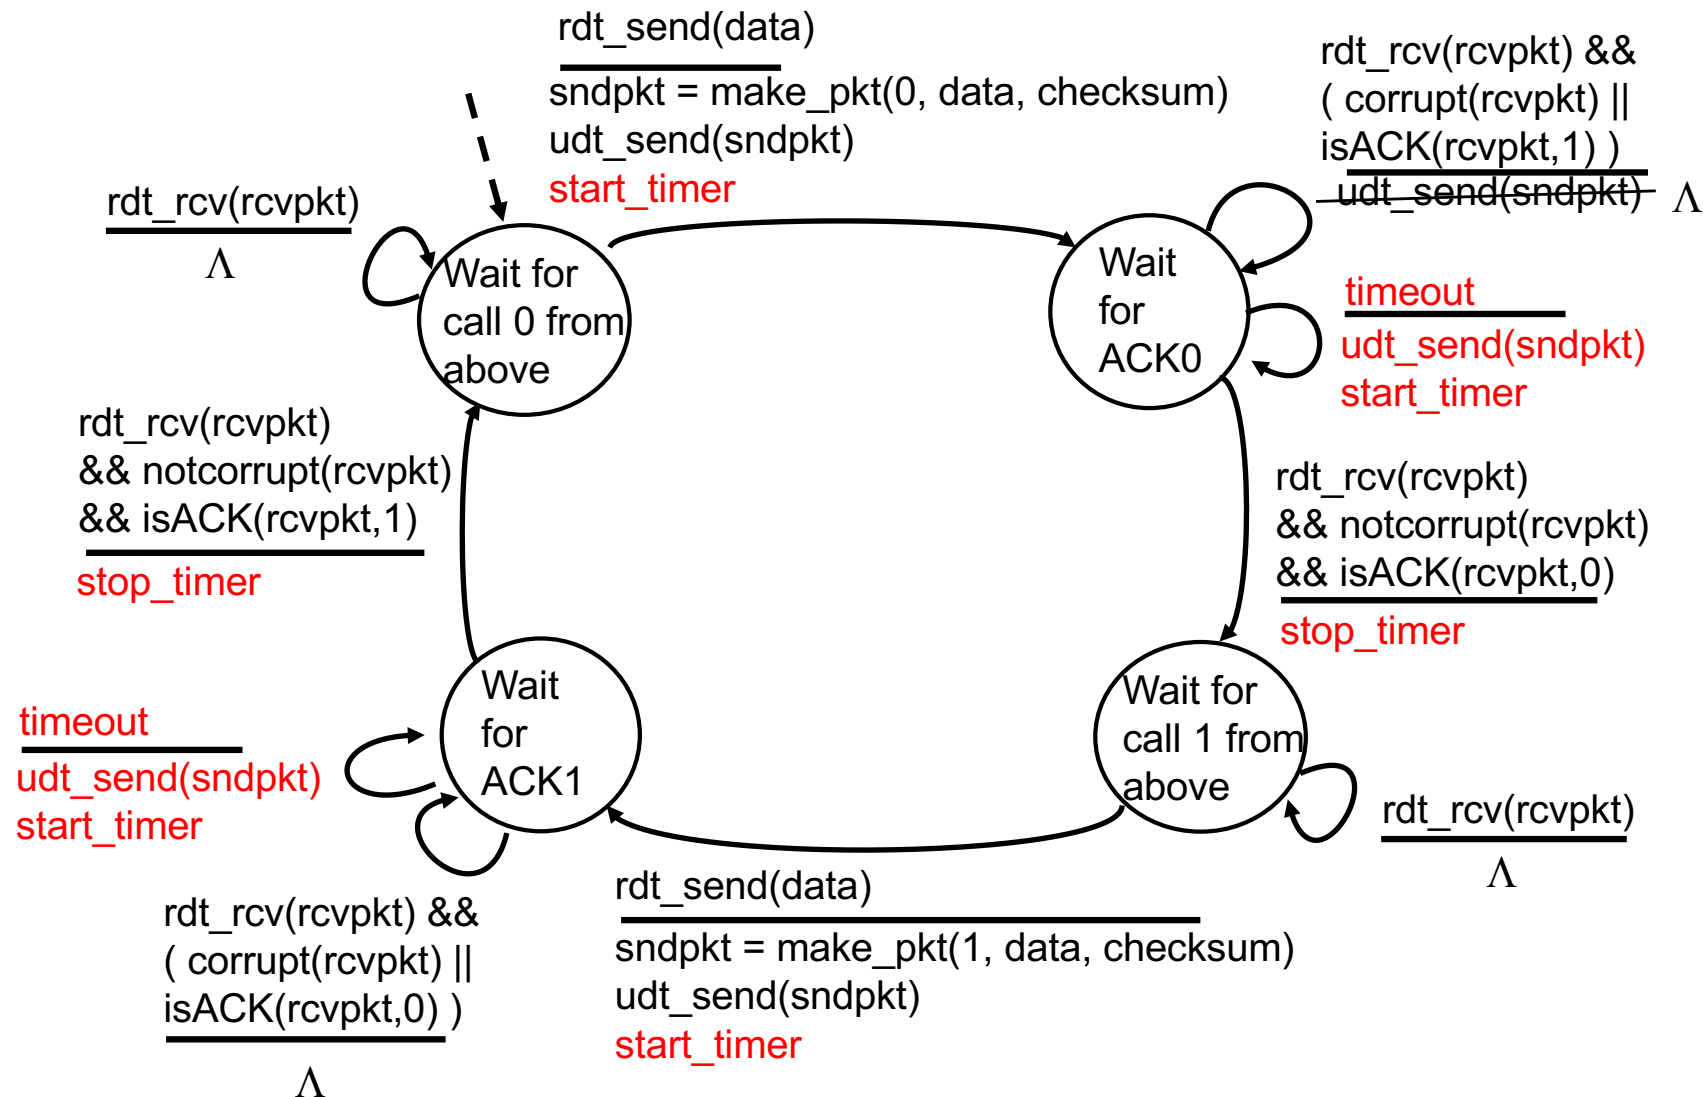
- checksum, seq. #, ACKs, retransmissions will be of help, but not enough

Q: Does rdt2.2 work under losses?


Approach: sender waits “reasonable” amount of time for ACK

- ❑ requires countdown timer
- ❑ retransmits if no ACK received in this time
- ❑ if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but use of seq. #'s already handles this
 - receiver must specify seq # of pkt being ACKed

rdt3.0 Sender



rdt3.0: Stop-and-Wait Performance



What is U_{sender} : **utilization** – fraction of time link busy sending?

Assume: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet

Performance of rdt3.0

- ❑ rdt3.0 works, but performance stinks
- ❑ Example: 1 Gbps link, 15 ms e-e prop. delay, 1KB packet:

$$T_{\text{transmit}} = \frac{L \text{ (packet length in bits)}}{R \text{ (transmission rate, bps)}} = \frac{8\text{kb/pkt}}{10^{**9} \text{ b/sec}} = 8 \text{ microsec}$$

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 1KB pkt every 30 msec -> 33kB/sec throughput over 1 Gbps link
- network protocol limits use of physical resources !


A Summary of Questions

- ❑ How to improve the performance of rdt3.0?
- ❑ What if there are reordering and duplication?
- ❑ How to determine the “right” timeout value?

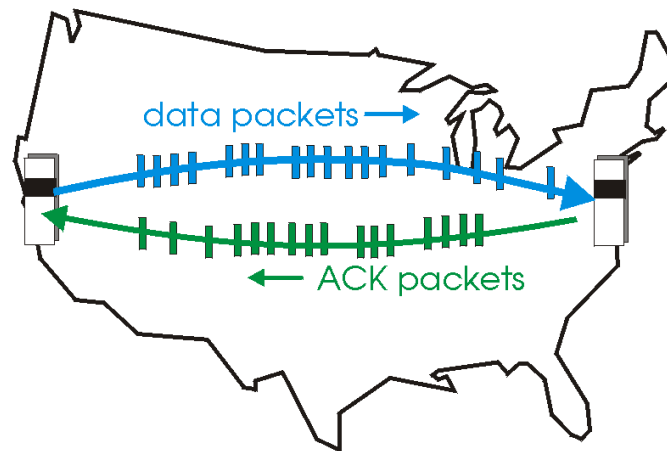
Sliding Window Protocols: Pipelining

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver

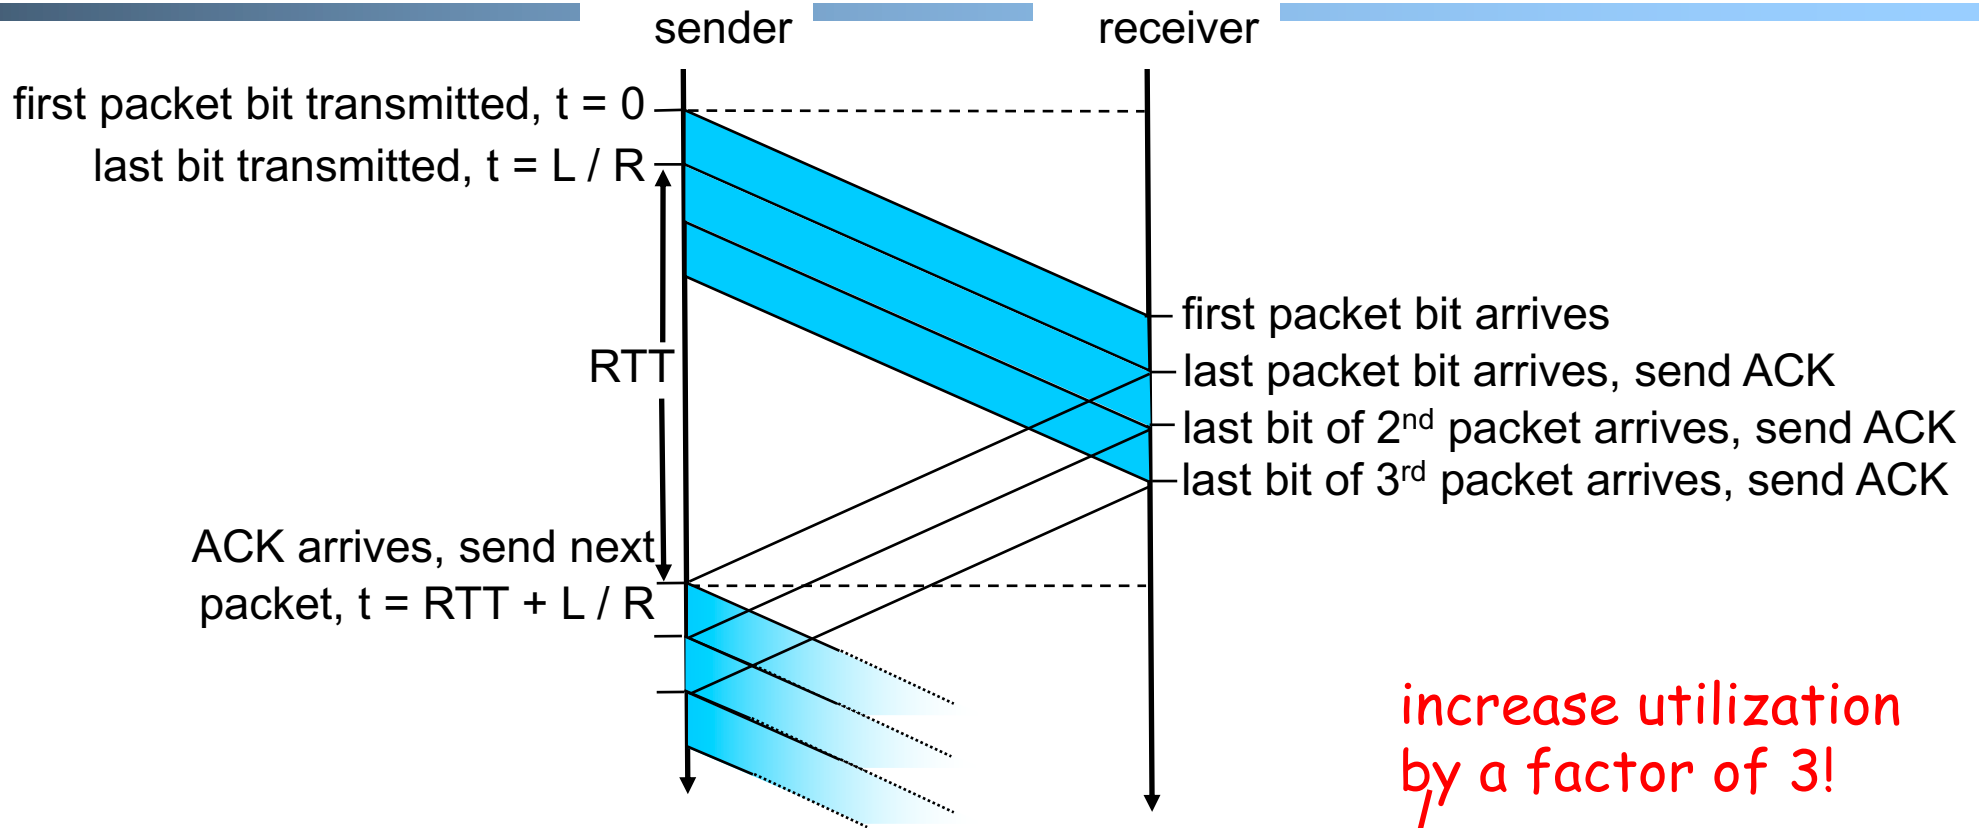


(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation

Pipelining: Increased Utilization




$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

Question: a rule-of-thumb window size?

Realizing Sliding Window: Go-Back-n

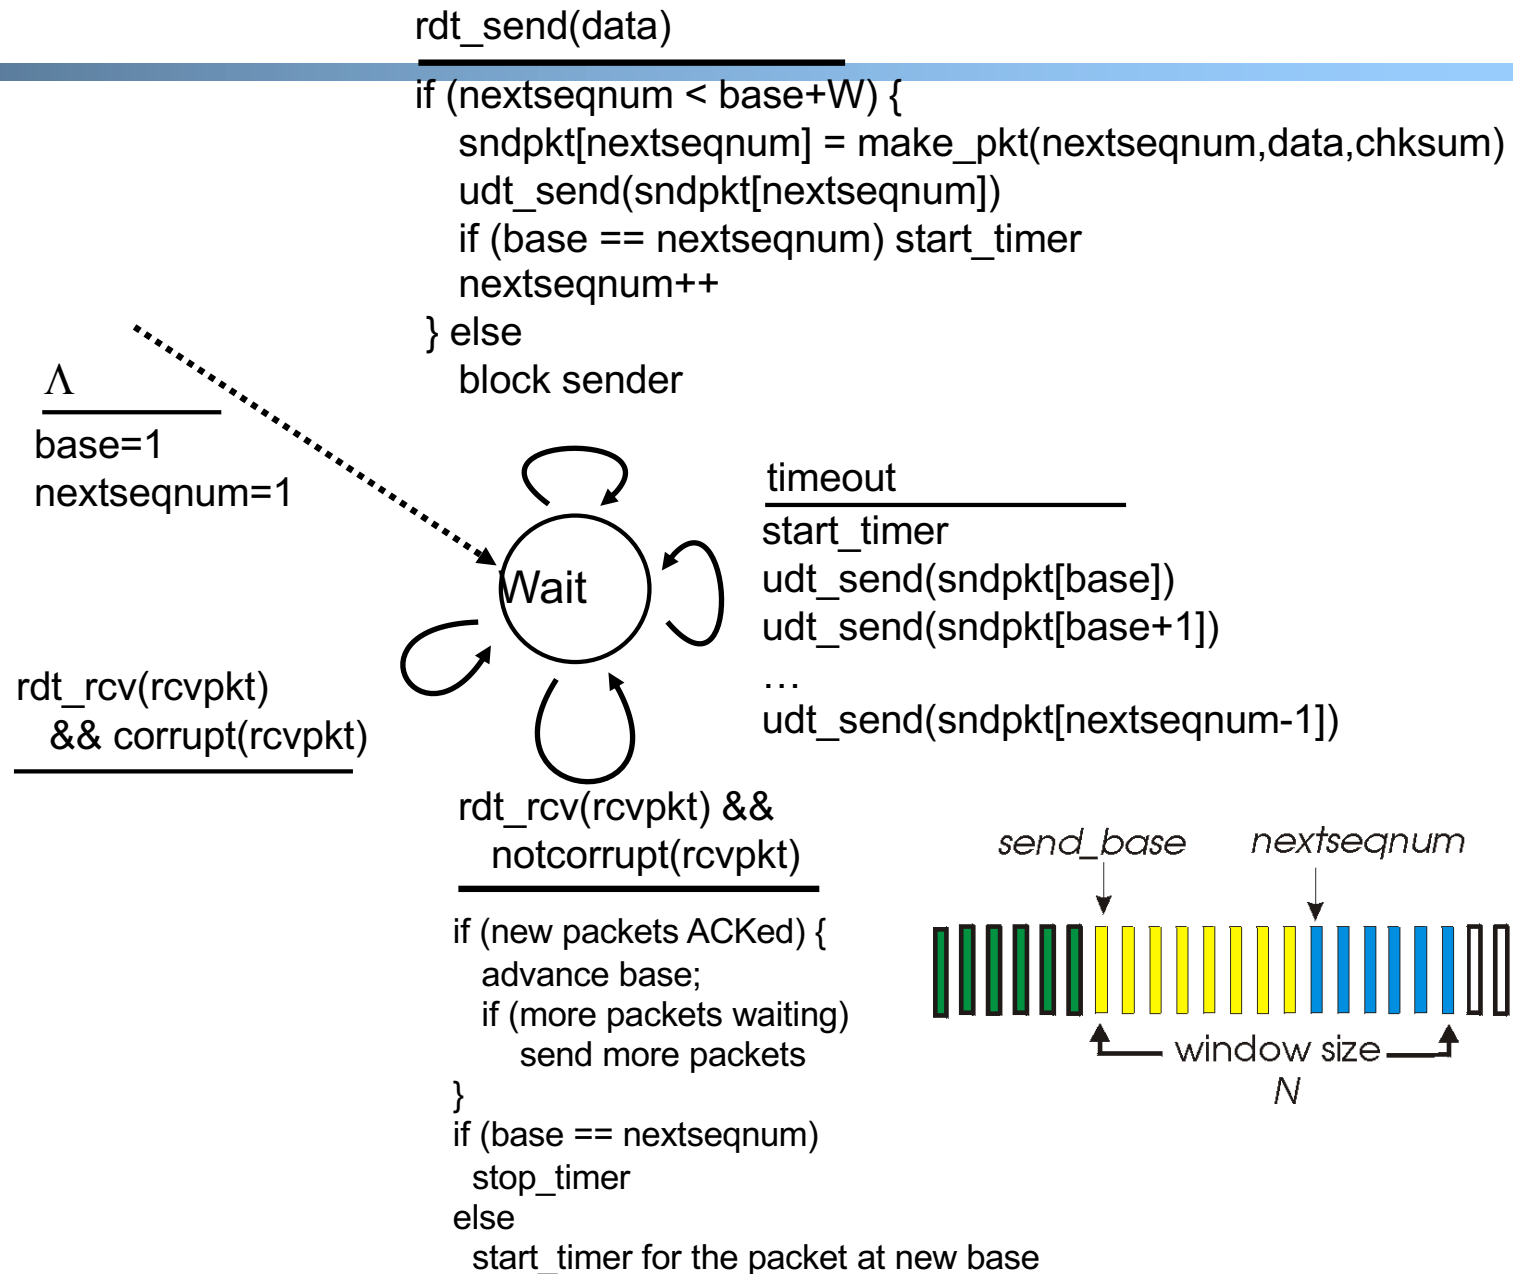
Sender:

- ❑ k-bit seq # in pkt header
- ❑ “window” of up to W , consecutive unack’ed pkts allowed

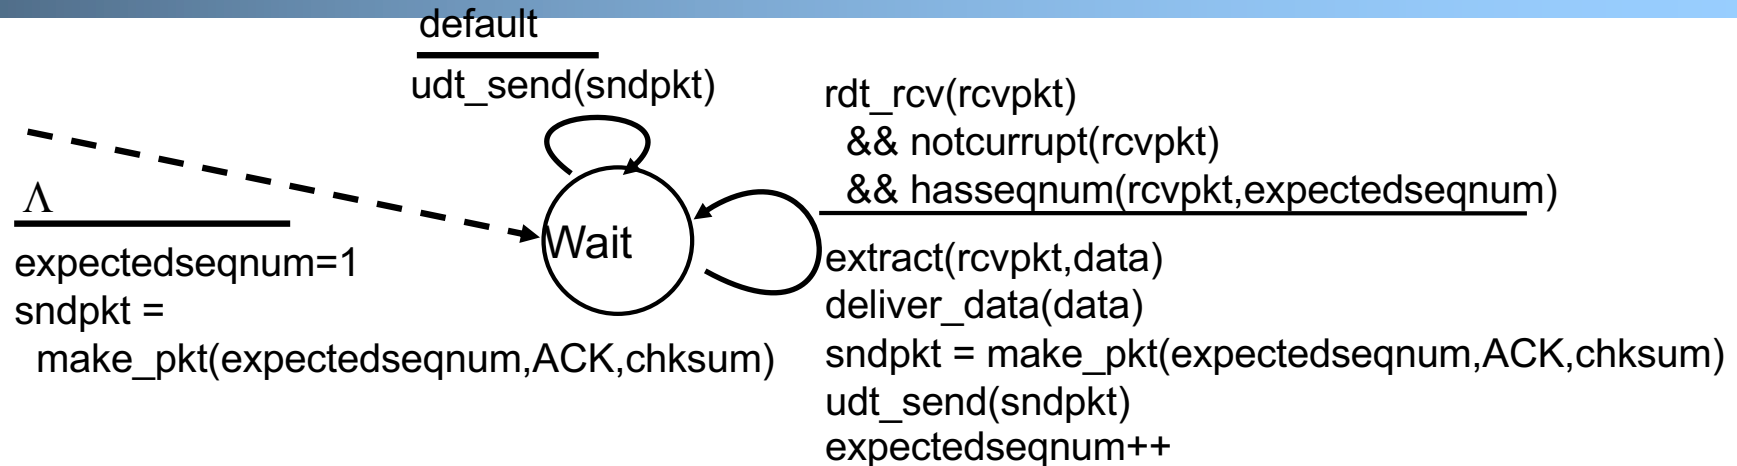


- ❑ **ACK(n): ACKs all pkts up to, including seq # n - “cumulative ACK”**
 - note: ACK(n) could mean two things: I have received **upto and include** n, or I am waiting for n
- ❑ timer for the packet at base
- ❑ *timeout(n)*: retransmit pkt n and all higher seq # pkts in window

GBN: Sender FSM



GBN: Receiver FSM



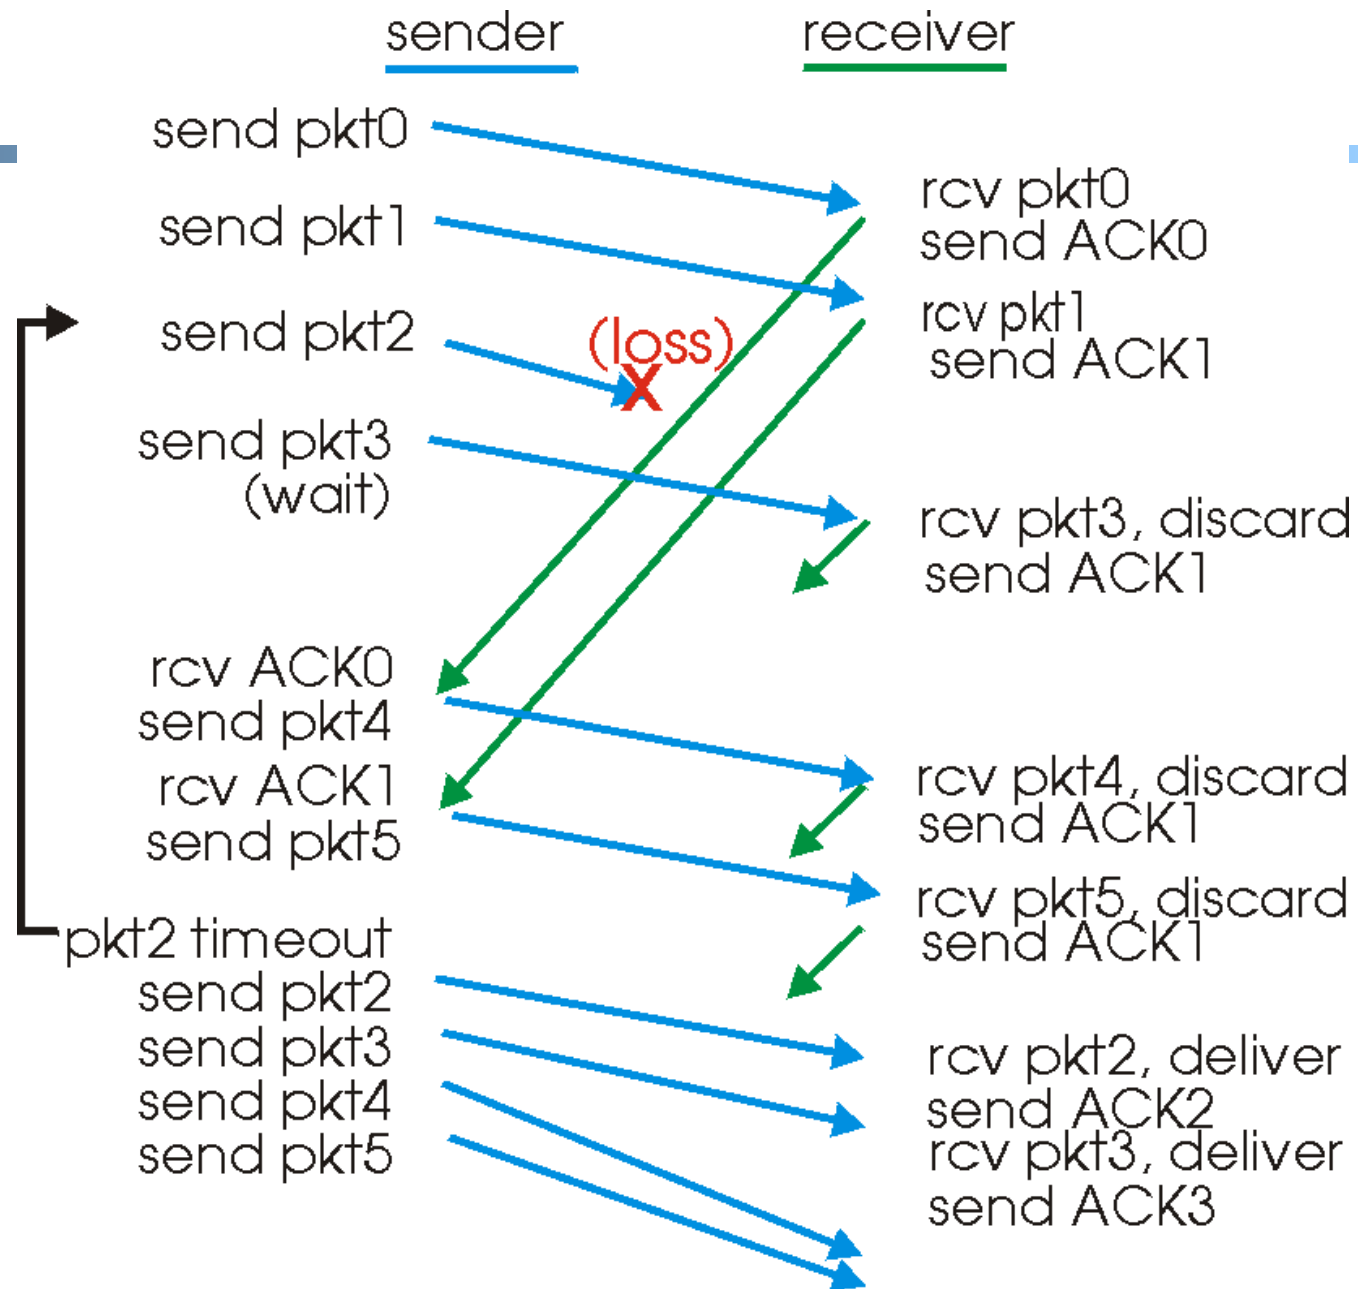
Only state: **expectedseqnum**

□ out-of-order pkt:

- discard (don't buffer) -> **no receiver buffering!**
- re-ACK pkt with highest in-order seq #
- may generate duplicate ACKs

GBN in Action

window
size = 4



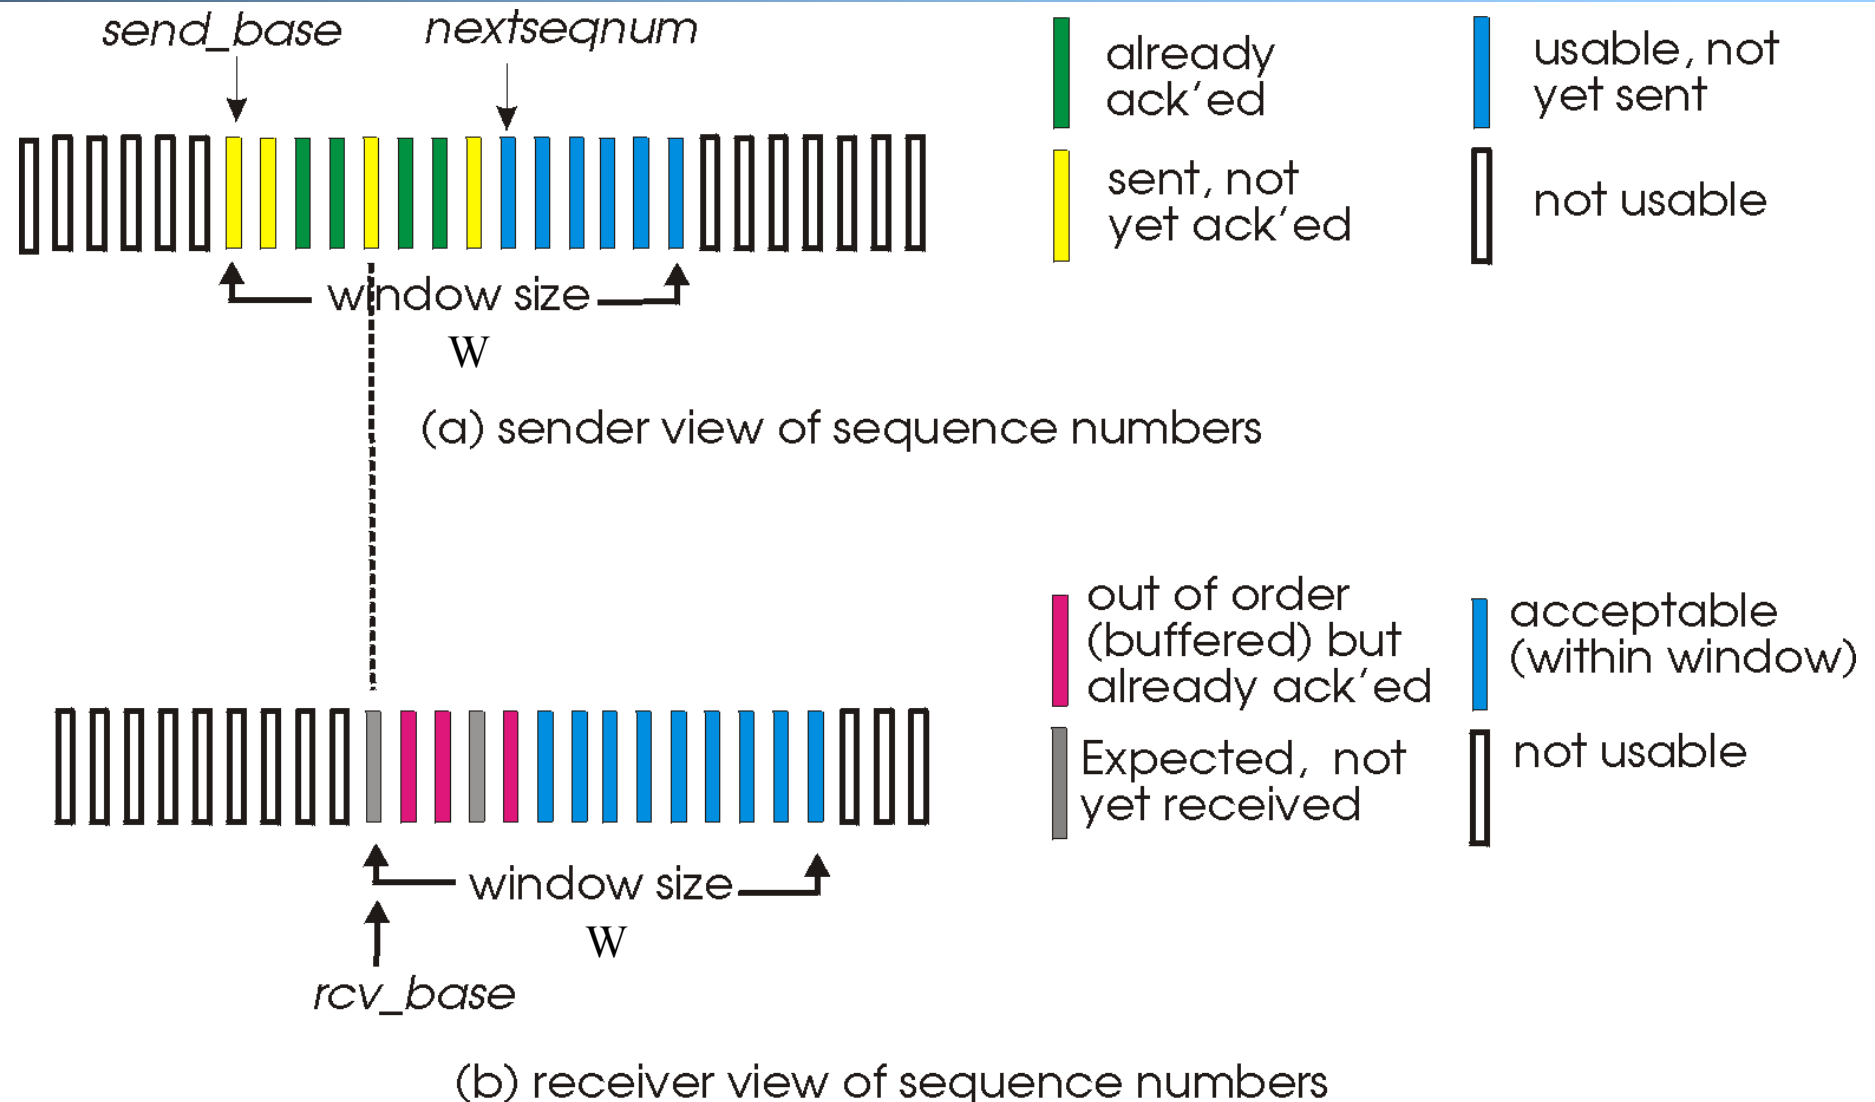
Analysis: Efficiency of Go-Back-n

- ❑ Assume window size W
- ❑ Assume each packet is lost with probability p
- ❑ On average, how many packets do we send for each data packet received?

Selective Repeat

- ❑ Sender window
 - Window size W : W consecutive unACKed seq #'s
- ❑ Receiver *individually* acknowledges correctly received pkts
 - *buffers out-of-order* pkts, for eventual in-order delivery to upper layer
 - *ACK(n) means received packet with seq# n only*
 - buffer size at receiver: window size
- ❑ Sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt

Selective Repeat: Sender, Receiver Windows



Selective Repeat

sender

data from above :

- unACKed packets is less than window size W , send; otherwise block app.

timeout(n):

- resend pkt n , restart timer

ACK(n) in $[\text{sendbase}, \text{sendbase} + W - 1]$:

- mark pkt n as received
- update sendbase to the first packet unACKed

receiver

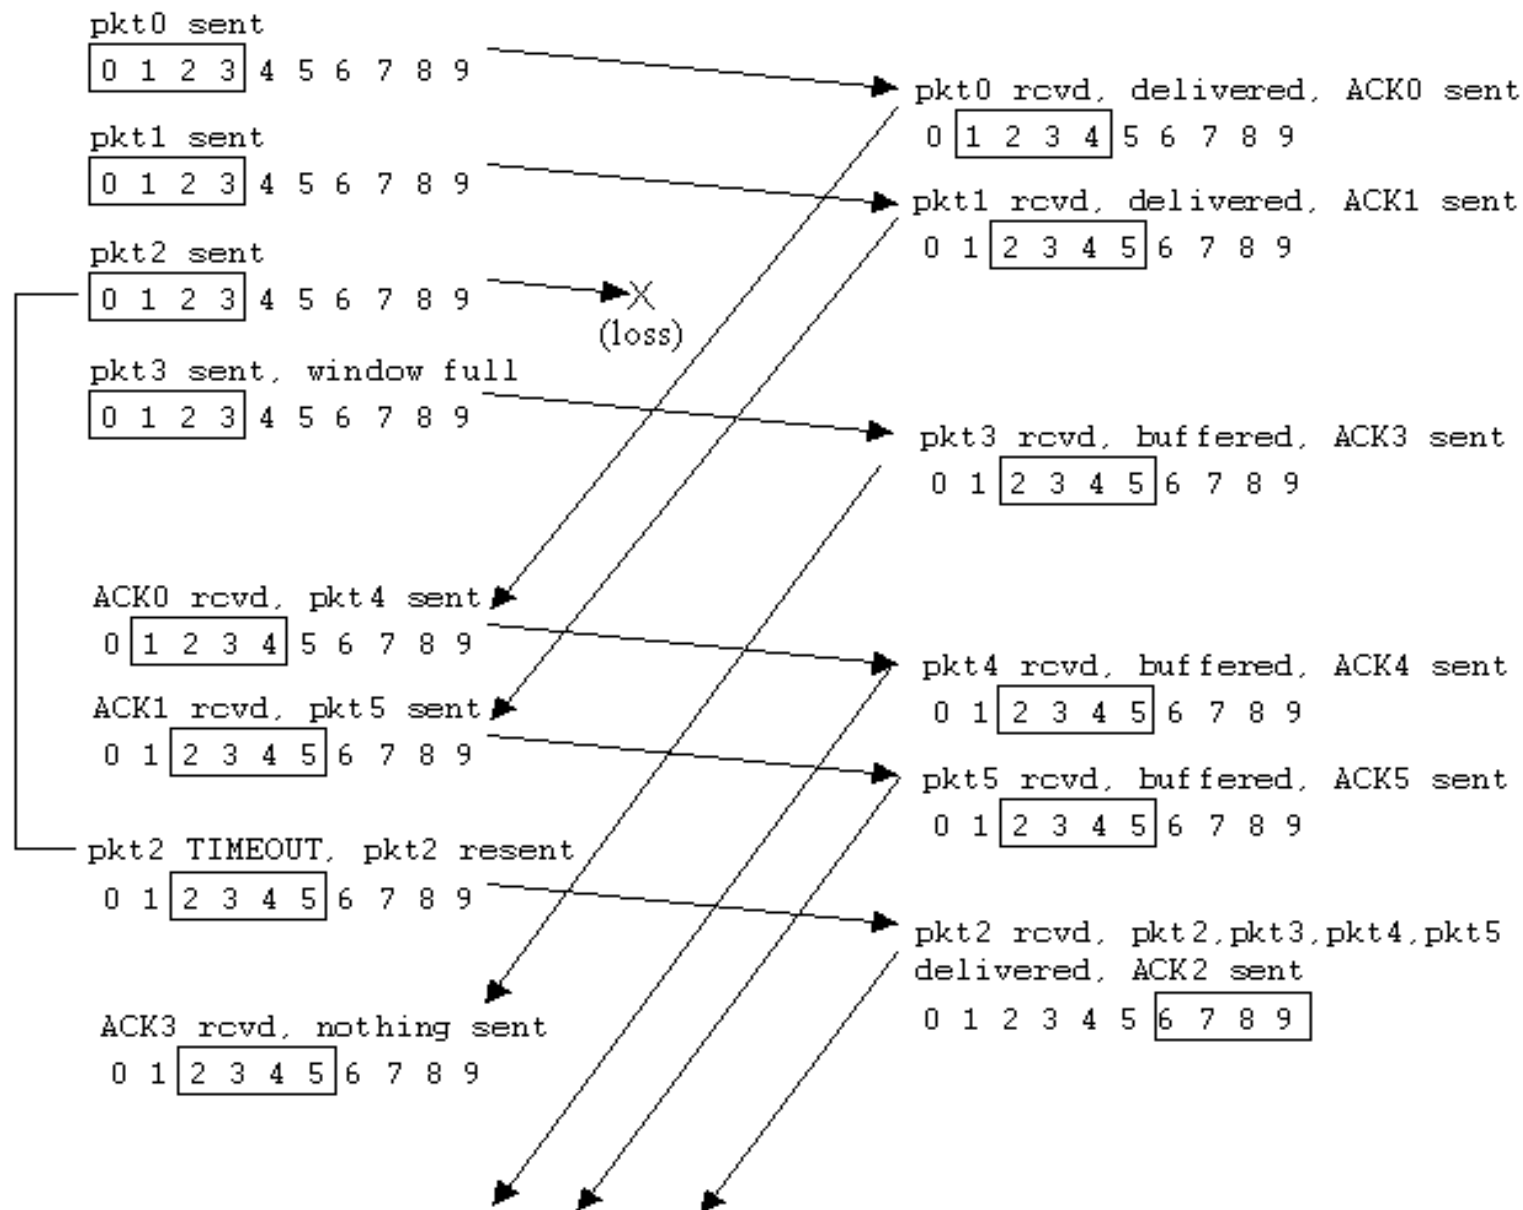
pkt n in $[\text{rcvbase}, \text{rcvbase} + W - 1]$

- send ACK(n)
- if (out-of-order)
 mark and buffer pkt n
 else /*in-order*/
 deliver any in-order packets

otherwise:

- ignore

Selective Repeat in Action




Discussion: Efficiency of Selective Repeat

- ❑ Assume window size W
- ❑ Assume each packet is lost with probability p
- ❑ On average, how many packets do we send for each data packet received?

Selective Repeat: Seq# Ambiguity


Example:

- ❑ seq #'s: 0, 1, 2, 3
- ❑ window size=3
- ❑ Error: incorrectly passes duplicate data as new.




State Invariant: Window Location

□ Go-back-n (GBN)




□ Selective repeat (SR)




Window Location

Q: what relationship between seq # size and window size?

□ Go-back-n (GBN)



□ Selective repeat (SR)



Selective Repeat

sender

data from above :

- unACKed packets is less than window size W , send; otherwise block app.

timeout(n):

- resend pkt n , restart timer

ACK(n) in $[\text{sendbase}, \text{sendbase} + W - 1]$:

- mark pkt n as received
- update sendbase to the first packet unACKed

receiver

pkt n in $[\text{rcvbase}, \text{rcvbase} + W - 1]$

- send ACK(n)
- if (out-of-order)
mark and buffer pkt n
else /*in-order*/
deliver any in-order packets

pkt n in $[\text{rcvbase} - W, \text{rcvbase} - 1]$

- send ACK(n)

otherwise:

- ignore

Sliding Window Protocols: Go-back-n and Selective Repeat

	Go-back-n	Selective Repeat
data bandwidth: sender to receiver (avg. number of times a pkt is transmitted)	Less efficient $\frac{1-p+pw}{1-p}$	More efficient $\frac{1}{1-p}$
ACK bandwidth (receiver to sender)	More efficient	Less efficient
Relationship between M (the number of seq#) and W (window size)	$M > W$	$M \geq 2W$
Buffer size at receiver	1	W
Complexity	Simpler	More complex

p: the loss rate of a packet; M: number of seq# (e.g., 3 bit M = 8); W: window size