
Network Layer: Distance Vector Protocols Variations Link-State Protocol

Qiao Xiang

<https://qiaoxiang.me/courses/cnns-xmuf22/index.shtml>

11/22/2022

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - Distance vector protocols (distributed computing)
 - Link state protocols (distributed state synchronization)

Admin

- ❑ Guest lectures are important 😊
- ❑ Upcoming guest lectures
 - December 1, Dr. Linghe Kong@SJTU, Internet of Things
 - December 8, Dr. Zaoxing (Alan) Liu@Boston Univ., Programmable Networks
 - December 15, Dr. Zhenhua Li@THU, 5G Network
- ❑ Schedule is tentative


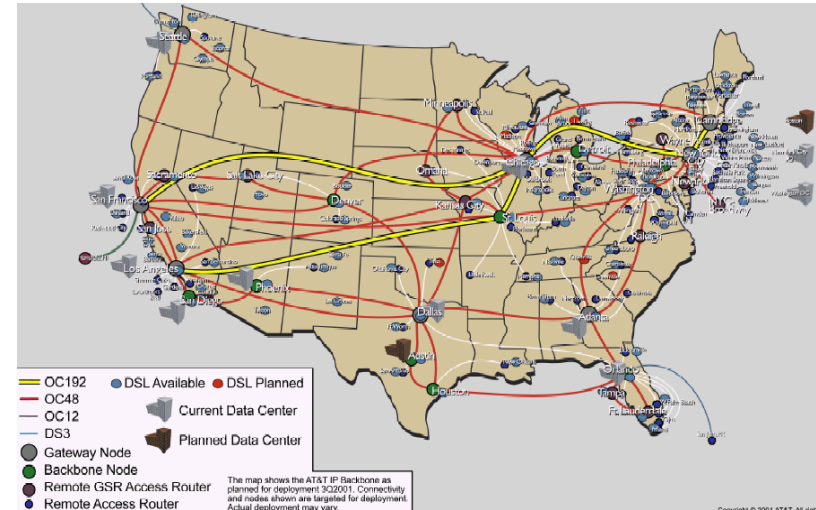
Recap: Routing Context

Routing

Goal: determine “good” paths (sequences of routers) thru networks from source to dest.


Often depends on a graph abstraction:

- graph nodes are routers
- graph edges are physical links
 - links have properties: delay, capacity, \$ cost, **policy**



Recap: Routing Design Space

- ❑ Routing has a large design space
 - who decides routing?
 - source routing: end hosts make decision
 - network routing: networks make decision
 - how many paths from source s to destination d ?
 - multi-path routing
 - single path routing
 - what does routing compute?
 - network cost minimization (shortest path routing)
 - QoS aware
 - will routing adapt to network traffic demand?
 - adaptive routing
 - static routing
 - ...




Recap: Distance Vector Routing: Basic Idea (Bellman-Ford Alg)

- At node i , the basic update rule

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

where

- d_i denotes the distance estimation from i to the destination,
- $N(i)$ is set of neighbors of node i , and
- d_{ij} is the distance of the direct link from i to j




Recap: Synchronous Bellman-Ford (SBF)


□ Nodes update in rounds:

- there is a global clock;
- at the beginning of each round, each node sends its estimate to all of its neighbors;
- at the end of the round, updates its estimation

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$



Recap: SBF/ ∞




□ Initialization (time 0):

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ \infty & \text{otherwise} \end{cases}$$

Example

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$



Consider D as destination; $d(t)$ is a vector consisting of estimation of each node at round t

	A	B	C	E	D
$d(0)$	∞	∞	∞	∞	0
$d(1)$	∞	∞	2	2	0
$d(2)$	12	3	2	2	0
$d(3)$	10	3	2	2	0
$d(4)$	10	3	2	2	0

Observation: $d(0) \geq d(1) \geq d(2) \geq d(3) \geq d(4) = d^*$

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

A Nice Property of SBF: Monotonicity

- Consider two configurations $d(t)$ and $d'(t)$
- If $d(t) \geq d'(t)$
 - i.e., each node has a higher estimate in one scenario (d) than in another scenario (d'),
- then $d(t+1) \geq d'(t+1)$
 - i.e., each node has a higher estimate in d than in d' after one round of synchronous update.

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

Correctness of SBF/ ∞

□ Claim: $d_i(h)$ is the length $L_i(h)$ of a shortest path from i to the destination using $\leq h$ hops

- base case: $h = 0$ is trivially true
- assume true for $\leq h$,
i.e., $L_i(h) = d_i(h)$, $L_i(h-1) = d_i(h-1)$, ...

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$

Correctness of SBF/ ∞

□ consider $\leq h+1$ hops:

$$\begin{aligned} L_i(h+1) &= \min(L_i(h), \min_{j \in N(i)} (d_{ij} + L_j(h))) \\ &= \min(d_i(h), \min_{j \in N(i)} (d_{ij} + d_j(h))) \\ &= \min(d_i(h), d_i(h+1)) \end{aligned}$$

since $d_i(h) \leq d_i(h-1)$


$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h)) \leq \min_{j \in N(i)} (d_{ij} + d_j(h-1)) = d_i(h)$$

$$L_i(h+1) = d_i(h+1)$$

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - *Distributed distance vector protocols*
 - *synchronous Bellman-Ford (SBF)*
 - SBF/ ∞
 - SBF/-1 SBF/ ∞

SBF at another Initial Configuration: SBF/-1




□ Initialization (time 0):

$$d_i(0) = \begin{cases} 0 & i = \text{dest} \\ -1 & \text{otherwise} \end{cases}$$

Example

$$d_i(h+1) = \min_{j \in N(i)} (d_{ij} + d_j(h))$$



Consider D as destination

	A	B	C	E	D
d(0)	-1	-1	-1	-1	0
d(1)	6	0	0	2	0
d(2)	7	1	1	2	0
d(3)	8	2	2	2	0
d(4)	9	3	3	2	0
d(5)	10	3	3	2	0
d(6)	10	3	3	2	0

Observation: $d(0) \leq d(1) \leq d(2) \leq d(3) \leq d(4) \leq d(5) = d(6) = d^*$

Correctness of SBF/-1

- ❑ SBF/-1 converges due to monotonicity
- ❑ Remaining question:
 - Can we guarantee that SBF/-1 converges to shortest path?

Correctness of SBF/-1

- Common between SBF/ ∞ and SBF/-1: they solve the Bellman equation

$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

where $d_D = 0$.

- We have proven SBF/ ∞ is the shortest path solution.
- SBF/-1 computes shortest path if Bellman equation has a unique solution.


$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

Uniqueness of Solution to BE

□ Assume another solution d , we will show that $d = d^*$

case 1: we show $d \geq d^*$

Since d is a solution to BE, we can construct paths as follows: for each i , pick a j which satisfies the equation; since d^* is shortest, $d \geq d^*$



$$d_i = \min_{j \in N(i)} (d_{ij} + d_j)$$

Uniqueness of Solution to BE

Case 2: we show $d \leq d^*$

assume we run SBF with two initial configurations:

- one is d
- another is $SBF/\infty (d^\infty)$,

-> monotonicity and convergence of SBF/∞ imply that $d \leq d^*$

Discussion

- ❑ Will SBF converge under other non-negative initial conditions?
- ❑ Problems of running *synchronous* BF?

Outline

- ❑ Admin and recap
- ❑ Network overview
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - *Distributed distance vector protocols*
 - synchronous Bellman-Ford (SBF)
 - *asynchronous Bellman-Ford (ABF)*

Asynchronous Bellman-Ford (ABF)

- No notion of global iterations
 - each node updates at its own pace
- Asynchronously each node i computes


$$d_i = \min_{j \in N(i)} (d_{ij} + d_j^i)$$

using last received value d_j^i from neighbor j .

- Asynchronously node j sends its estimate to its neighbor i :
 - We assume that there is an upper bound on the delay of estimate packet

ABF: Example

Below is just one step! The protocol repeats forever!




		distance tables from neighbors			computation			E's distance table	distance table E sends to its neighbors
$d_E()$		A	B	D	A	B	D		
destinations	A	0	7	∞	10	15	∞	A: 10	A: 10
	B	7	0	∞	17	8	∞	B: 8	B: 8
	C	∞	1	2	∞	9	4	D: 4	C: 4
	D	∞	∞	0	∞	∞	2	D: 2	D: 2
		10	8	2					E: 0

Asynchronous Bellman-Ford (ABF)


- ABF will eventually converge to the shortest path
 - links can go down and come up - but if topology is stabilized after some time t and connected, ABF will eventually converge to the shortest path !

ABF Convergence Proof Complexity: Complex System State



What is system state?

System State



three types of distance state from node j :

- d_j : current distance estimate state at node j
- d_j^i : last d_j that neighbor i received
- d_j^i : those d_j that are still in transit to neighbor i

ABF Convergence Proof: The Sandwich Technique

□ Basic idea:

- bound system state using extreme states

□ Extreme states:

- SBF/∞ ; call the sequence $U()$
- $SBF/-1$; call the sequence $L()$

ABF Convergence

- Consider the time when the topology is stabilized as time 0
- $U(0)$ and $L(0)$ provide upper and lower bounds at time 0 on all corresponding elements of states
 - $L_j(0) \leq d_j \leq U_j(0)$ for all d_j state at node j
 - $L_j(0) \leq d_{ij}^i \leq U_j(0)$
 - $L_j(0) \leq \text{update messages } d_{ij}^i \leq U_j(0)$


ABF Convergence

□ d_j

- after at least one update at node j :
 d_j falls between $L_j(1) \leq d_j \leq U_j(1)$

□ d_j^i :

- eventually all d_j^i that are only bounded by $L_j(0)$ and $U_j(0)$ are replaced with in $L_j(1)$ and $U_j(1)$



Asynchronous Bellman-Ford: Summary

□ Distributed

- each node communicates its routing table to its directly-attached neighbors

□ Iterative

- continues periodically or when link changes, e.g. detects a link failure

□ Asynchronous

- nodes need *not* exchange info/iterate in lock step!

□ Convergence

- in finite steps, independent of initial condition if network is connected

Summary: Distributed Distance-Vector


- Tool box: a key technique for proving convergence (**liveness**) of distributed protocols: **monotonicity** and **bounding-box (sandwich)** design
 - Consider two configurations $d(t)$ and $d'(t)$:
 - if $d(t) \leq d'(t)$, then $d(t+1) \leq d'(t+1)$
 - Identify two extreme configurations to sandwich any real configurations

Outline

- ❑ Admin and recap
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - Distance vector protocols (distributed computing)
 - synchronous Bellman-Ford (SBF)
 - asynchronous Bellman-Ford (ABF)
 - *properties of DV*

Properties of Distance-Vector Algorithms

- Good news propagate fast



Properties of Distance-Vector Algorithms

❑ Bad news propagate slowly


	A	B	C	D	E	
	●	●	●	●	●	
		1	2	3	4	Initially
A-B link down		3	2	3	4	After 1 exchange
		3	4	3	4	After 2 exchanges
		5	4	5	4	After 3 exchanges
		5	6	5	6	After 4 exchanges
		7	6	7	6	After 5 exchanges
		7	8	7	8	After 6 exchanges
		⋮				
		∞	∞	∞	∞	

❑ This is called the *counting-to-infinity* problem

❑ Q: what causes counting-to-infinity?

Counting-To-Infinity is Because of Routing Loop

- Counting-to-infinity is caused by a routing loop, which is a **global state** (consisting of the nodes' local states) at a global moment (observed by an oracle) such that there exist nodes A, B, C, ... E such that A (locally) thinks B as next hop, B thinks C as next hop, ... E thinks A as next hop



Discussion


- ❑ Why avoid routing loops is hard?
- ❑ Any proposals to avoid routing loops?

Outline

- ❑ Admin and recap
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - Distance vector protocols (distributed computing)
 - synchronous Bellman-Ford (SBF)
 - asynchronous Bellman-Ford (ABF)
 - properties of DV
 - DV w/ loop prevention
 - *reverse poison*

The Reverse-Poison (Split-horizon) Hack

If the path to dest is through neighbor h, report ∞ to neighbor h for dest.



		distance tables from neighbors			computation			E's distance table		distance table E sends to its neighbors		
destinations	D ^E ()	A	B	D	A	B	D			To A	To B	To D
	A	0	7	∞	1	15	∞	1, A		A: ∞	A: 1	A: 1
	B	7	0	∞	8	8	∞	8, B		B: 8	B: ∞	B: 8
	C	∞	1	2	∞	9	4	4, D		C: 4	C: 4	C: ∞
	D	∞	∞	0	∞	∞	2	2, D		D: 2	D: 2	D: ∞
		1	8	2						E: 0	E: 0	E: 0
		c(E,A)	c(E,B)	c(E,D)								


distance through neighbor

DV+RP => RIP

(Routing Information Protocol)

- ❑ Included in BSD-UNIX Distribution in 1982
- ❑ Link cost: 1
- ❑ Distance metric: # of hops
- ❑ Distance vectors

- exchanged every 30 sec via Response Message (also called **advertisement**) using UDP
- each advertisement: route to up to 25 destination nets




RIP: Link Failure and Recovery

If no advertisement heard after 180 sec -->
neighbor/link declared dead

- routes via neighbor invalidated
- new advertisements sent to neighbors
- neighbors in turn send out new advertisements (if tables changed)
- link failure info quickly propagates to entire net
- reverse-poison used to prevent **ping-pong** loops
- set infinite distance = 16 hops (why?)


General Routing Loops and Reverse-poison

- Exercise: Can Reverse-poison guarantee no loop for this network?



General Routing Loops and Reverse-poison

- ❑ Reverse-poison removes two-node loops but may not remove more-node loops



- ❑ Unfortunate timing can lead to a loop
 - When the link between C and D fails, C will set its distance to D as ∞
 - A receives the bad news (∞) from C, A will use B to go to D
 - A sends the news to C
 - C sends the news to B


Outline

- ❑ Admin and recap
- ❑ Network control plane
 - Routing
 - Link weights assignment
 - Routing computation
 - Distance vector protocols (distributed computing)
 - synchronous Bellman-Ford (SBF)
 - asynchronous Bellman-Ford (ABF)
 - properties of DV
 - DV w/ loop prevention
 - reverse poison
 - *destination-sequenced DV (DSDV)*

Destination-Sequenced Distance Vector protocol (DSDV)


- ❑ Basic idea: use sequence numbers to partition computation
 - tags each route with a sequence number
 - each destination node D periodically advertises monotonically increasing even-numbered sequence numbers
 - when a node realizes that **the link it uses to reach destination D is broken**, it advertises an **infinite** metric and a **sequence number which is one greater** than the previous route (i.e., an odd seq. number)
 - the route is repaired by a later even-number advertisement from the destination

DSDV: More Detail




- ❑ Let's assume the destination node is D
- ❑ There are optimizations but we present a simple version:
 - each node B maintains (S^B, d^B) , where S^B is the sequence number at B for destination D and d^B is the best distance using a neighbor from B to D
- ❑ Both periodical and triggered updates
 - periodically: D increases its seq. by 2 and broadcasts with $(S^D, 0)$
 - if B is using C as next hop to D and B discovers that C is no longer reachable
 - B increases its sequence number S^B by 1, sets d^B to ∞ , and sends (S^B, d^B) to all neighbors

Example




Will this trigger an update?

Example



Will this trigger an update?

DSDV: Update Alg.



- ❑ Consider simple version, no optimization
- ❑ Update after receiving a message
 - assume B sends to A its current state (S^B, d^B)
 - when A receives (S^B, d^B)
 - if $S^B > S^A$, then
 - // always update if a higher seq#**
 - » $S^A = S^B$
 - » if ($d^B == \infty$) $d^A = \infty$; else $d^A = d^B + d(A, B)$
 - else if $S^A == S^B$, then
 - » if $d^A > d^B + d(A, B)$
 - // update for the same seq# only if better route**
 - $d^A = d^B + d(A, B)$ and uses B as next hop